



Beyond Entitlements for Cloud Native

Scalable Responsibility Management with Spring Boot and Open Policy Agent

Chandra Guntur

August 2020

Bank of New York Mellon

Disclosure

BNY Mellon is the corporate brand of The Bank of New York Mellon Corporation and may be used as a generic term to reference the corporation as a whole and/or its various subsidiaries generally. Products and services may be provided under various brand names in various countries by duly authorized and regulated subsidiaries, affiliates, and joint ventures of The Bank of New York Mellon Corporation. Not all products and services are offered in all countries.

BNY Mellon will not be responsible for updating any information contained within this material and opinions and information contained herein are subject to change without notice.

BNY Mellon assumes no direct or consequential liability for any errors in or reliance upon this material. This material may not be reproduced or disseminated in any form without the express prior written permission of BNY Mellon.

©2020 The Bank of New York Mellon Corporation. All rights reserved.

About :: Chandra Guntur



- Technologist in the financial services industry since 2003 and is programming with Java since 1998.
- Chandra Guntur is a **Director** and **Sr. Principal Architect** in Java Platform Engineering, BNY Mellon.
- BNY Mellon representative in the **Java Community Process (JCP) Executive Committee**.
- Java Champion
- **JUG Leader**, at **NYJavaSIG** (New York Java Special Interest Group) and **NJ Java SIG**.
- Speaker at: **Oracle CodeOne**, **Oracle CodeNY**, **QCon NY**, **Devnexus**, **DawsCon** and **GIDS India**.

Agenda

- Responsibility Management
 - Scenarios
 - Common (existing) solutions for both Data and Logic
 - A Right Solution - described and visualized
- Architecture
 - Technology choices
 - Federated model
 - Distributed model
 - Comparison
- A Sample Policy in action
- IDE Productivity tool

Responsibility Management for the Enterprise

Background and existing solutions

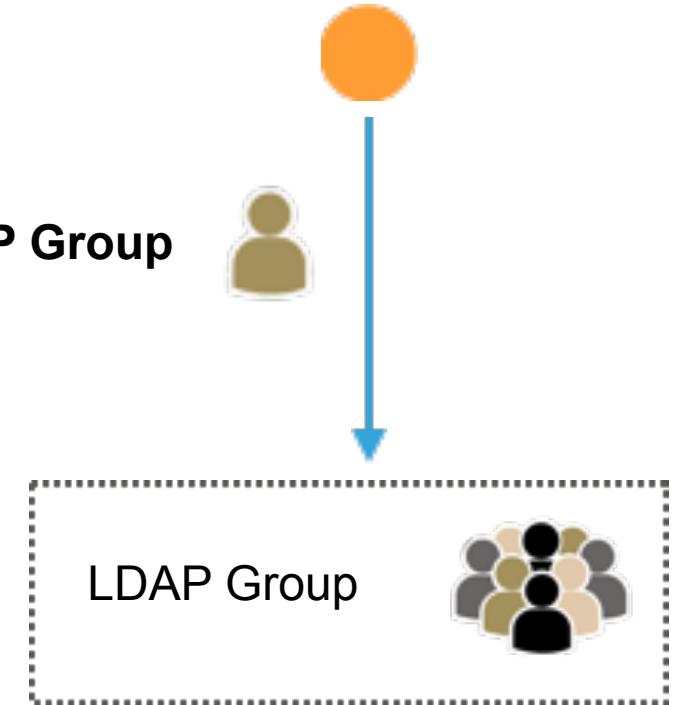


Responsibility Management System

Scenarios

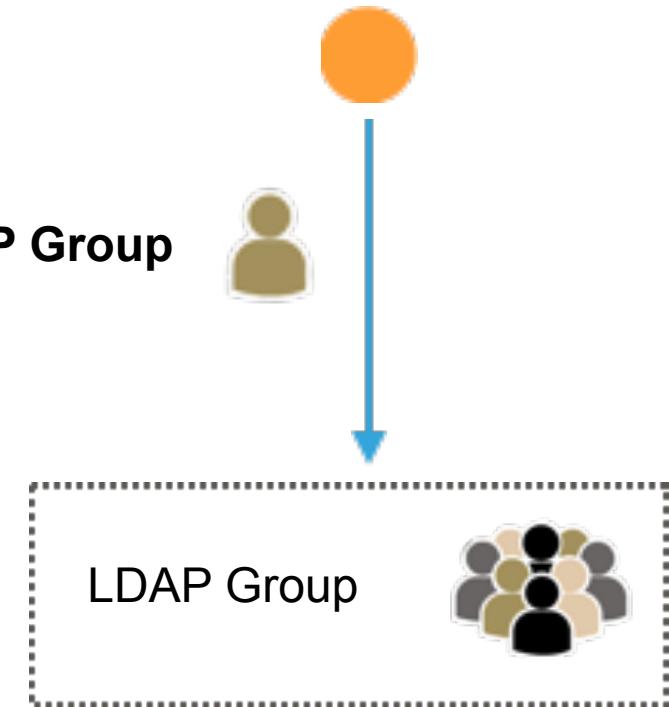
Why Responsibility Management – Scenario 1

- Service A needs to know if a user is a member of an enterprise **LDAP Group**
 - Access may be granted based on membership.
 - Access may be denied based on membership.
 - Access may be granted based on lack of membership.
 - Access may be denied based on lack of membership.



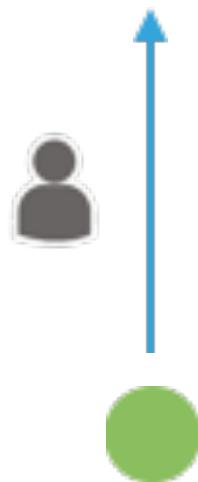
Why Responsibility Management – Scenario 1

- Service A needs to know if a user is a member of an enterprise **LDAP Group**
 - Access may be granted based on membership.
 - Access may be denied based on membership.
 - Access may be granted based on lack of membership.
 - Access may be denied based on lack of membership.



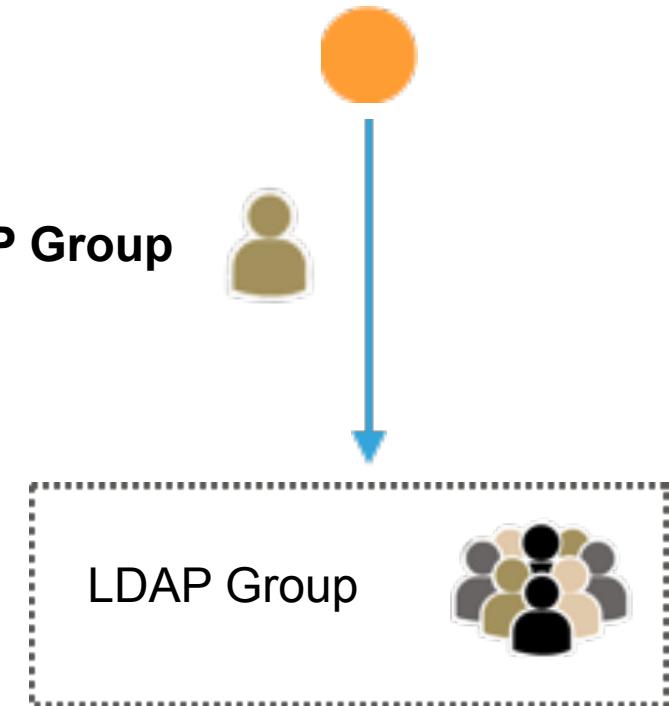
Then ...

- Service B needs to know if a user is a member of an enterprise **LDAP Group**



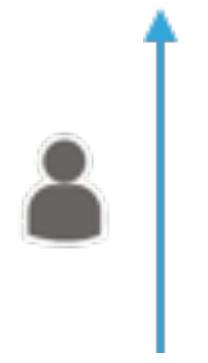
Why Responsibility Management – Scenario 1

- Service A needs to know if a user is a member of an enterprise **LDAP Group**
 - Access may be granted based on membership.
 - Access may be denied based on membership.
 - Access may be granted based on lack of membership.
 - Access may be denied based on lack of membership.



Then ...

- Service B needs to know if a user is a member of an enterprise **LDAP Group**



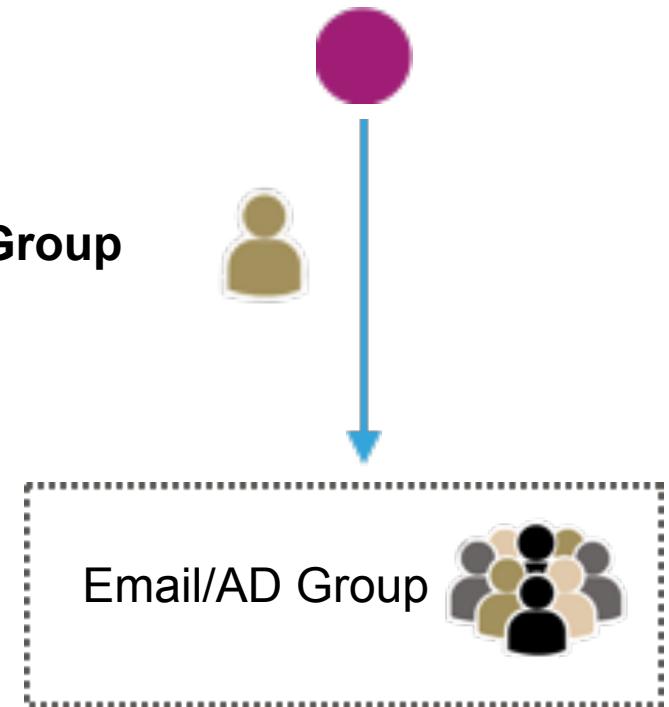
Questions

- How about Service/Application C, D or E ?
- Who manages people who **move/leave/join the department/org/company** (Movers/Leavers/Joiners)



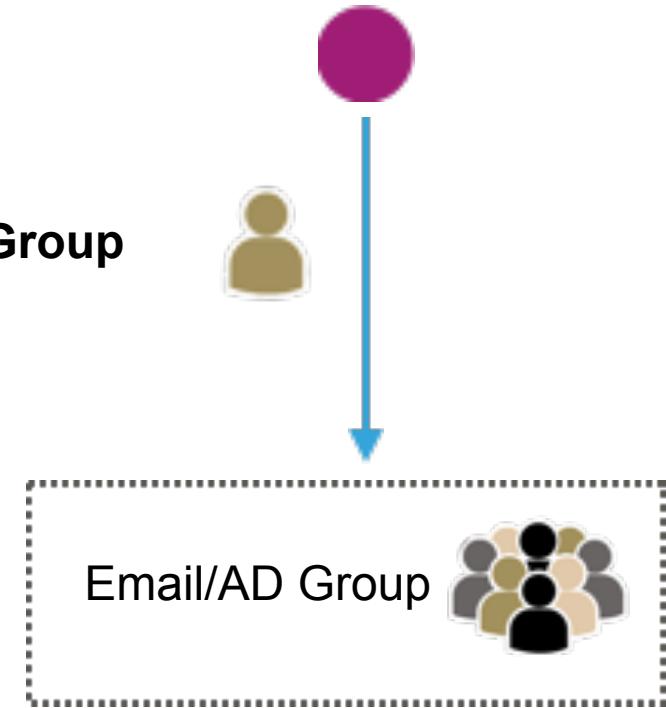
Why Responsibility Management – Scenario 2

- Service M needs to know if a user is a member of an enterprise **AD Group**
 - Access may be granted based on membership.
 - Access may be denied based on membership.
 - Access may be granted based on lack of membership.
 - Access may be denied based on lack of membership.



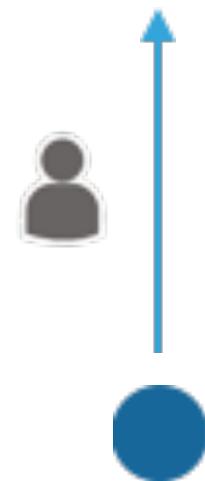
Why Responsibility Management – Scenario 2

- Service M needs to know if a user is a member of an enterprise **AD Group**
 - Access may be granted based on membership.
 - Access may be denied based on membership.
 - Access may be granted based on lack of membership.
 - Access may be denied based on lack of membership.



Then ...

- Service N needs to know if a user is a member of an enterprise **AD Group**



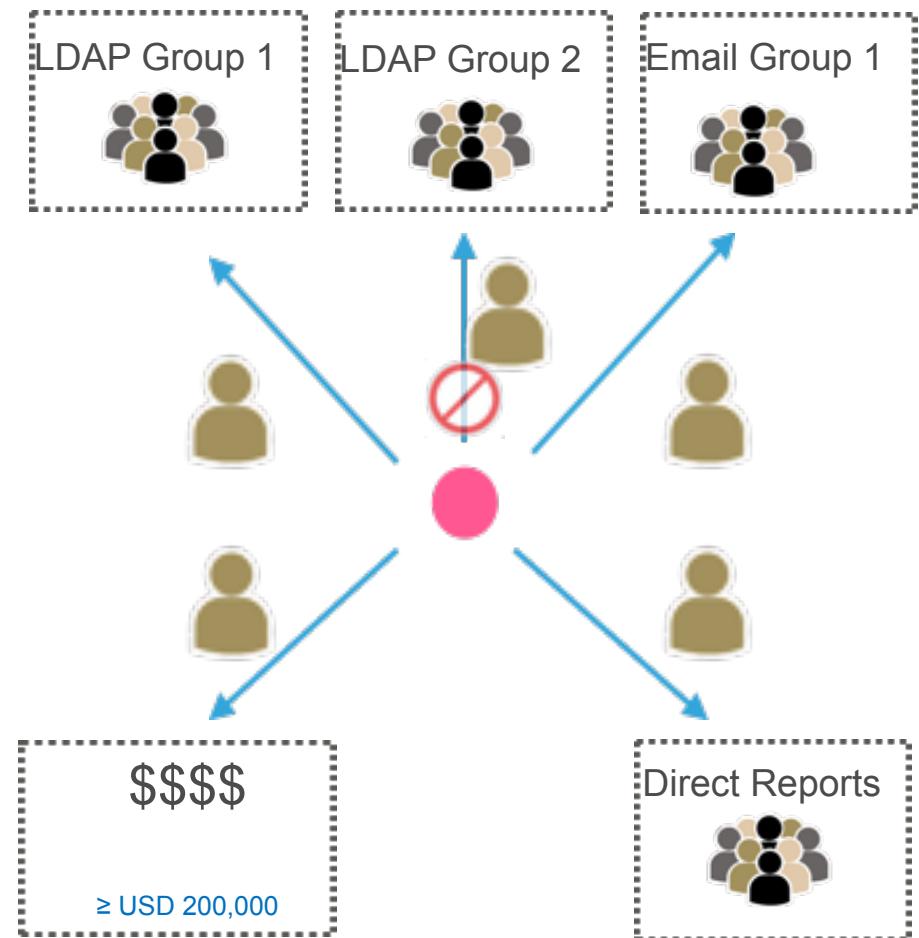
Questions

- How about Service/Application O, P or Q with different groups?
- Who manages people who **move/leave/join the department/org/company** (Movers/Leavers/Joiners)

Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

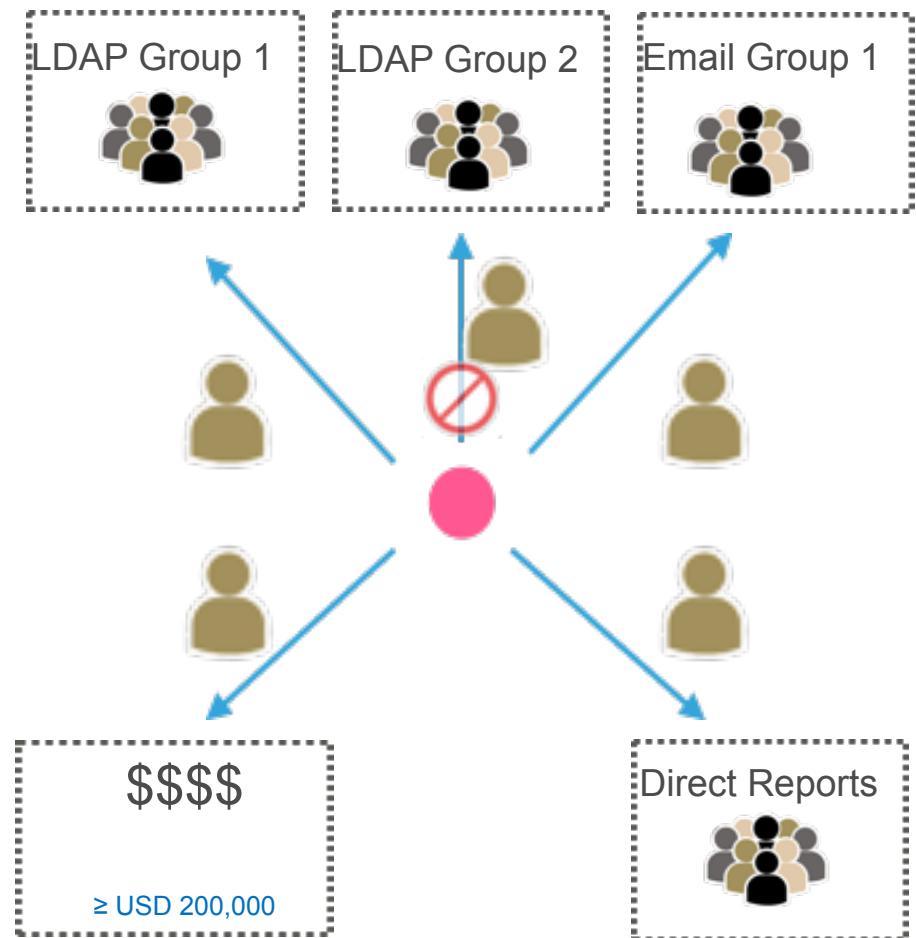


Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

- is member of LDAP Group 1

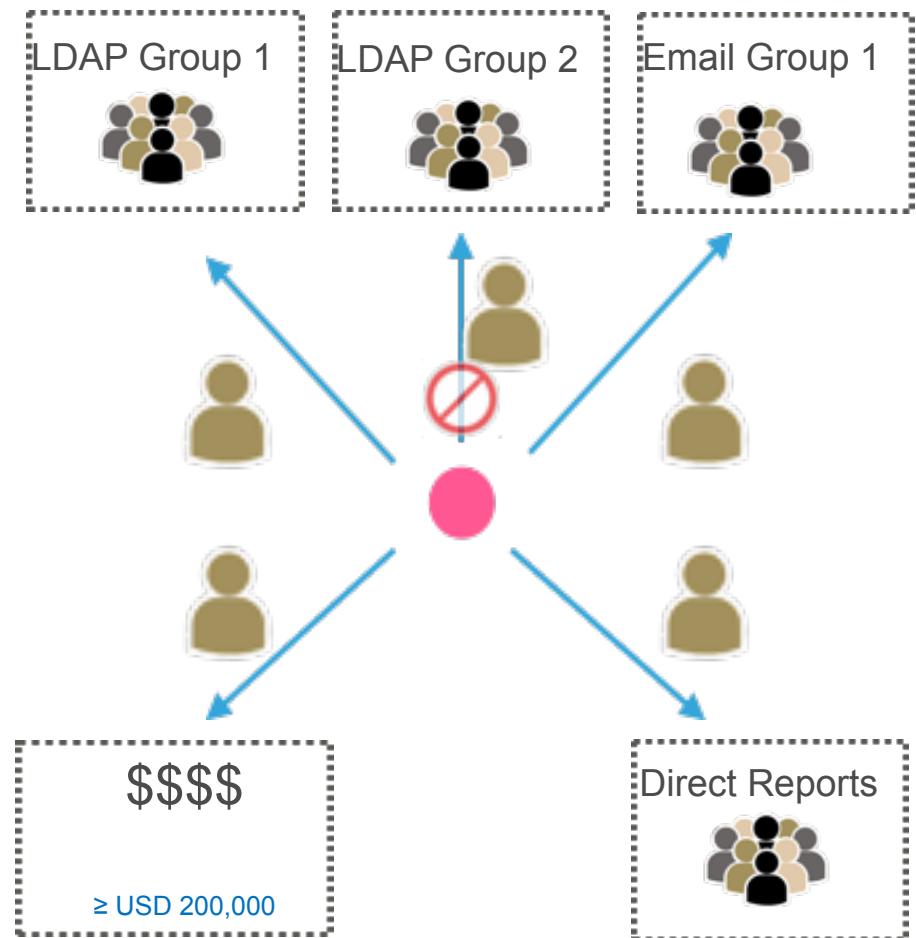


Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

- is member of LDAP Group 1
- is **not** member of LDAP Group 2

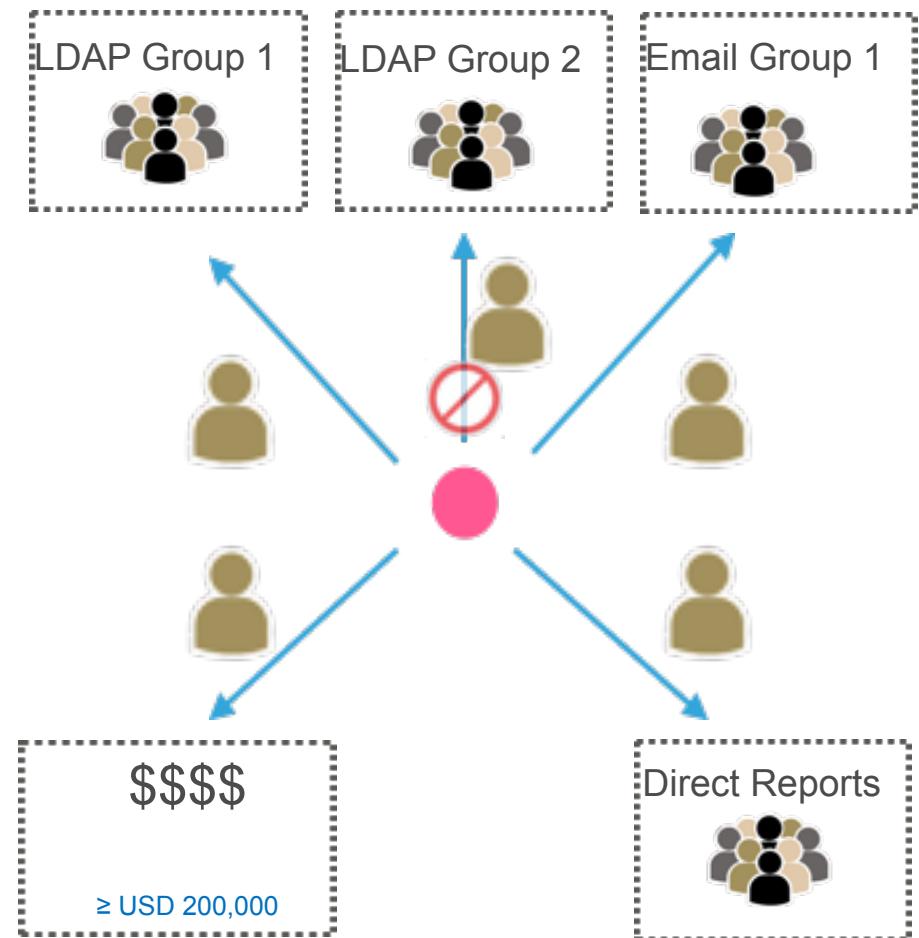


Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

- is member of LDAP Group 1
- is **not** member of LDAP Group 2
- is member of Email Group 1

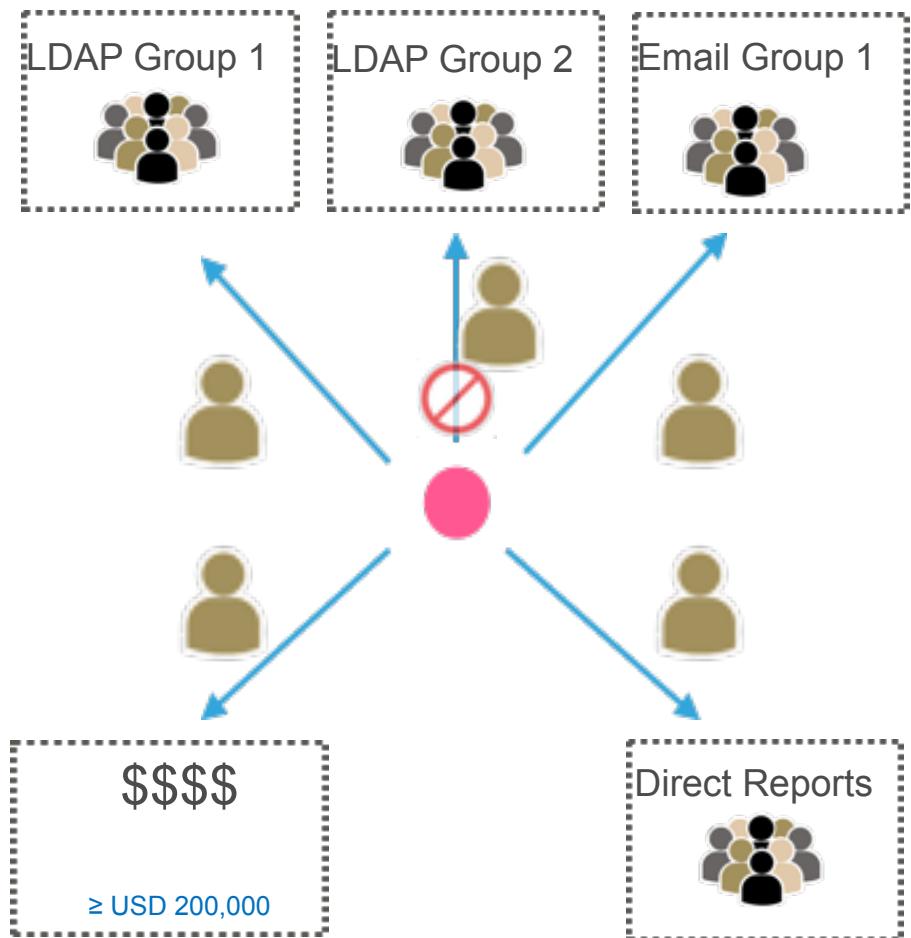


Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

- is member of LDAP Group 1
- is **not** member of LDAP Group 2
- is member of Email Group 1
- is allowed procurement amount USD 200,000

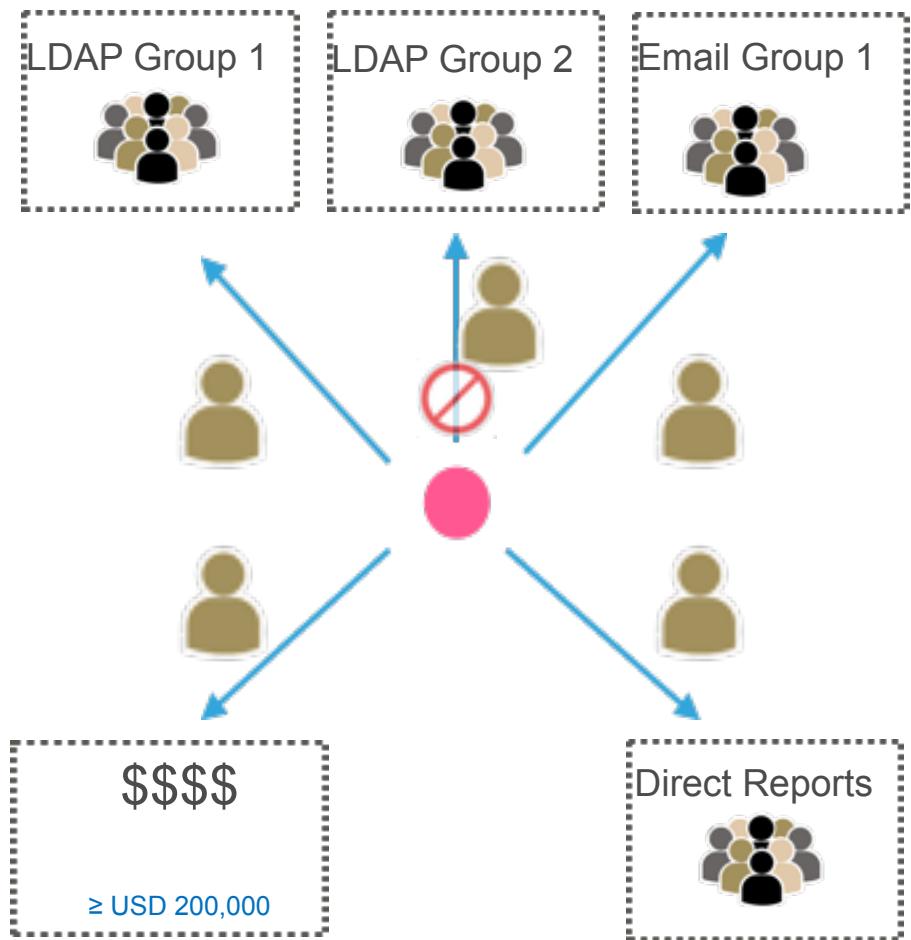


Why Responsibility Management – Scenario 3

More complex evaluations occur as well.

Service X needs to check if all of the below are true for a user:

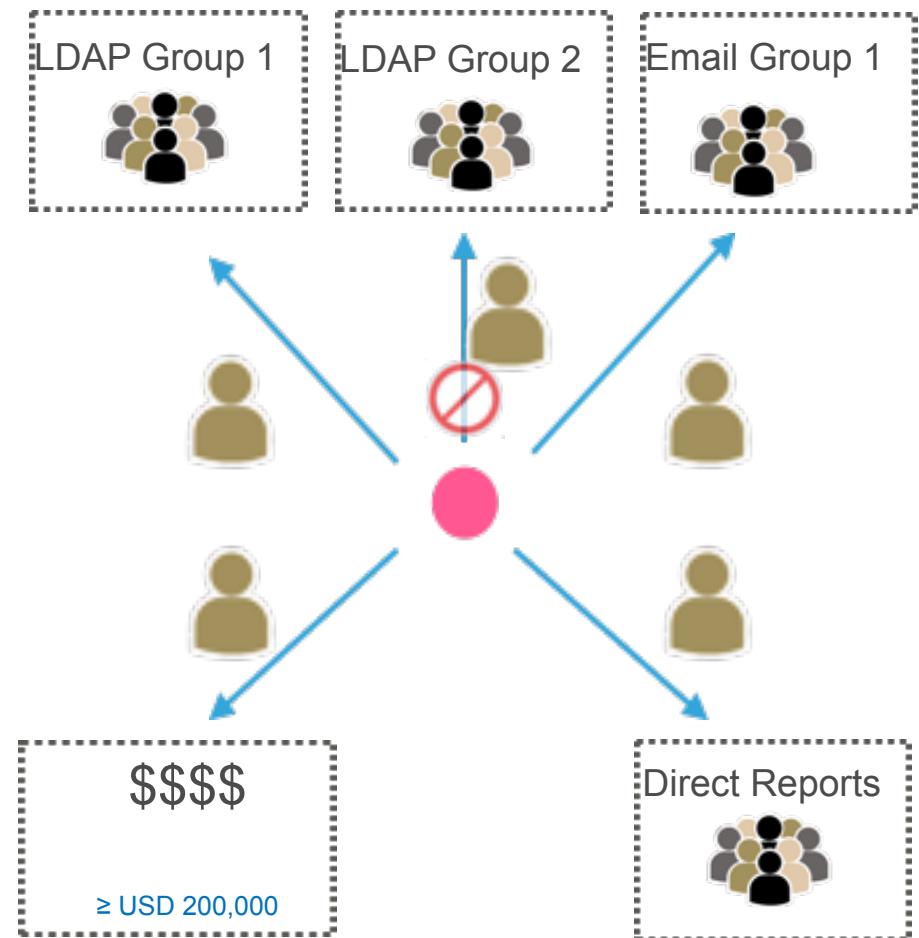
- is member of LDAP Group 1
- is **not** member of LDAP Group 2
- is member of Email Group 1
- is allowed procurement amount USD 200,000
- has at least two direct reports



Why Responsibility Management – Scenario 3 - Continued

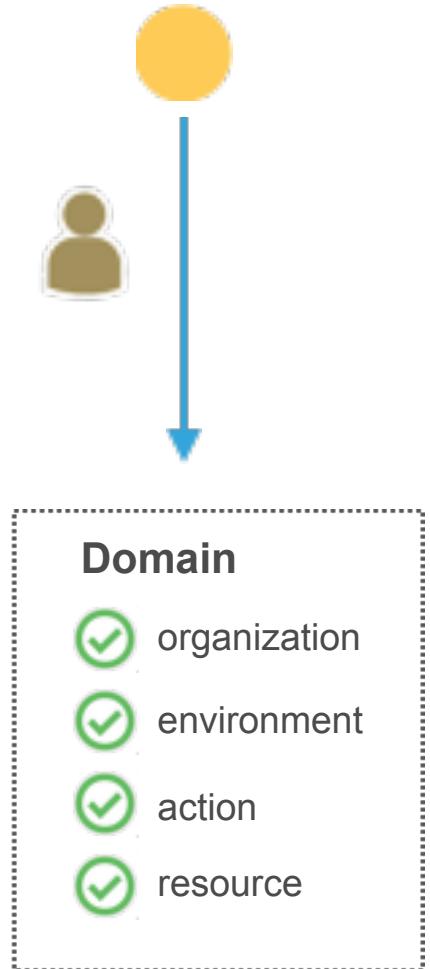
Questions

- What if each request is for different sets of groups and/or amounts?
- What if other services have similar functional constraints with different values?
- Where are such “policies” maintained?
Are they auditable?
Follow SDLC and Config Management?
- Who manages Mover/Leaver/Joiner logic?



Why Responsibility Management – Scenario 4

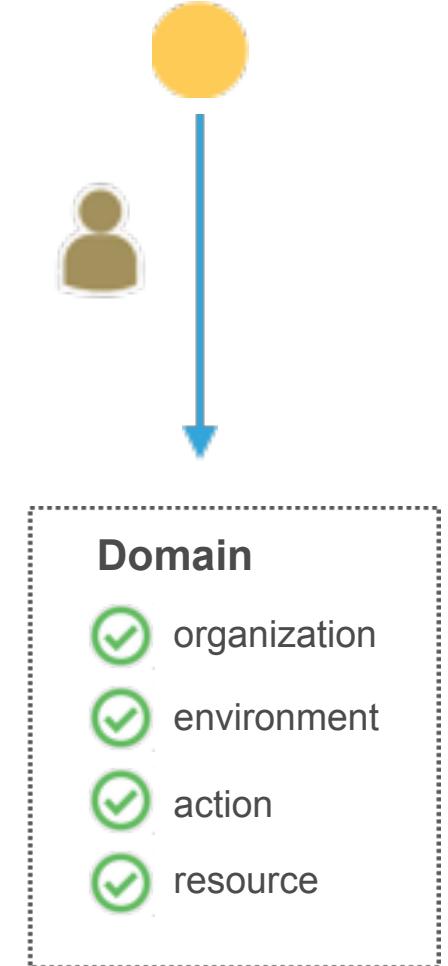
Service Y needs to check responsibility privileges for a user/subject:



Why Responsibility Management – Scenario 4

Service Y needs to check responsibility privileges for a user/subject:

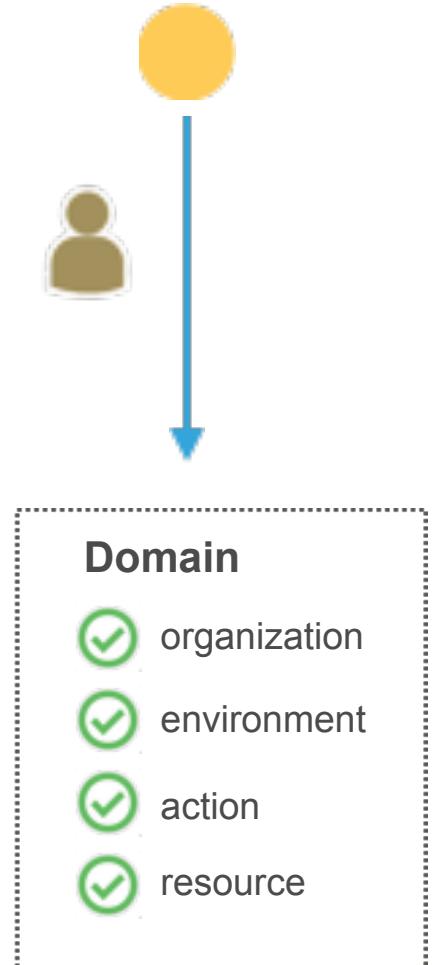
- in a given **domain** (Infra or Shared - service or tool)



Why Responsibility Management – Scenario 4

Service Y needs to check responsibility privileges for a user/subject:

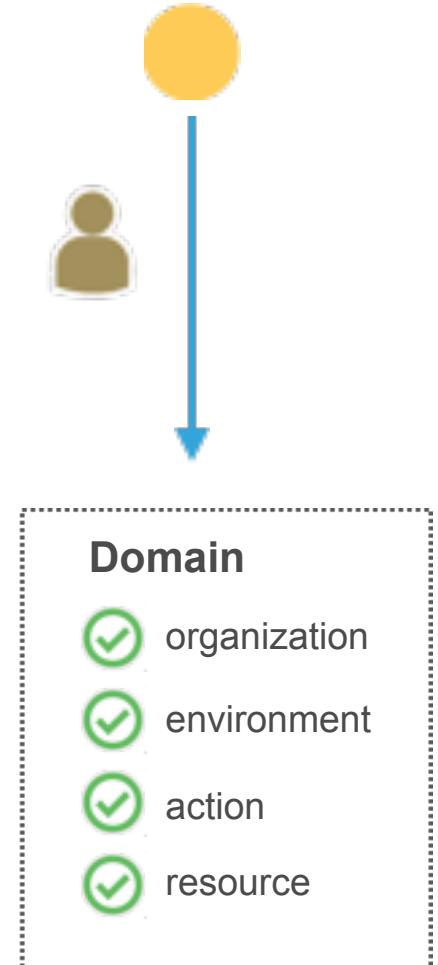
- in a given **domain** (Infra or Shared - service or tool)
- for a given **cost code identifier** or org. business unit (\$)



Why Responsibility Management – Scenario 4

Service Y needs to check responsibility privileges for a user/subject:

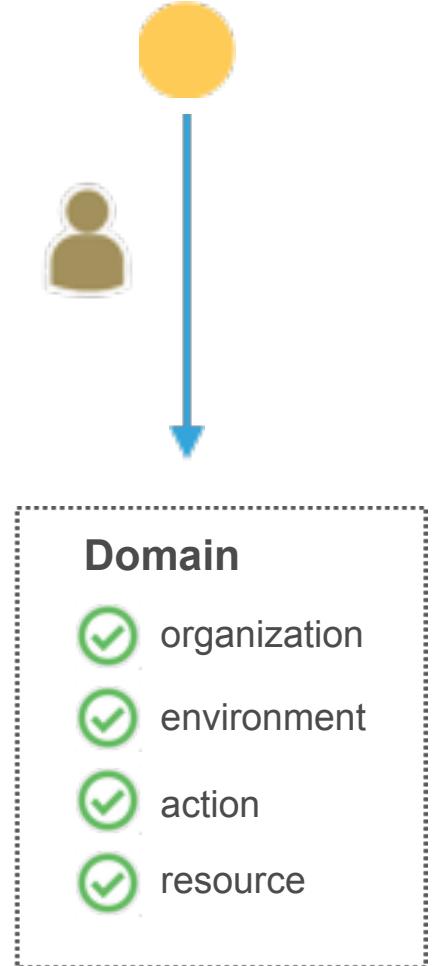
- in a given **domain** (Infra or Shared - service or tool)
- for a given **cost code identifier** or org. business unit (\$)
- for a given **environment** (e.g. 'PROD', 'QA', 'DEV' ...)



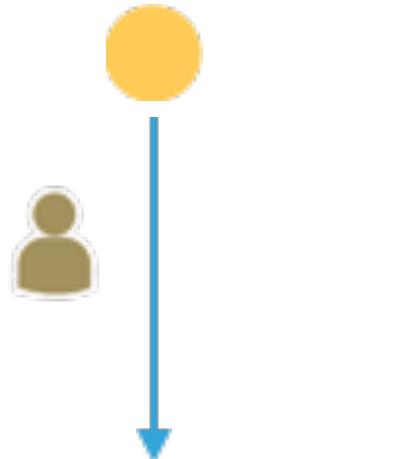
Why Responsibility Management – Scenario 4

Service Y needs to check responsibility privileges for a user/subject:

- in a given **domain** (Infra or Shared - service or tool)
- for a given **cost code identifier** or org. business unit (\$)
- for a given **environment** (e.g. 'PROD', 'QA', 'DEV' ...)
- for a given **action** (e.g. EDIT, DELETE, CREATE ...)



Why Responsibility Management – Scenario 4



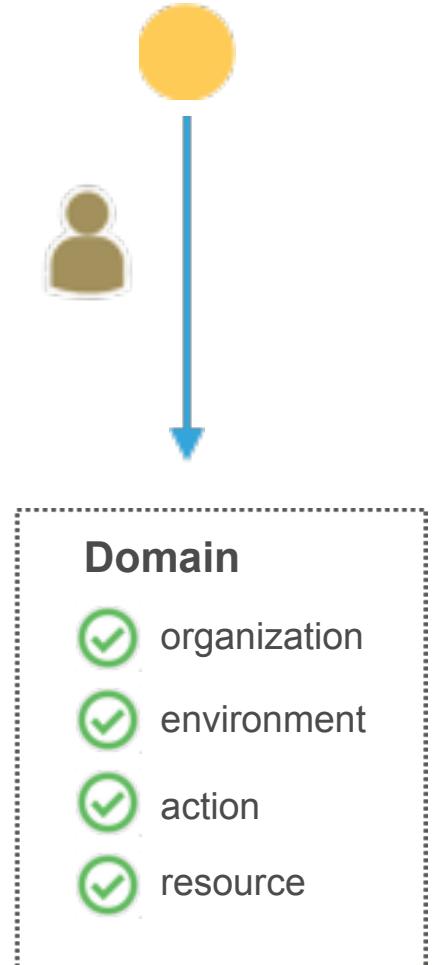
Service Y needs to check responsibility privileges for a user/subject:

- in a given **domain** (Infra or Shared - service or tool)
- for a given **cost code identifier** or org. business unit (\$)
- for a given **environment** (e.g. 'PROD', 'QA', 'DEV' ...)
- for a given **action** (e.g. EDIT, DELETE, CREATE ...)
- for a given **resource** (e.g. org.databases.prod.instance1.schema1)

Why Responsibility Management – Scenario 4 - Continued

Questions

- What if each request is for different sets of values for the given domain?
- What if other services have similar functional constraints with different values?
- Who manages **Role-Responsibility** per domain and **User-Role Mappings**?
- Who manages Mover/Leaver/Joiner logic ?





Responsibility Management System

Common Solutions

Responsibility Management – Common Solutions – For Data

DATA - Reference or Static Data for evaluation

- **LDAP/Active directory** queried by the application/service via **direct connections**.
- **User approver/manager** is queried via **proprietary corporate directory services**.
- **Role-Responsibility mappings** usually stored in **local persistence, proprietary systems**.
- **User-Role mappings** usually stored in any of: **local persistence, proprietary systems**.

Responsibility Management – Common Solutions – For Functions

LOGIC - Calculations / Functions (Policies)

- Complex functions/calculations were **coded into the application/service**.
- Newer applications/services may separate such as an **independent [micro]service**.
- Some applications/services utilize **embedded rule engines** such as Drools.
- Some applications/services utilize **proprietary entitlement systems** for evaluations.



Responsibility Management System

A solution to manage dynamic privileges and entitlements



Responsibility Management System

A Right Solution

Responsibility Management System (RMS) – The Right Solution

A Responsibility Management System that:

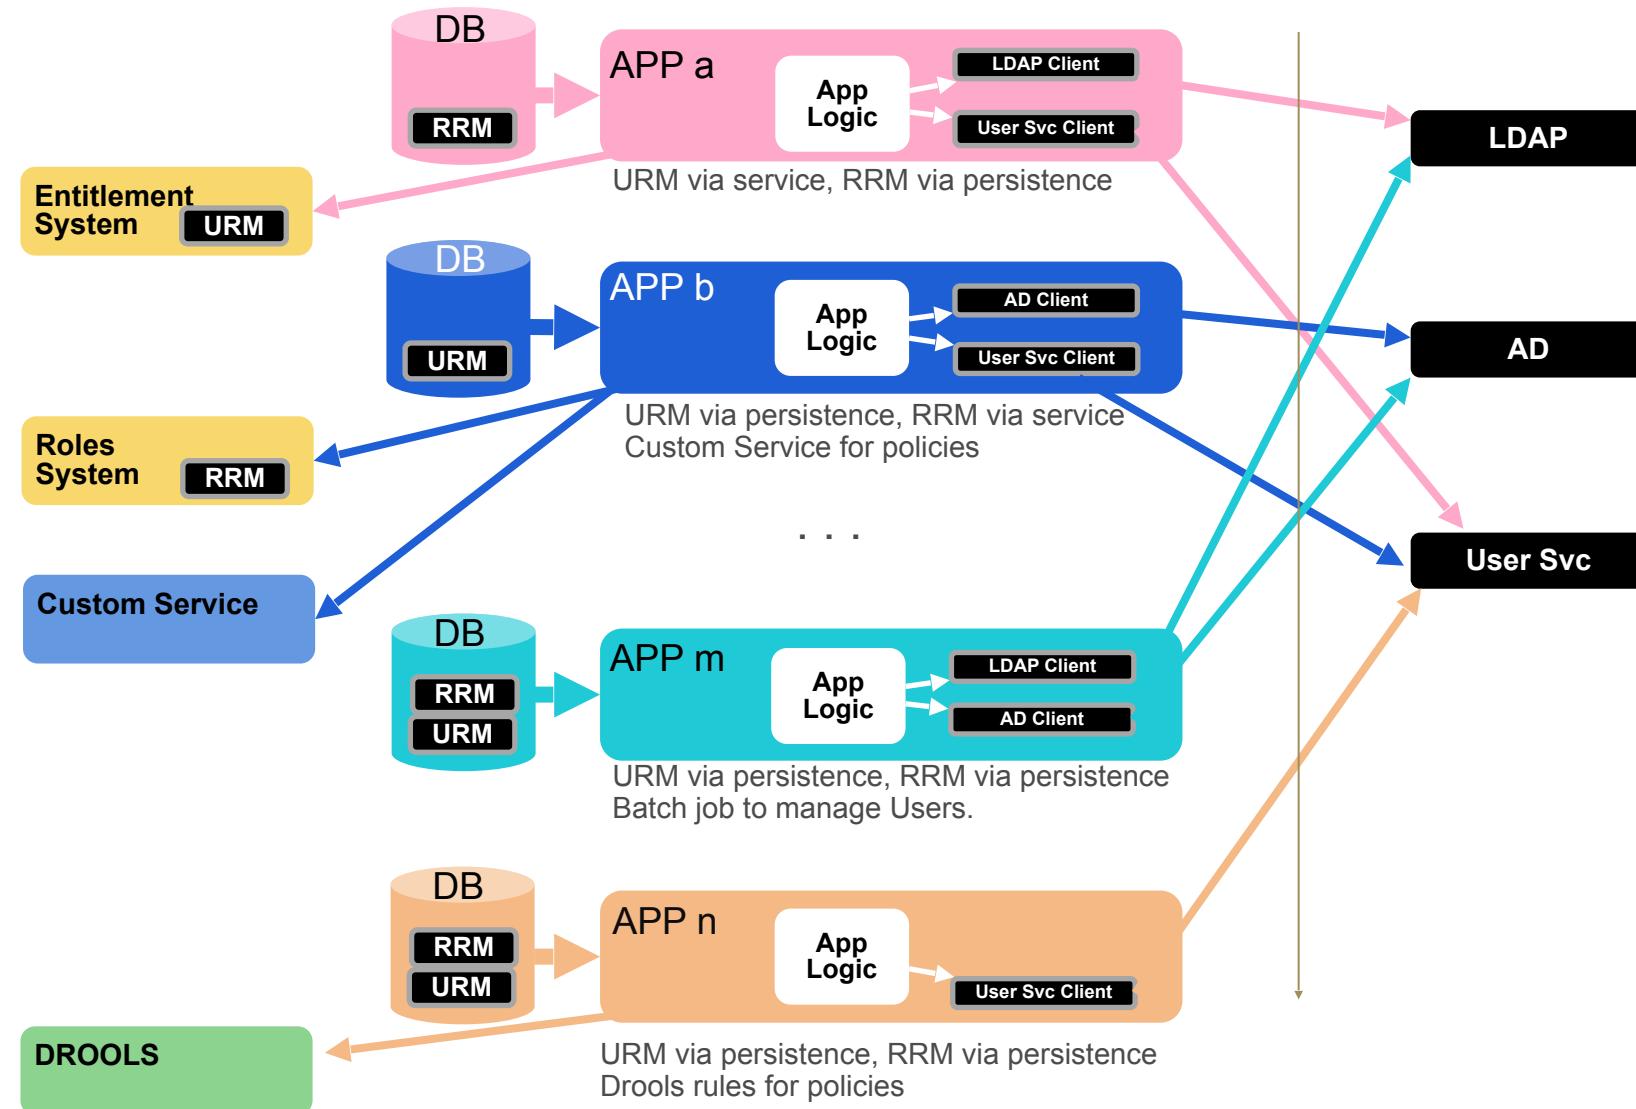
- federates the calls to LDAP, Active Directory, and other services as **integrated services**
- provides appropriate **mapping of roles and responsibilities**, per domain
- provides for **user to role mapping, per organization per domain**
- provides **proper SDLC and audit mechanism** for policies per domain, to author and deploy
- • •

Responsibility Management System (RMS) – The Right Solution - Continued

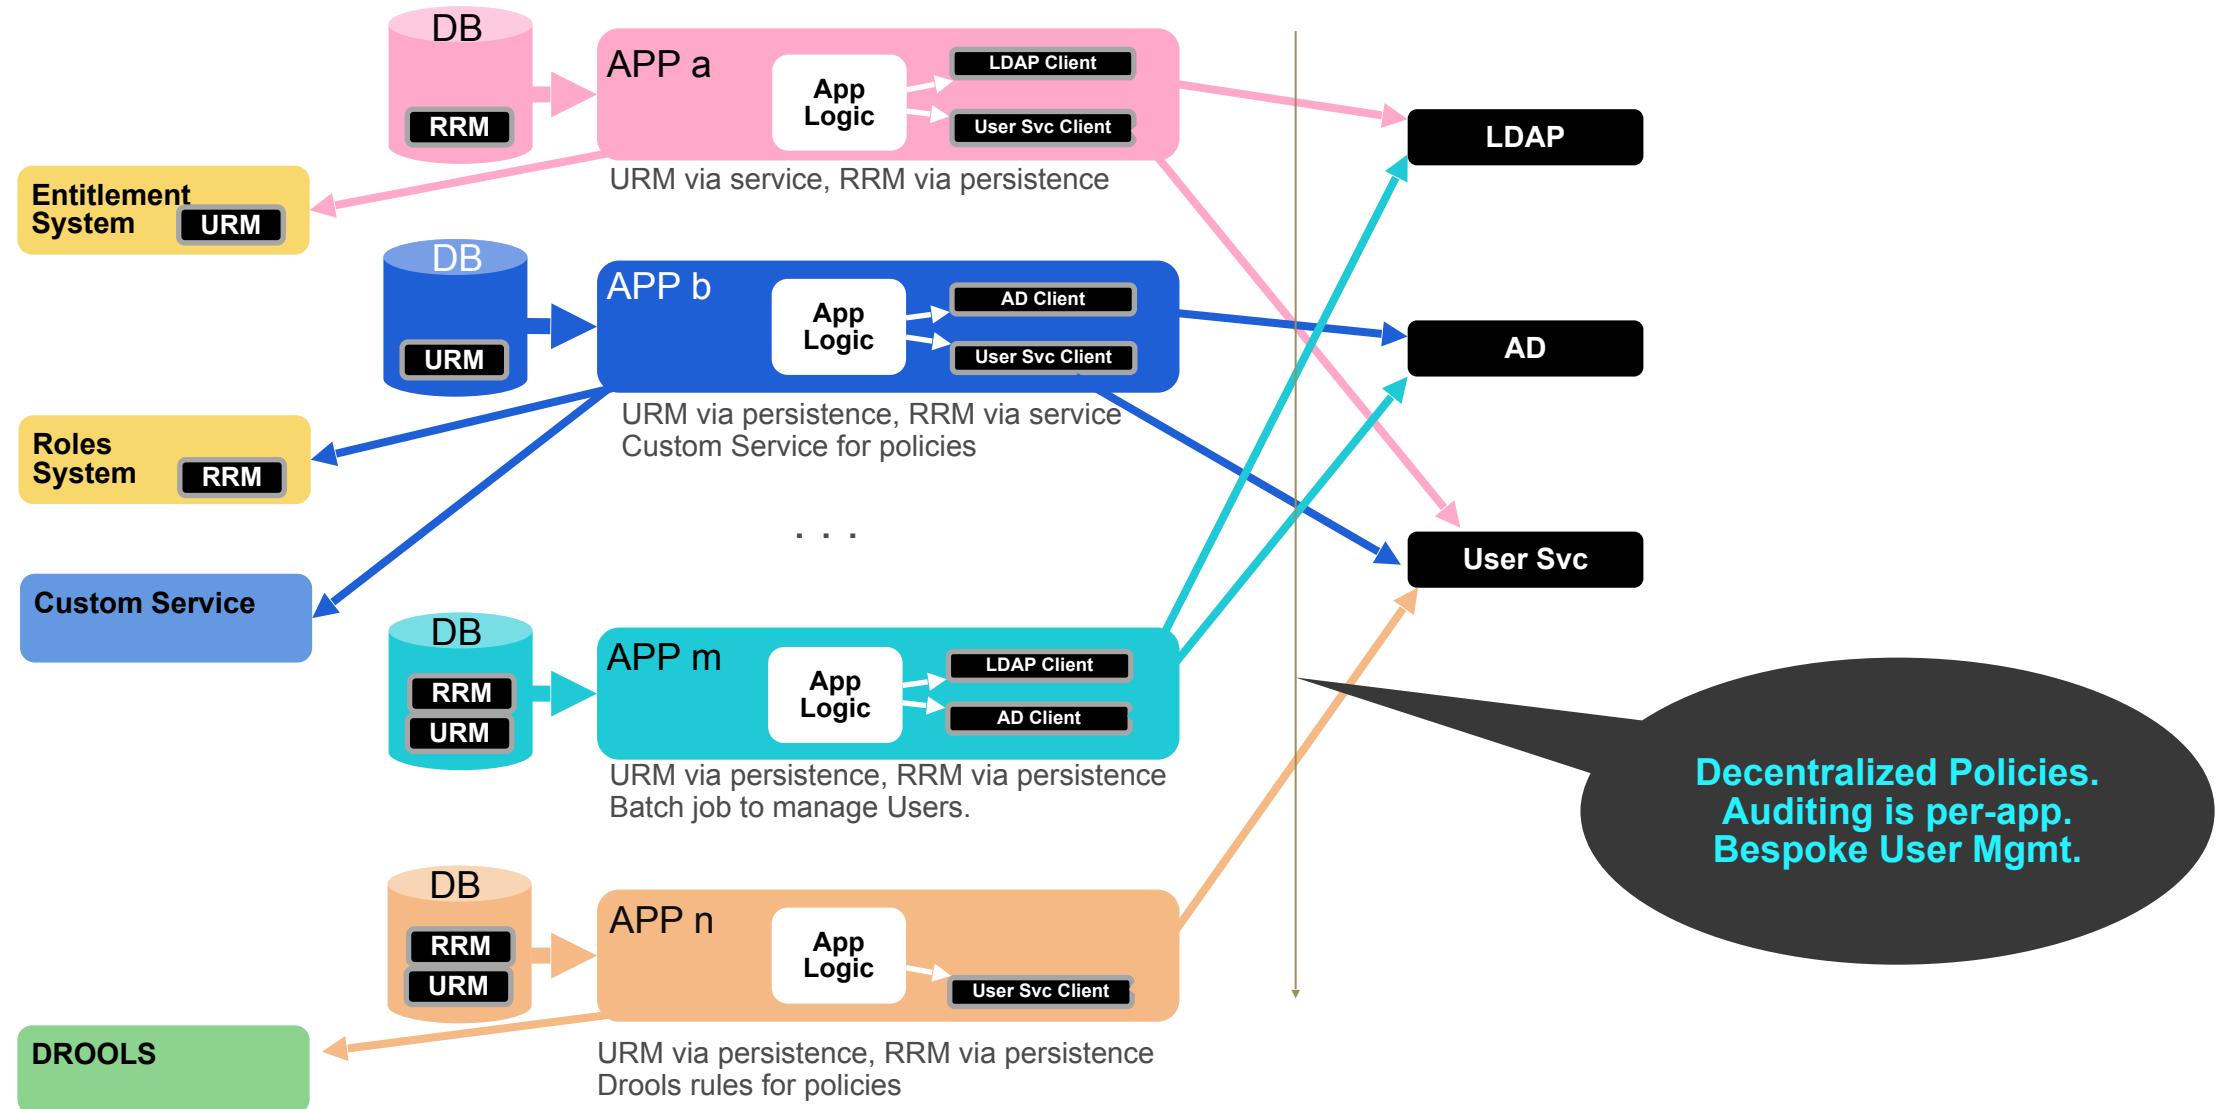
A Responsibility Management System that:

- provides a **policy engine** to evaluate complex calculations/functions using:
 - data provided as inputs by service-consumer
 - data queried from integrated services
 - policies provided by the domains
- caters to **applying a mover/leaver/joiner logic** to all controlled datasets
- provides **horizontal scaling** and thus, **high availability** for varying request volumes

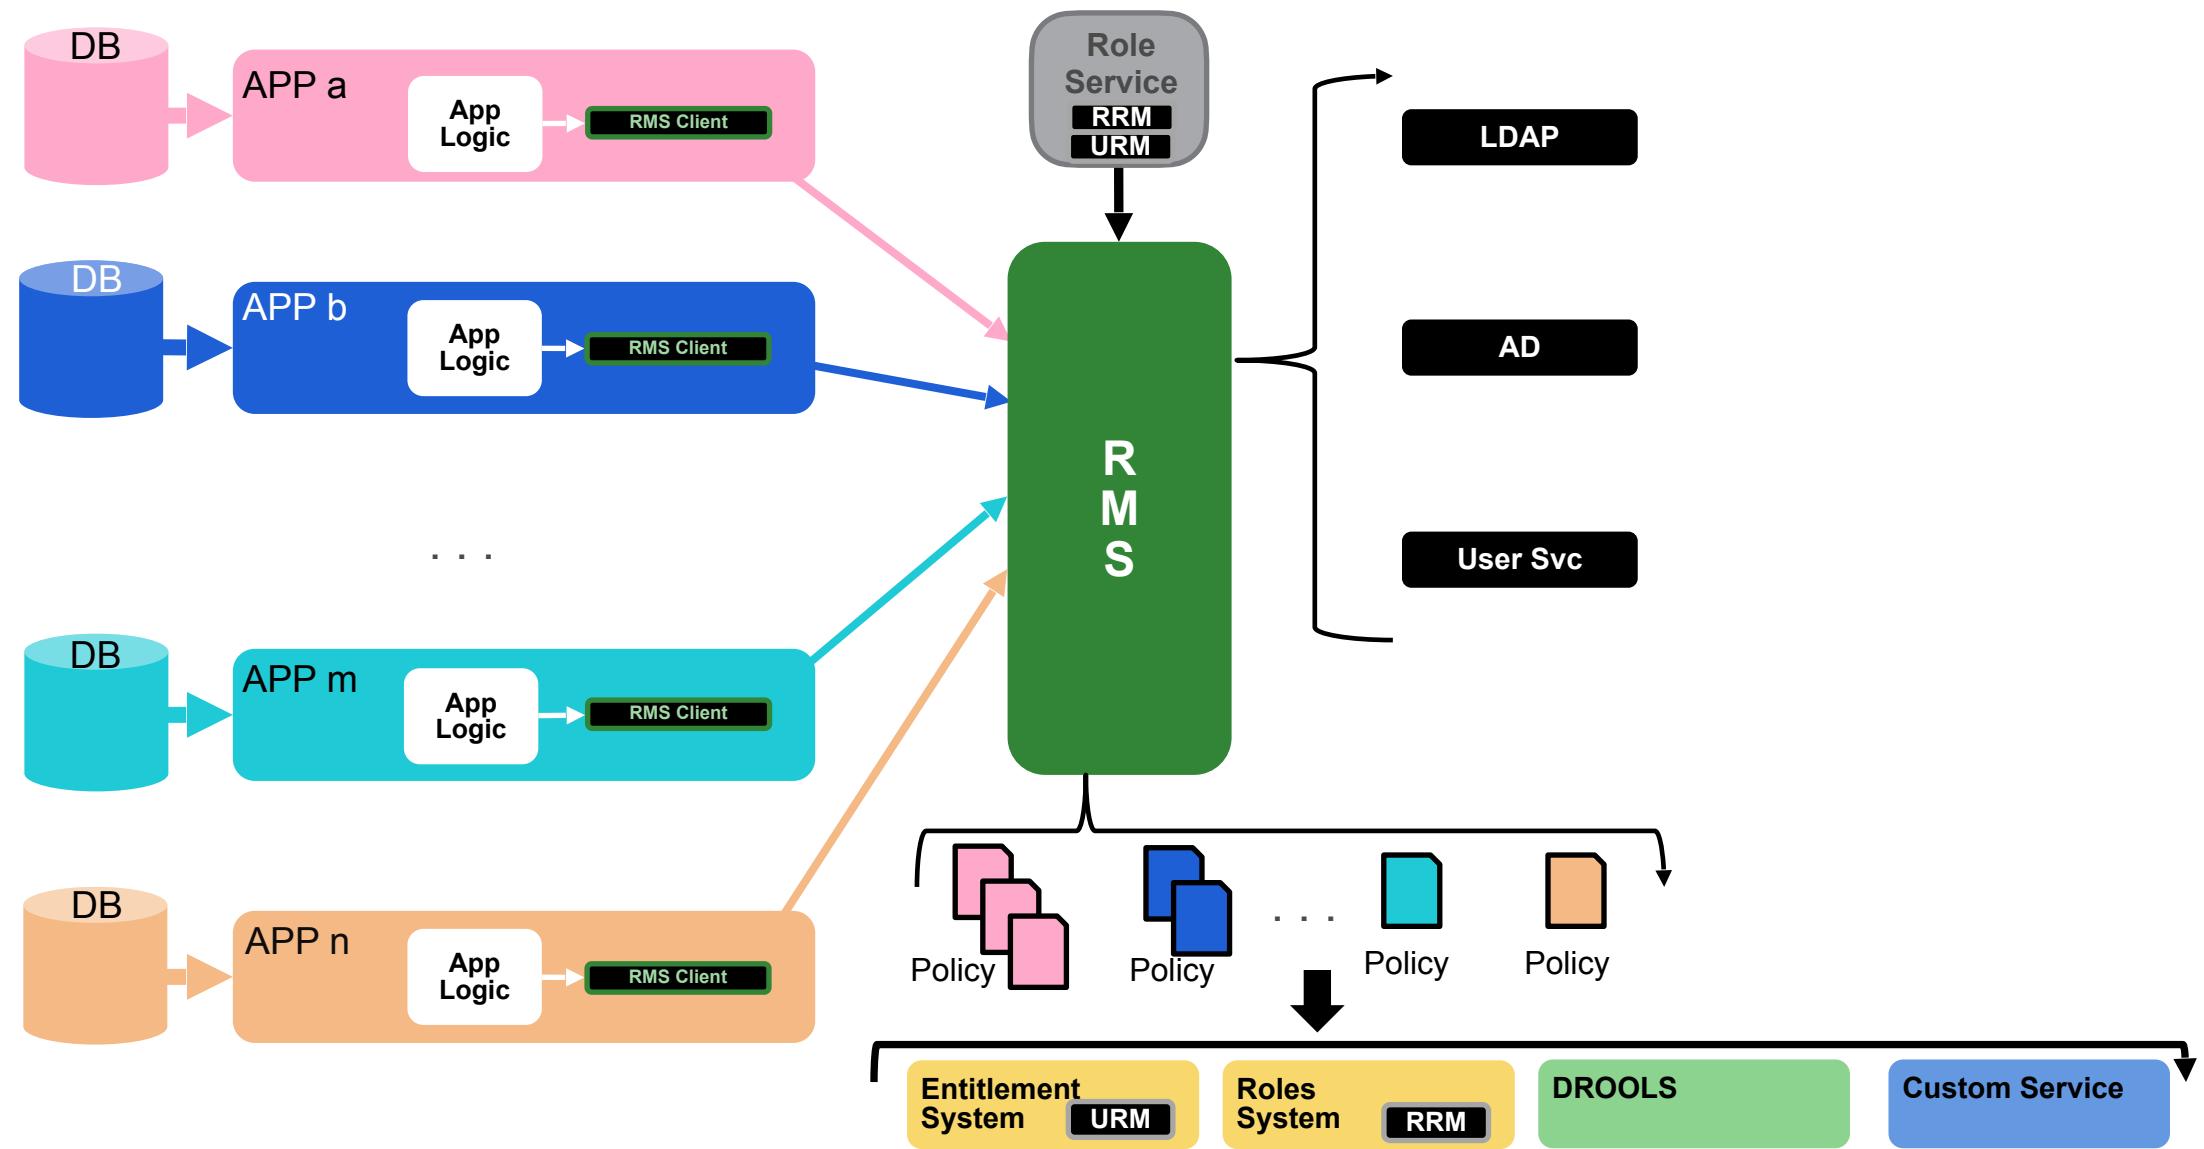
Before RMS



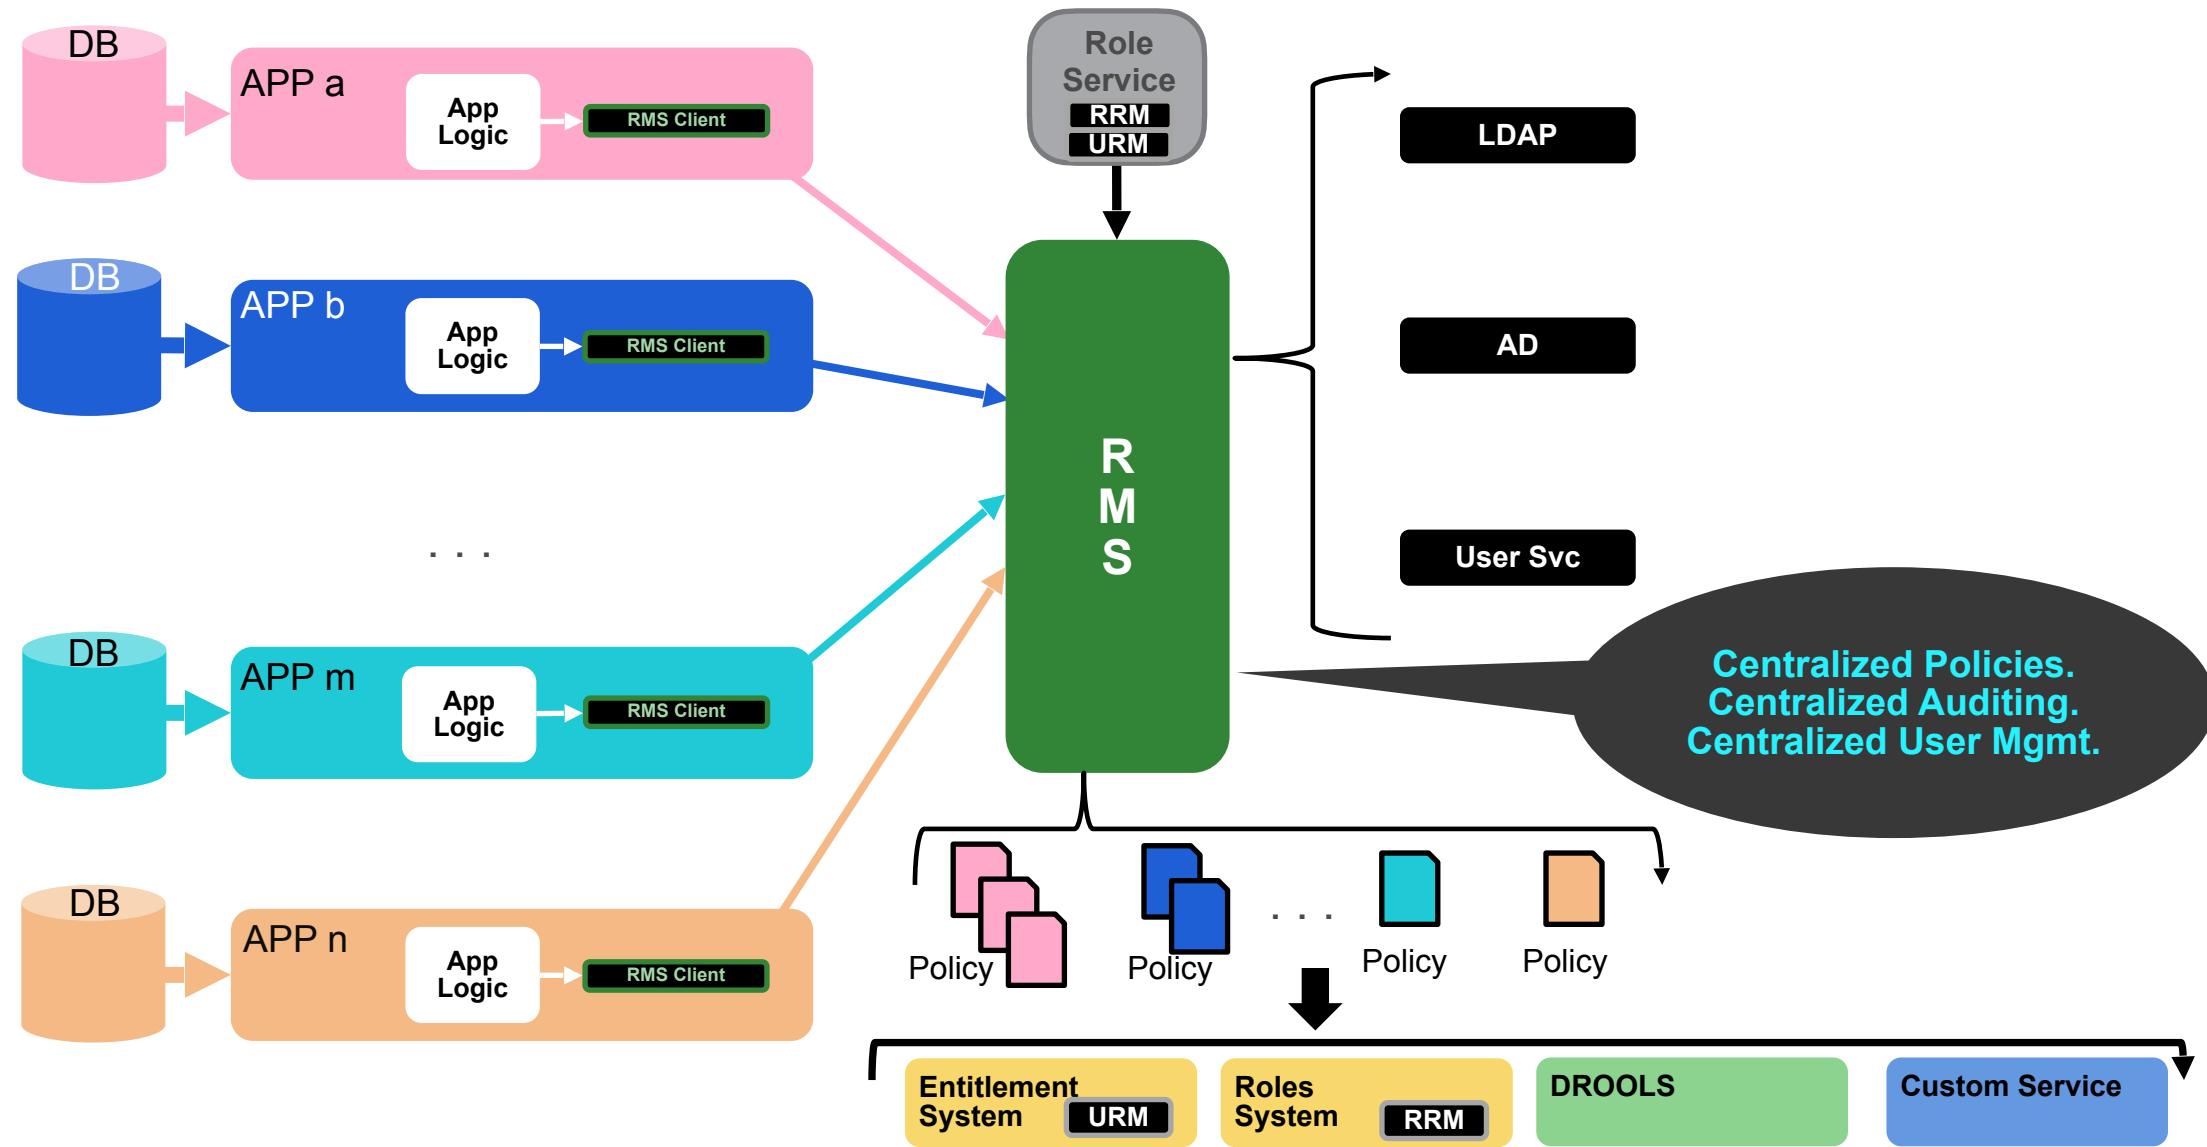
Before RMS



Post-RMS visualized



Post-RMS visualized





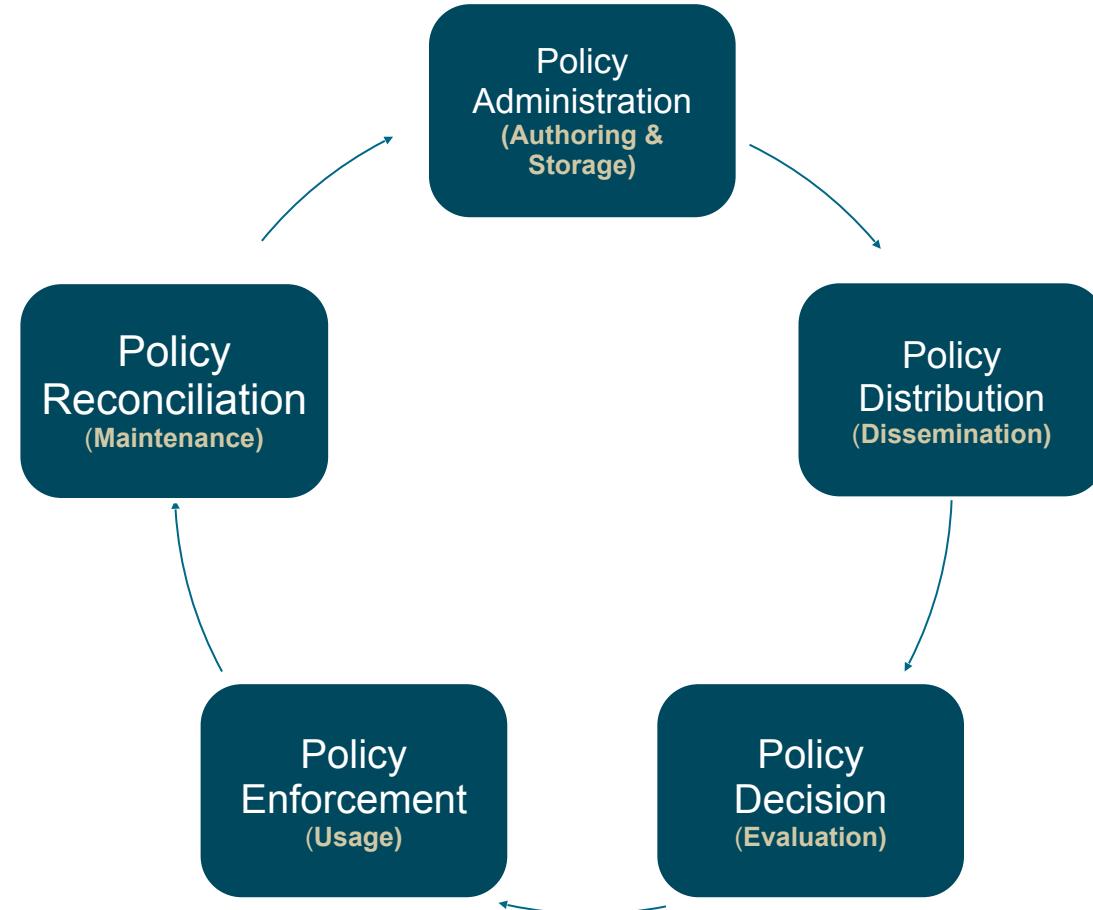
Architecting the Responsibility Management System

A platform solution for Responsibility Management

Responsibility Management - Policy Lifecycle

Responsibility Management is performed via policies

Policies have a lifecycle



More detailed flow: Appendix

Technology Choices

Technology choices for building the Responsibility Management System

Payload format: HOCON: Use case

A Case for using Human-Optimized Configuration Object Notation

- Intent is to expose query-only operations.
- Policy decisions should be queried (non-mutating), thus logically GET operations.
 - GET operations do not support a request body.
 - GET operations may be exposed to character limits, large content not possible.
 - ☒ POST operations allow for a request body but do not support meaningful caching.
 - JSON and individual query parameters are quite verbose.
 - ☒ HOCON * trims the parameter verbosity by a significant amount.

<https://github.com/lightbend/config/blob/master/HOCON.md>

Payload format: HOCON: Benefits

Benefits of using Human-Optimized Configuration Object Notation

HOCON (see link below)

- syntax is quite simple and has low ambiguity.
- is a superset of JSON. JSON is parsed properly by HOCON parsers.
- allows the use of comments.
- allows multi-line strings.
- allows for includes and substitutions (next slide has more on this).
- has built-in durations (5d or 100ms)

<https://github.com/lightbend/config/blob/master/HOCON.md>

Payload format: HOCON: Features

Human-Optimized Configuration Object Notation - using *includes* and *substitutions*

generic.conf

```
{x: 10, y: ${x}, z: 5s}
```

my.conf

```
{a : { include "generic.conf" } }
```

Substitution

Inclusion

```
a.x = 10  
a.y = 10  
a.z = 5s
```

<https://github.com/lightbend/config/blob/master/HOCON.md>

Payload format: HOCON: Comparing to JSON

Human-Optimized Configuration Object Notation - sample comparisons to JSON

Sample JSON

```
foo : {  
    bar : {  
        baz: myvalue  
    }  
}  
  
employee: {  
    firstname: "Jane"  
    lastname: "Doe"  
    nested: {  
        loginTimeoutInMilliseconds: 5000  
    }  
    fullname: "Jane Doe"  
}  
  
standard-policy: {  
    developer: "yes"  
    operator: false  
}
```

Sample HOCON

```
foo.bar.baz: myvalue  
--- Or ---  
foo { bar { baz: myvalue}}  
  
employee {  
    firstname: "Jane"  
    lastname: "Doe"  
    nested {  
        loginTimeout: 5s  
    }  
    fullname: ${employee.firstname} ${employee.lastname}  
}  
standard-policy {  
    developer: "yes"  
    operator: false  
}
```

Java Collections Library: Eclipse Collections

Key highlights

Eclipse Collections (see link below)

- Rich, concise and readable APIs.
- Clear mutable and immutable hierarchies for collection types.
- Memory efficient containers.
- Optimized eager APIs instead of Java Collection Framework's lazy APIs.
- Improved code readability.
- Ease of learning thanks to several Code Katas.

<https://www.eclipse.org/collections/>

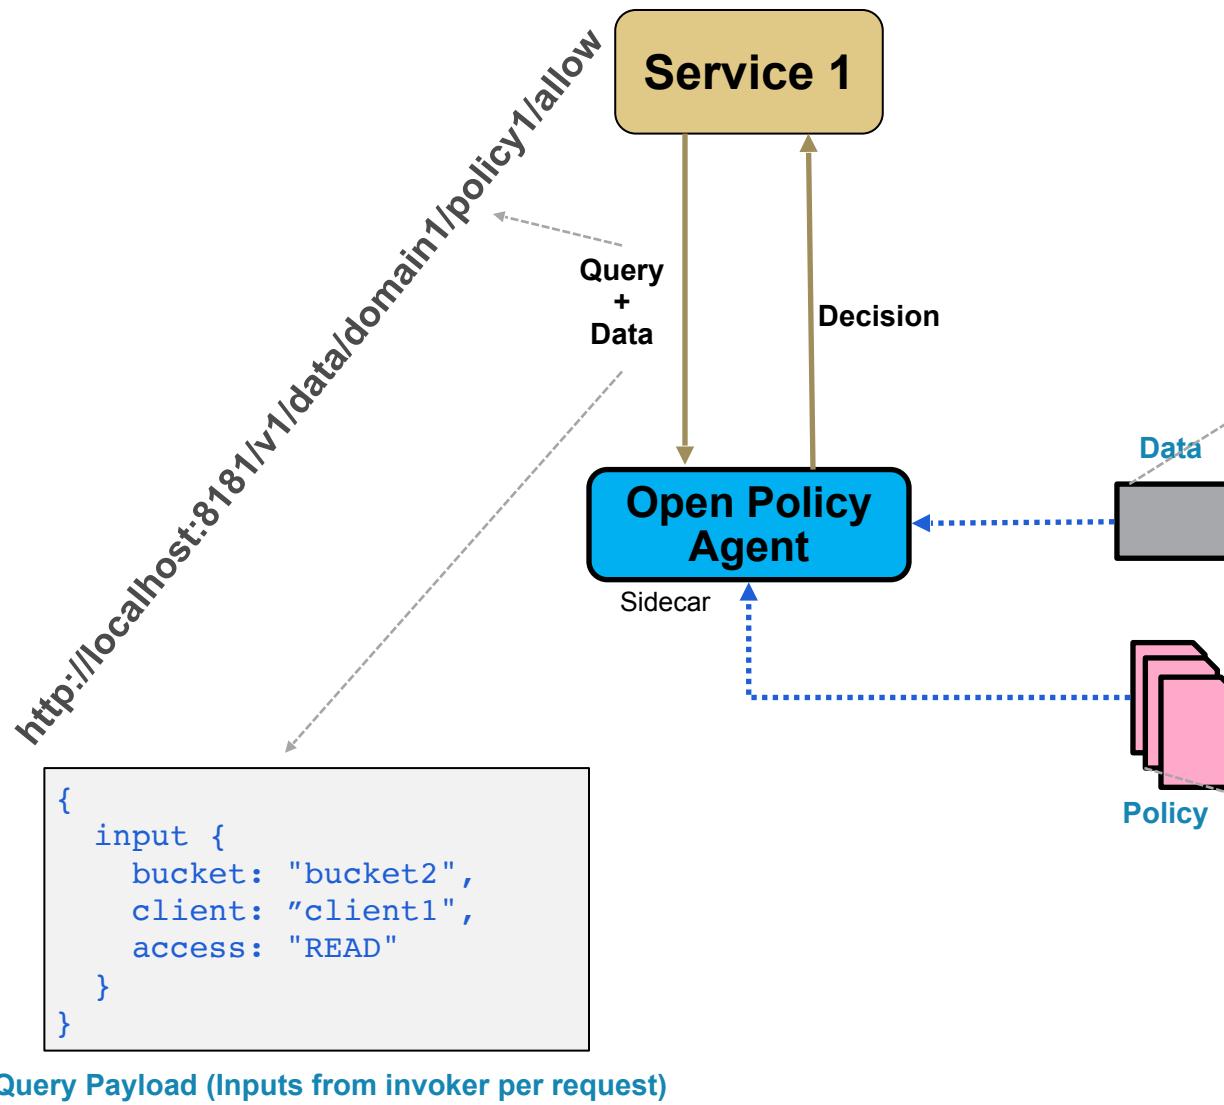
Policy Engine: Open Policy Agent

Key highlights

- Open Policy Agent (OPA) (see link below) is an open source general purpose policy engine.
- Uses “rego” (inspired by Datalog) as a declarative native query language.
- Policies are written as rulesets (similar to functions).
- Policies can be queried as RESTful POST operations.
- Data and policy publishing is via RESTful PUT operations.
- Can be launched as a **library** for a **service**, an **independent daemon** or **sidecar**.
- Decision in RMS was to use OPA as a **sidecar**.

<https://www.openpolicyagent.org/>

Open Policy Agent: Usage Pattern



```
[  
  {  
    "name": "bucket1",  
    "clients": [  
      {  
        "name": "client1",  
        "access": ["READ", "WRITE"]  
      },  
      {  
        "name": "client2",  
        "access": ["WRITE"]  
      }  
    ]  
  },  
  {  
    "name": "bucket2",  
    "clients": [  
      {  
        "name": "client1",  
        "access": ["READ"]  
      }  
    ]  
  }  
]
```

data.json

```
package domain1.policy1  
  
import data.domain1.policy1.buckets  
  
default allow = false  
  
allow {  
  buckets[i].name == input.bucket  
  buckets[i].clients[j].name == input.client  
  buckets[i].clients[j].access[k] == input.access  
}
```

policy.rego

Policy Logic (Available from Policy Distribution Point)

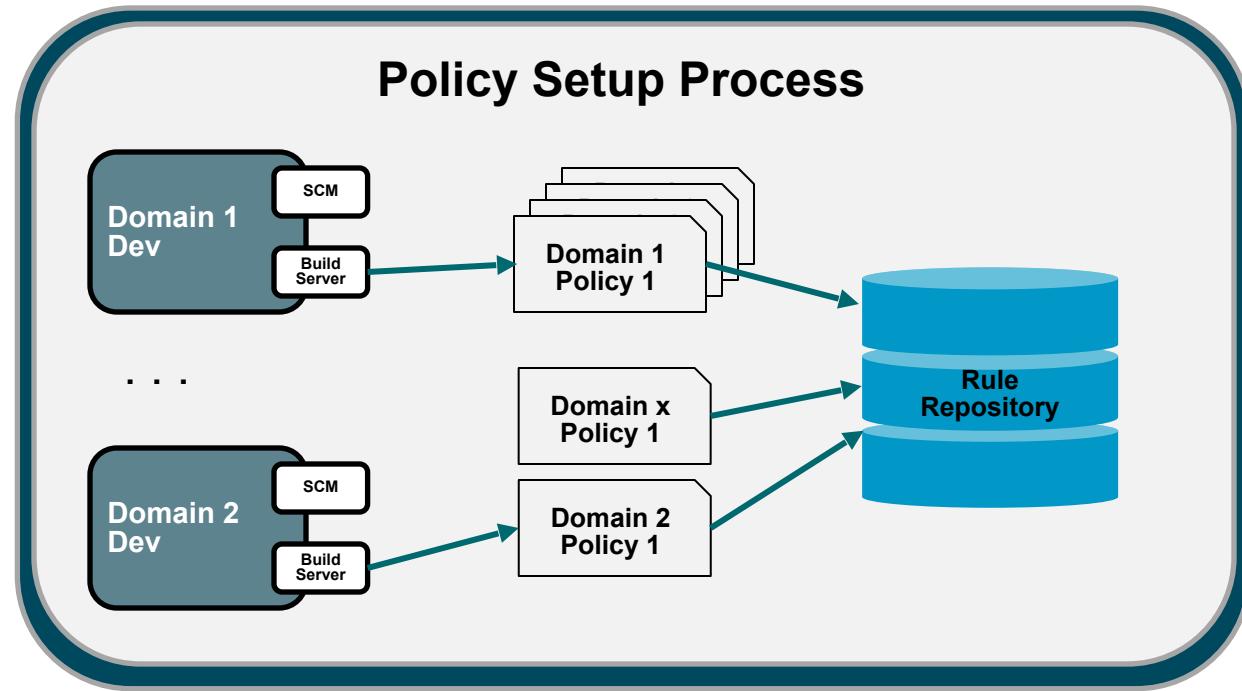


Responsibility Management System

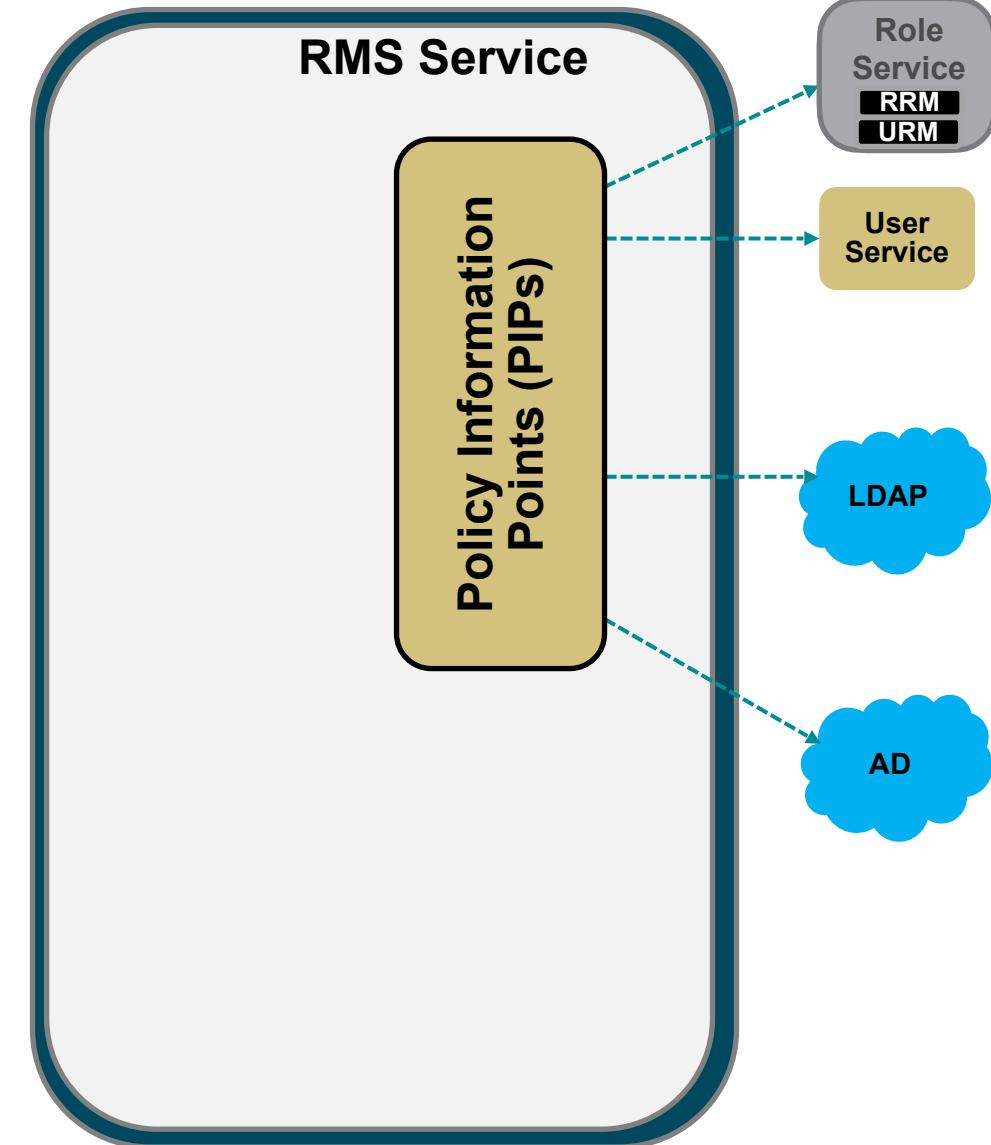
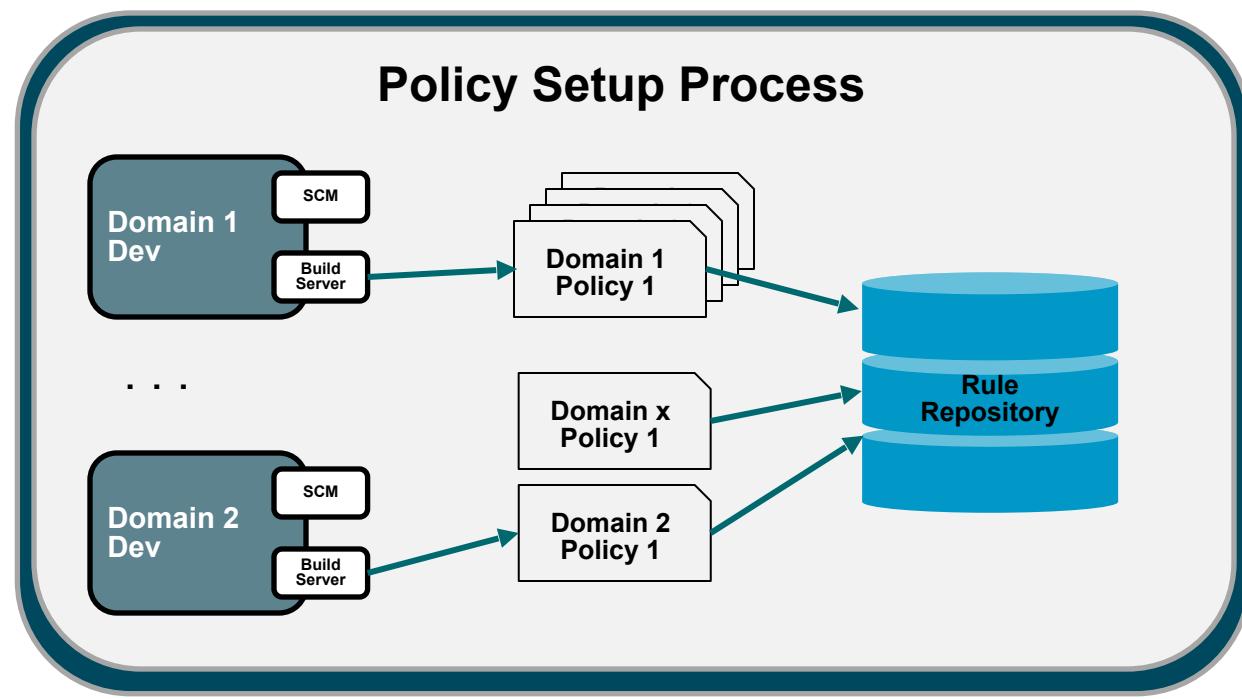
Architecture (Version 1)

A *Federated* Responsibility Management Service

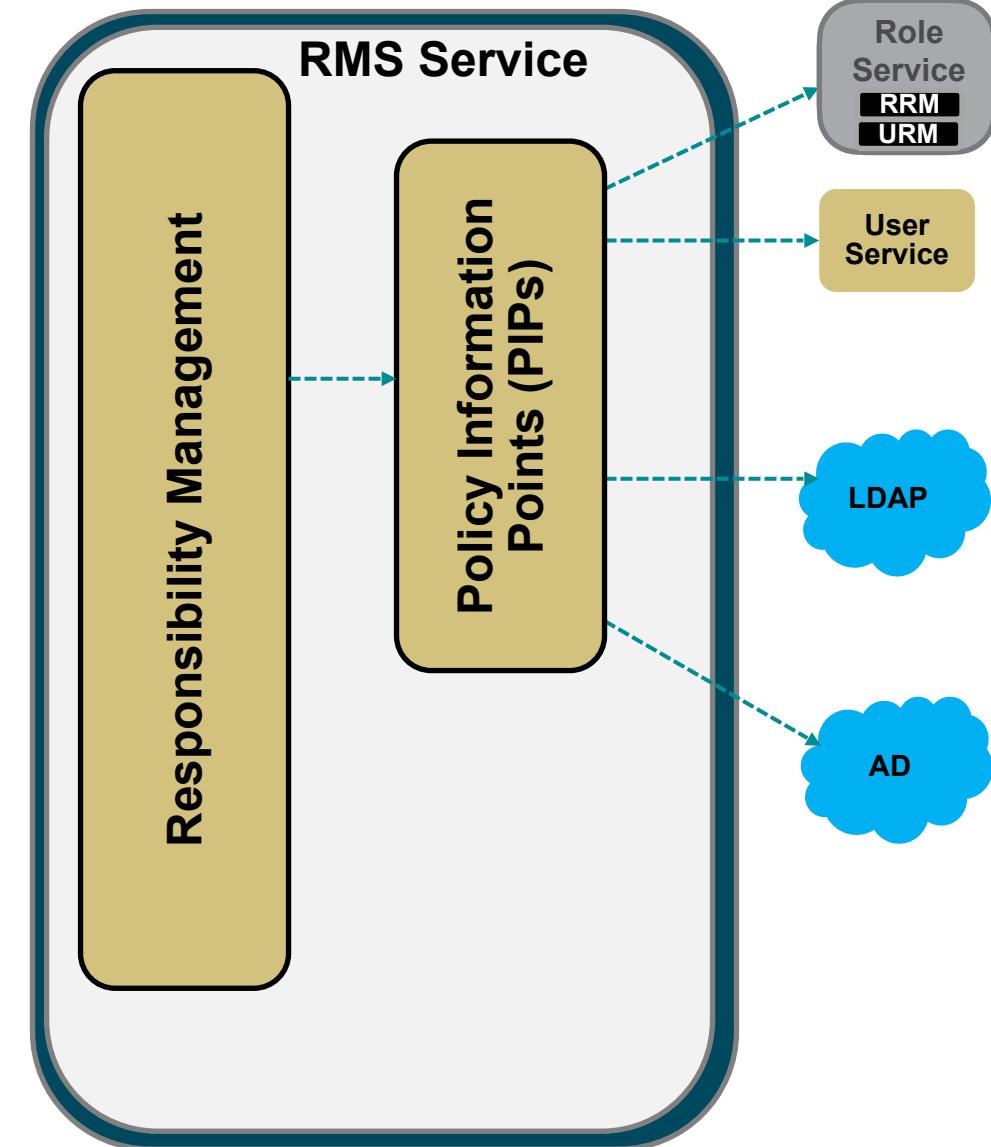
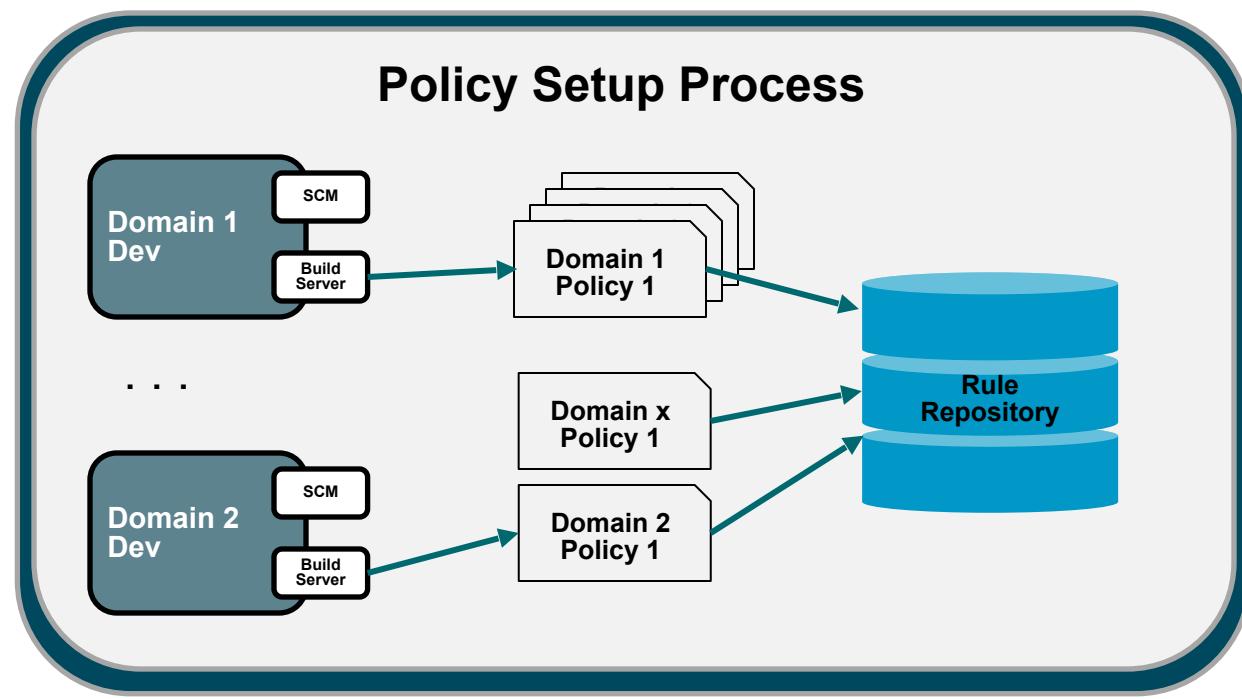
RMS Architecture - Version 1: Federated



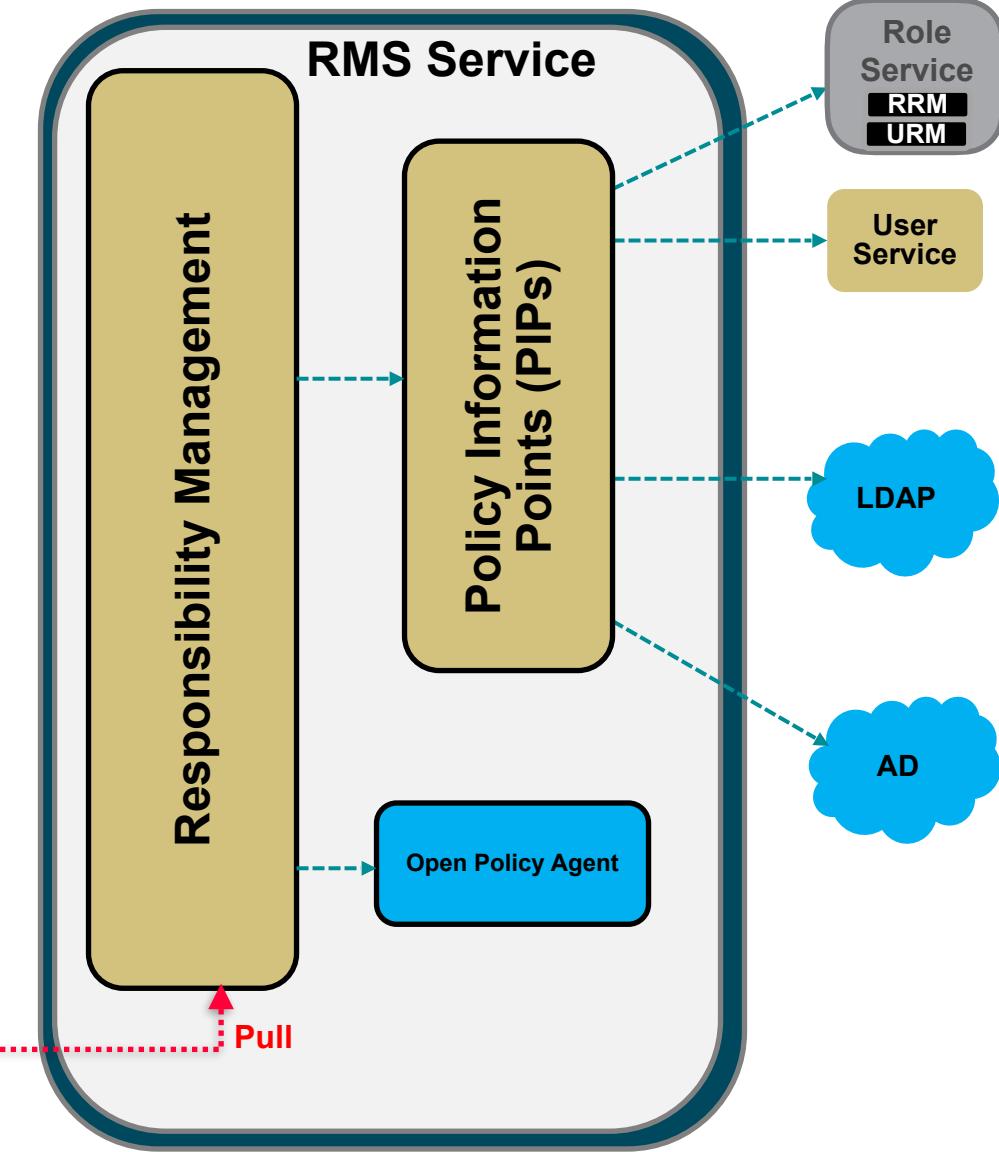
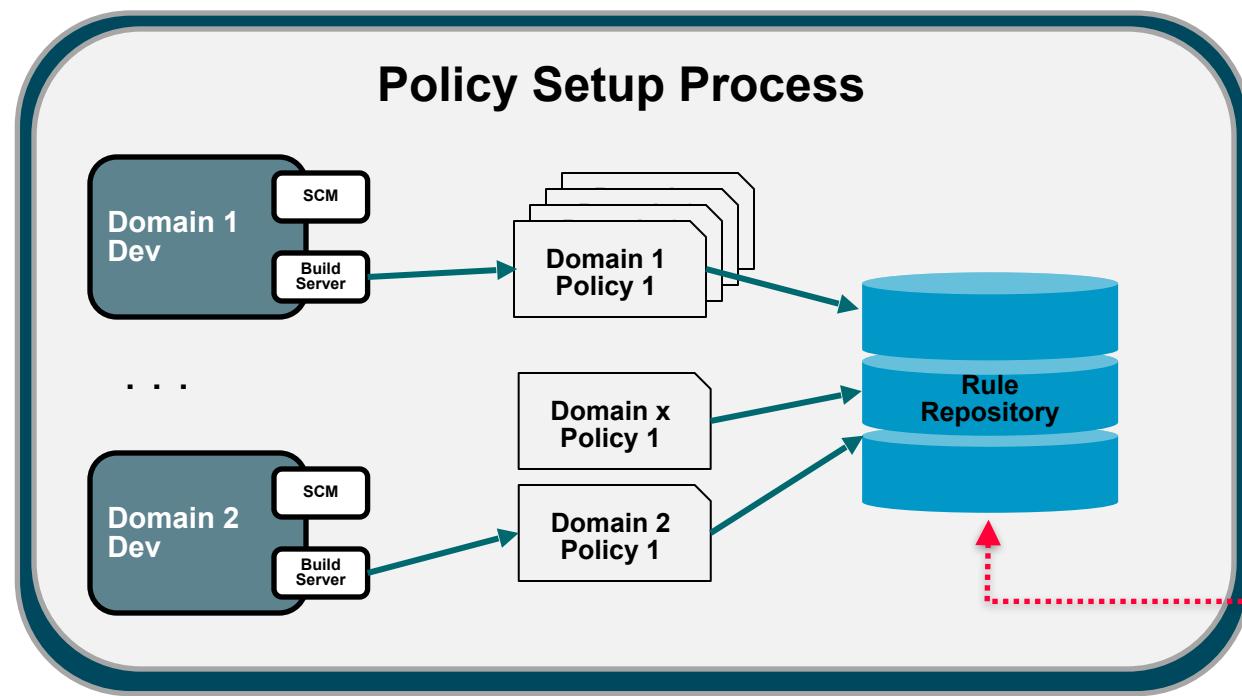
RMS Architecture - Version 1: Federated



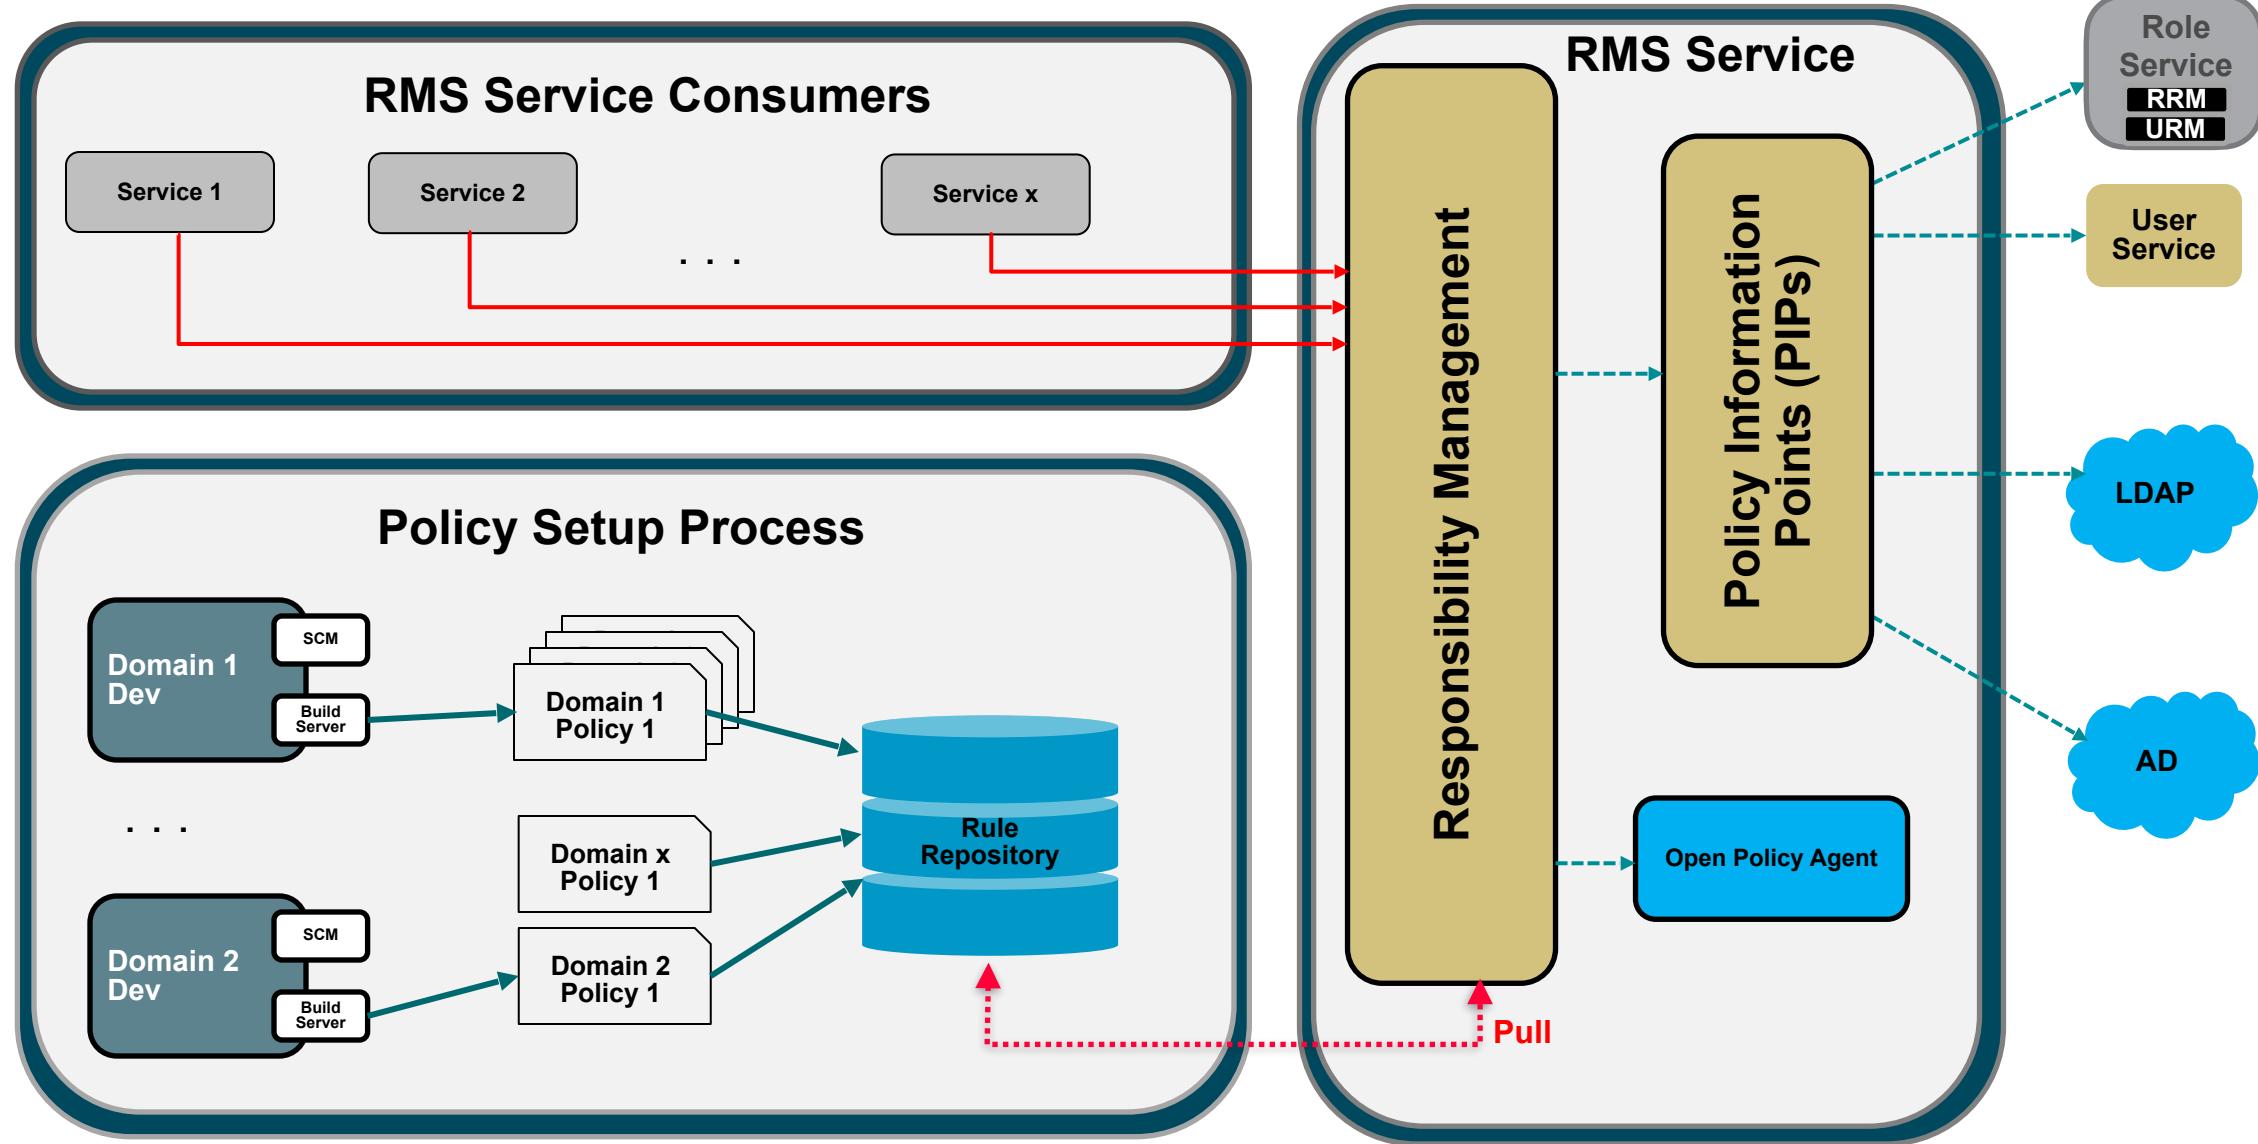
RMS Architecture - Version 1: Federated



RMS Architecture - Version 1: Federated



RMS Architecture - Version 1: Federated



Federated Architecture: Issues Faced

Key issues

- Segregation and information-barrier needs **implied more work**.
- A rogue policy script could lead to **loss of service for all domains**.
- RM Service **became the gatekeeper for testing and coverage**.
- RM Service **had to establish a release-train model** to pick up new policies.
- Out-of-band policy changes lead to **intermittent service-unavailability**.
- **Observation:** Policy changes were more frequent when a new domain onboards.

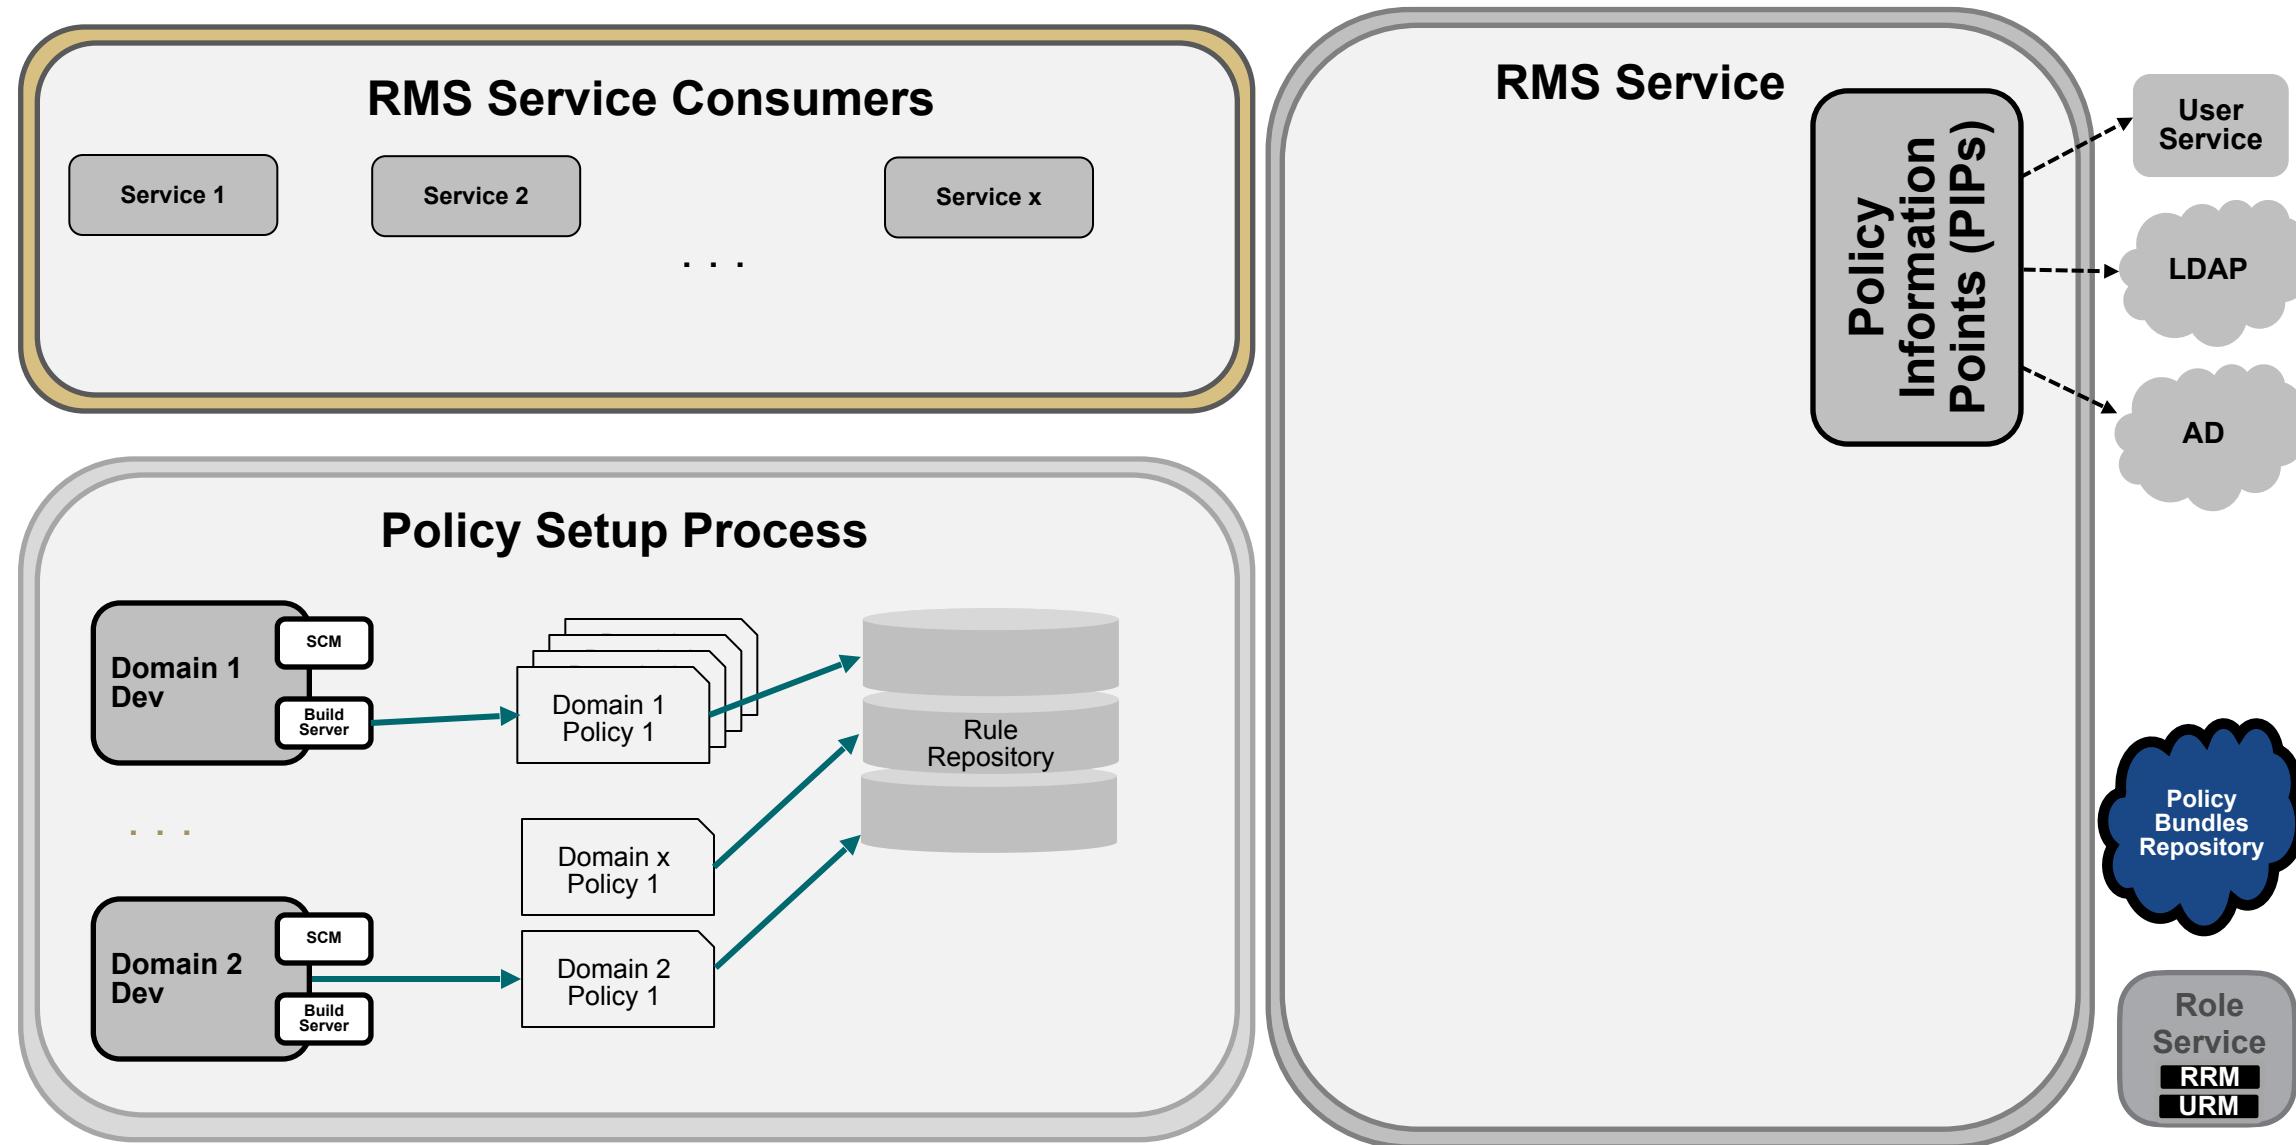


Responsibility Management System

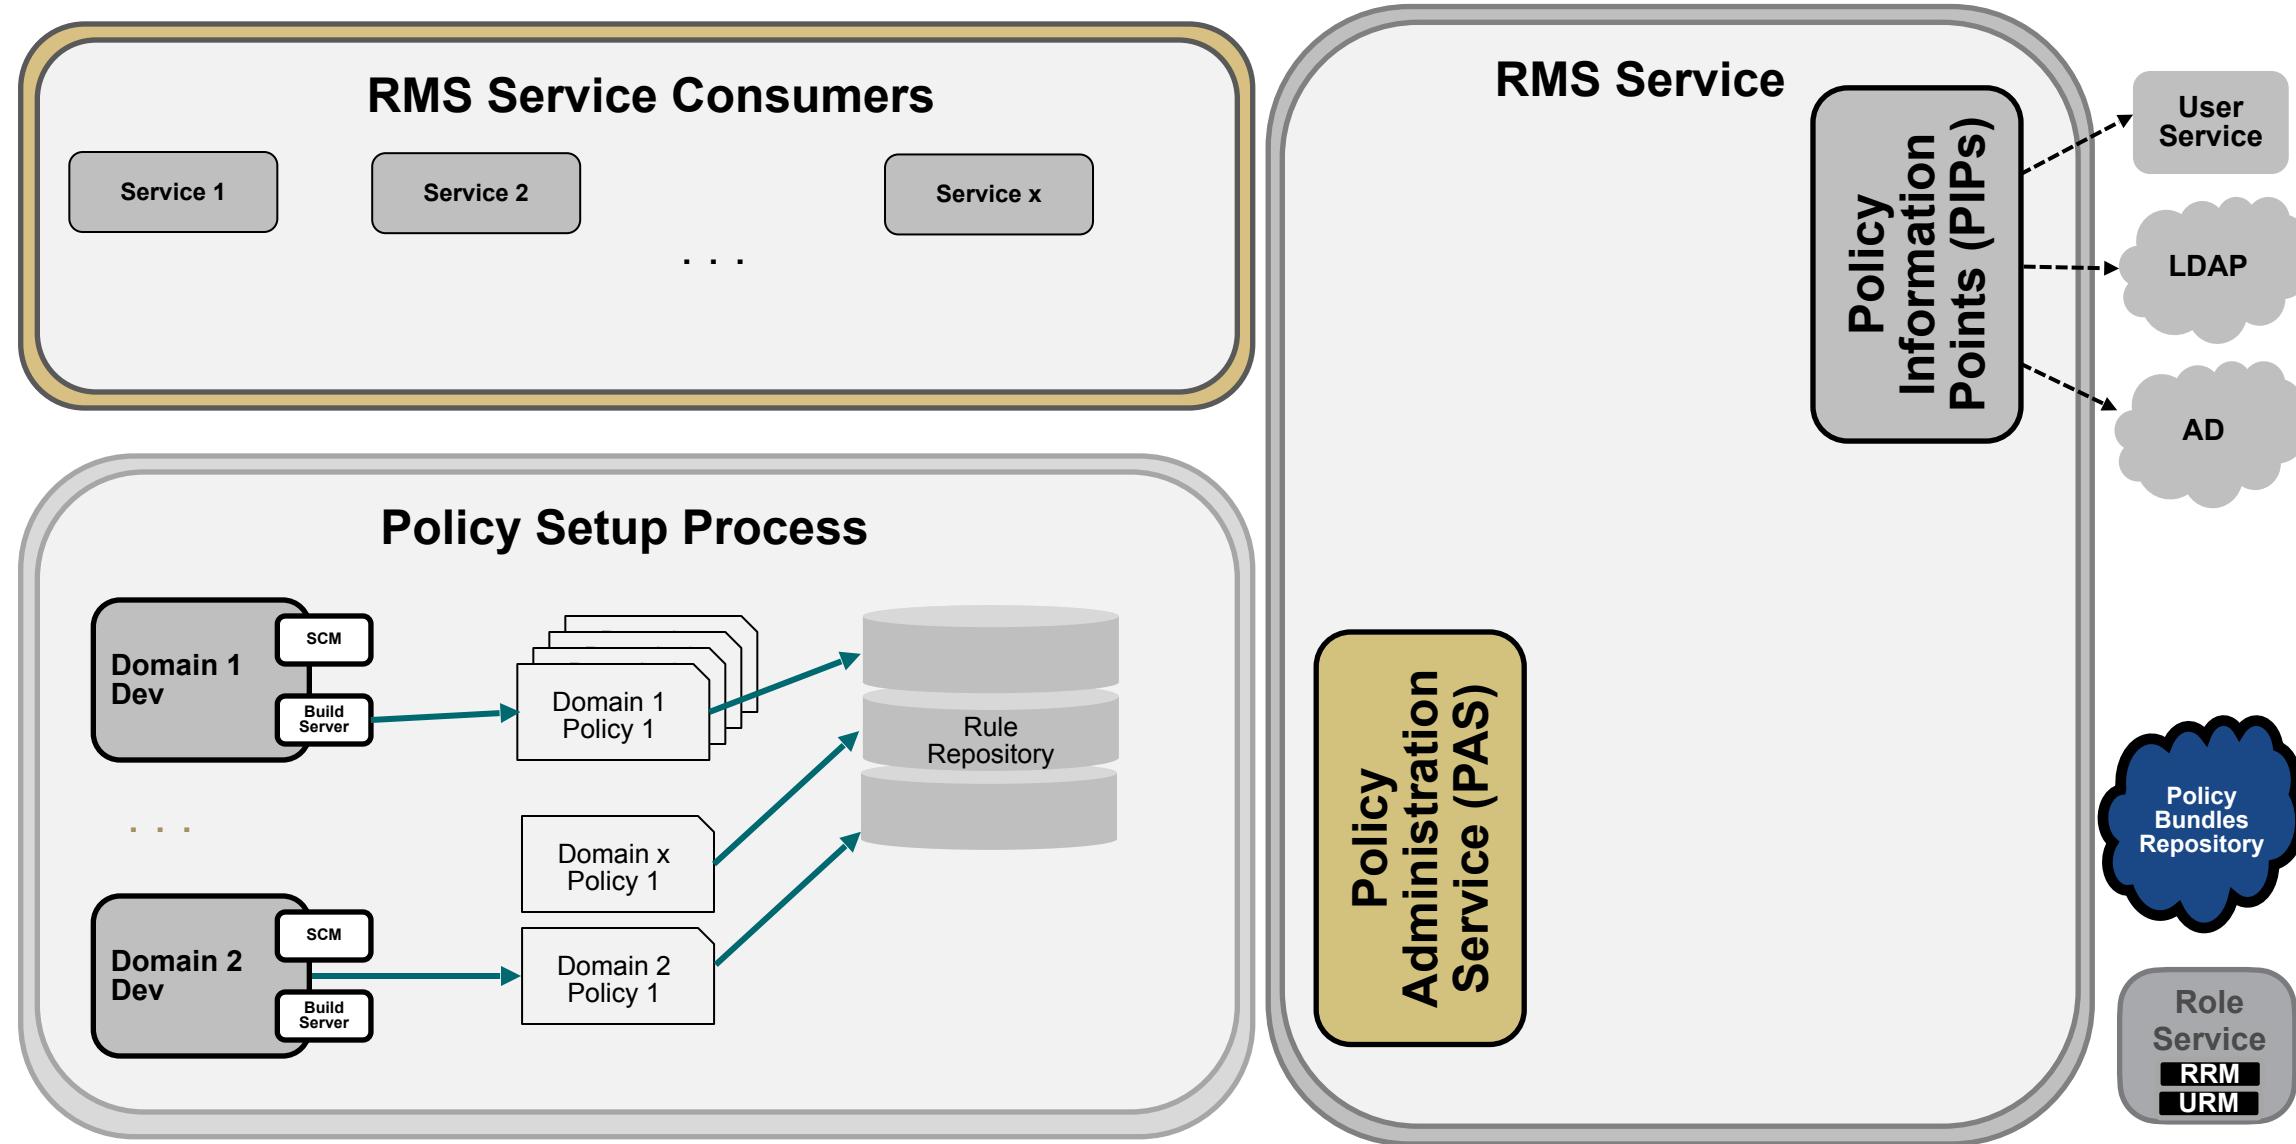
Architecture (Version 2)

A *Distributed* Responsibility Management Service

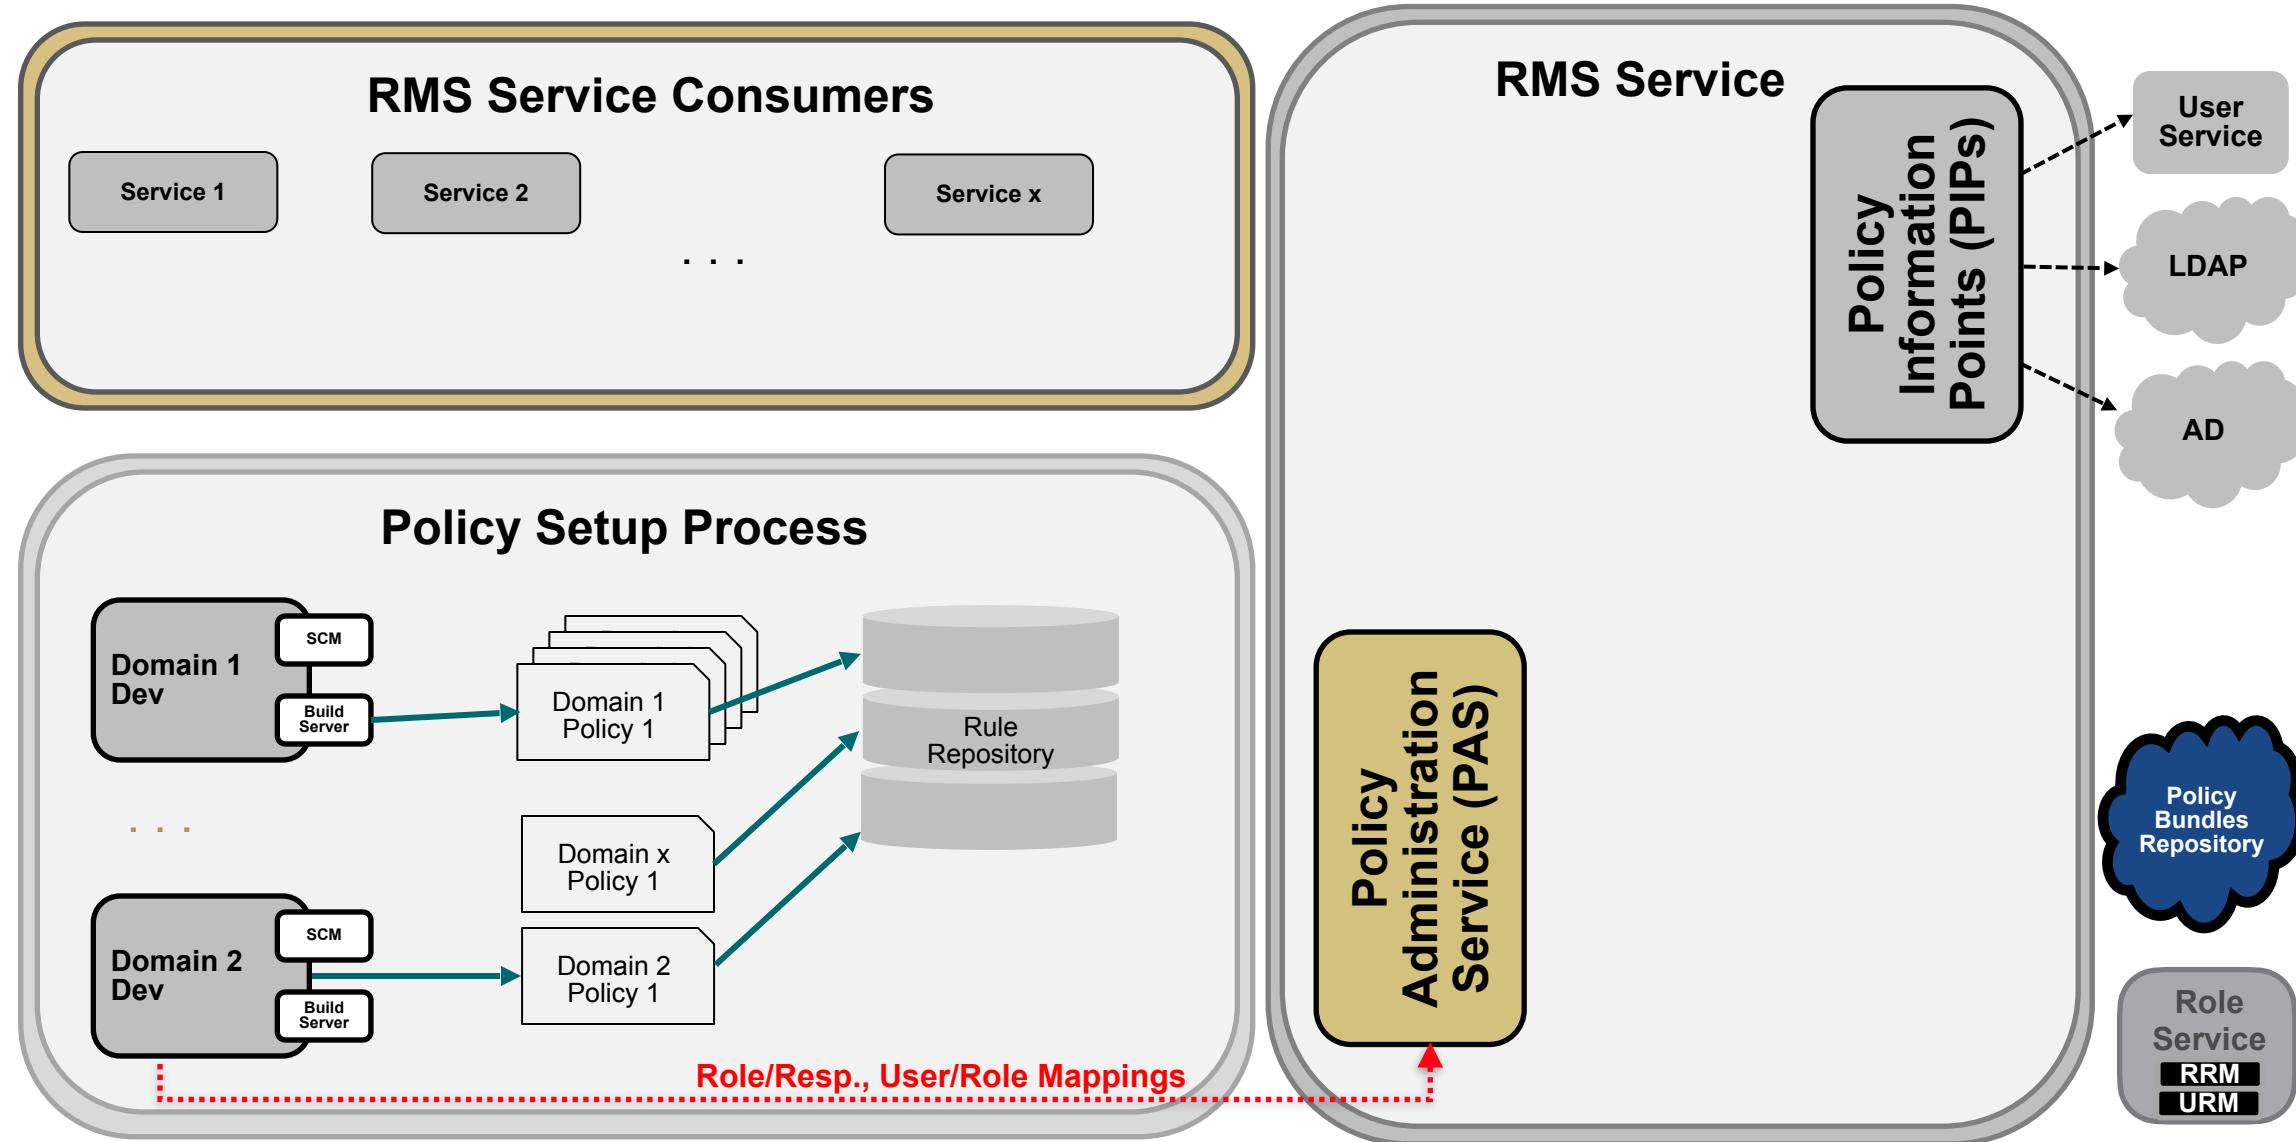
RMS Architecture - Version 2: Distributed



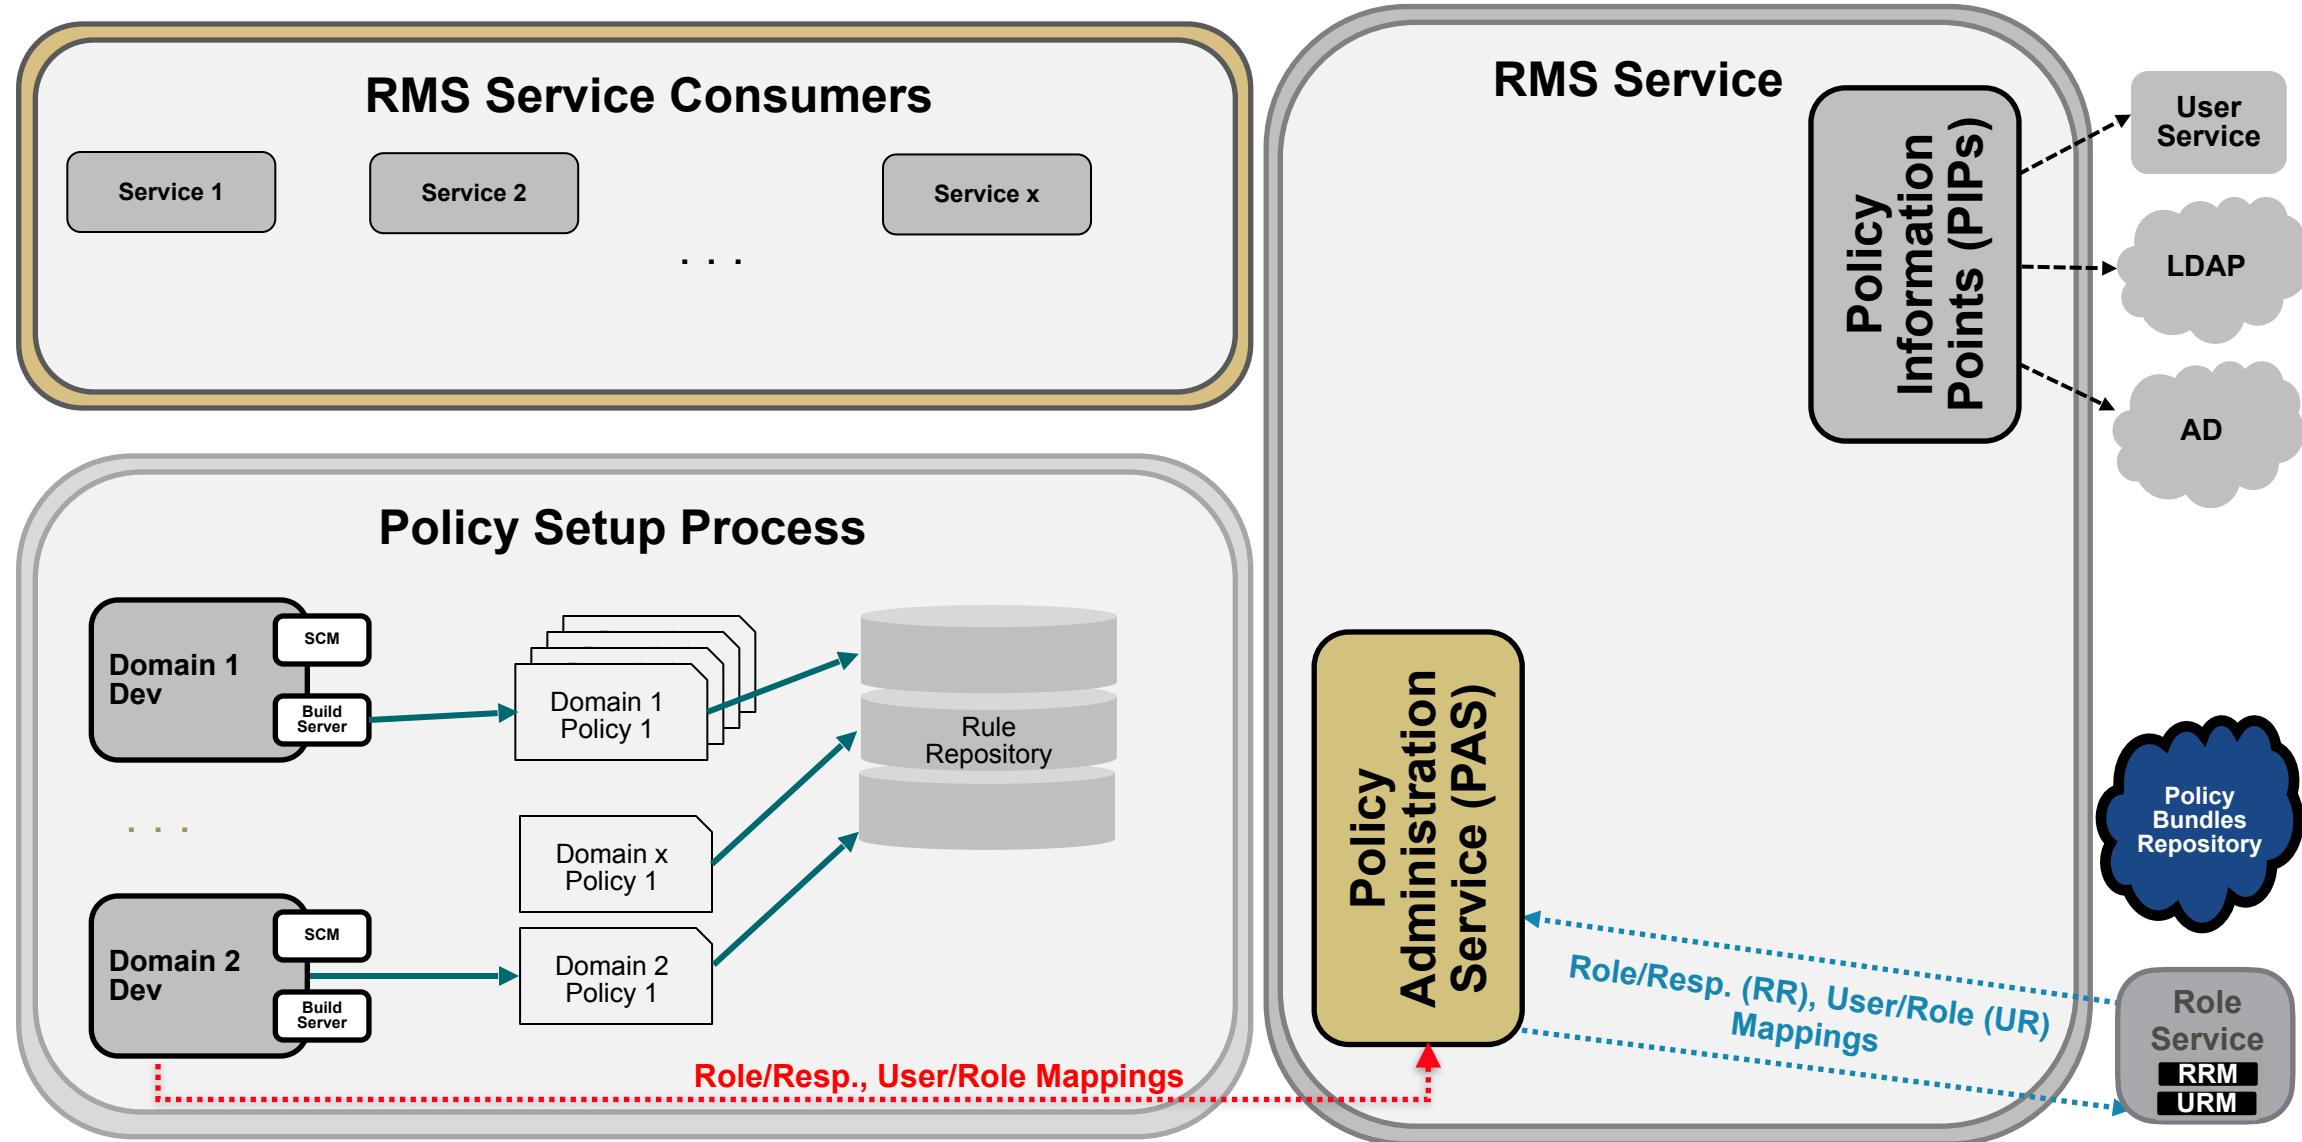
RMS Architecture - Version 2: Distributed



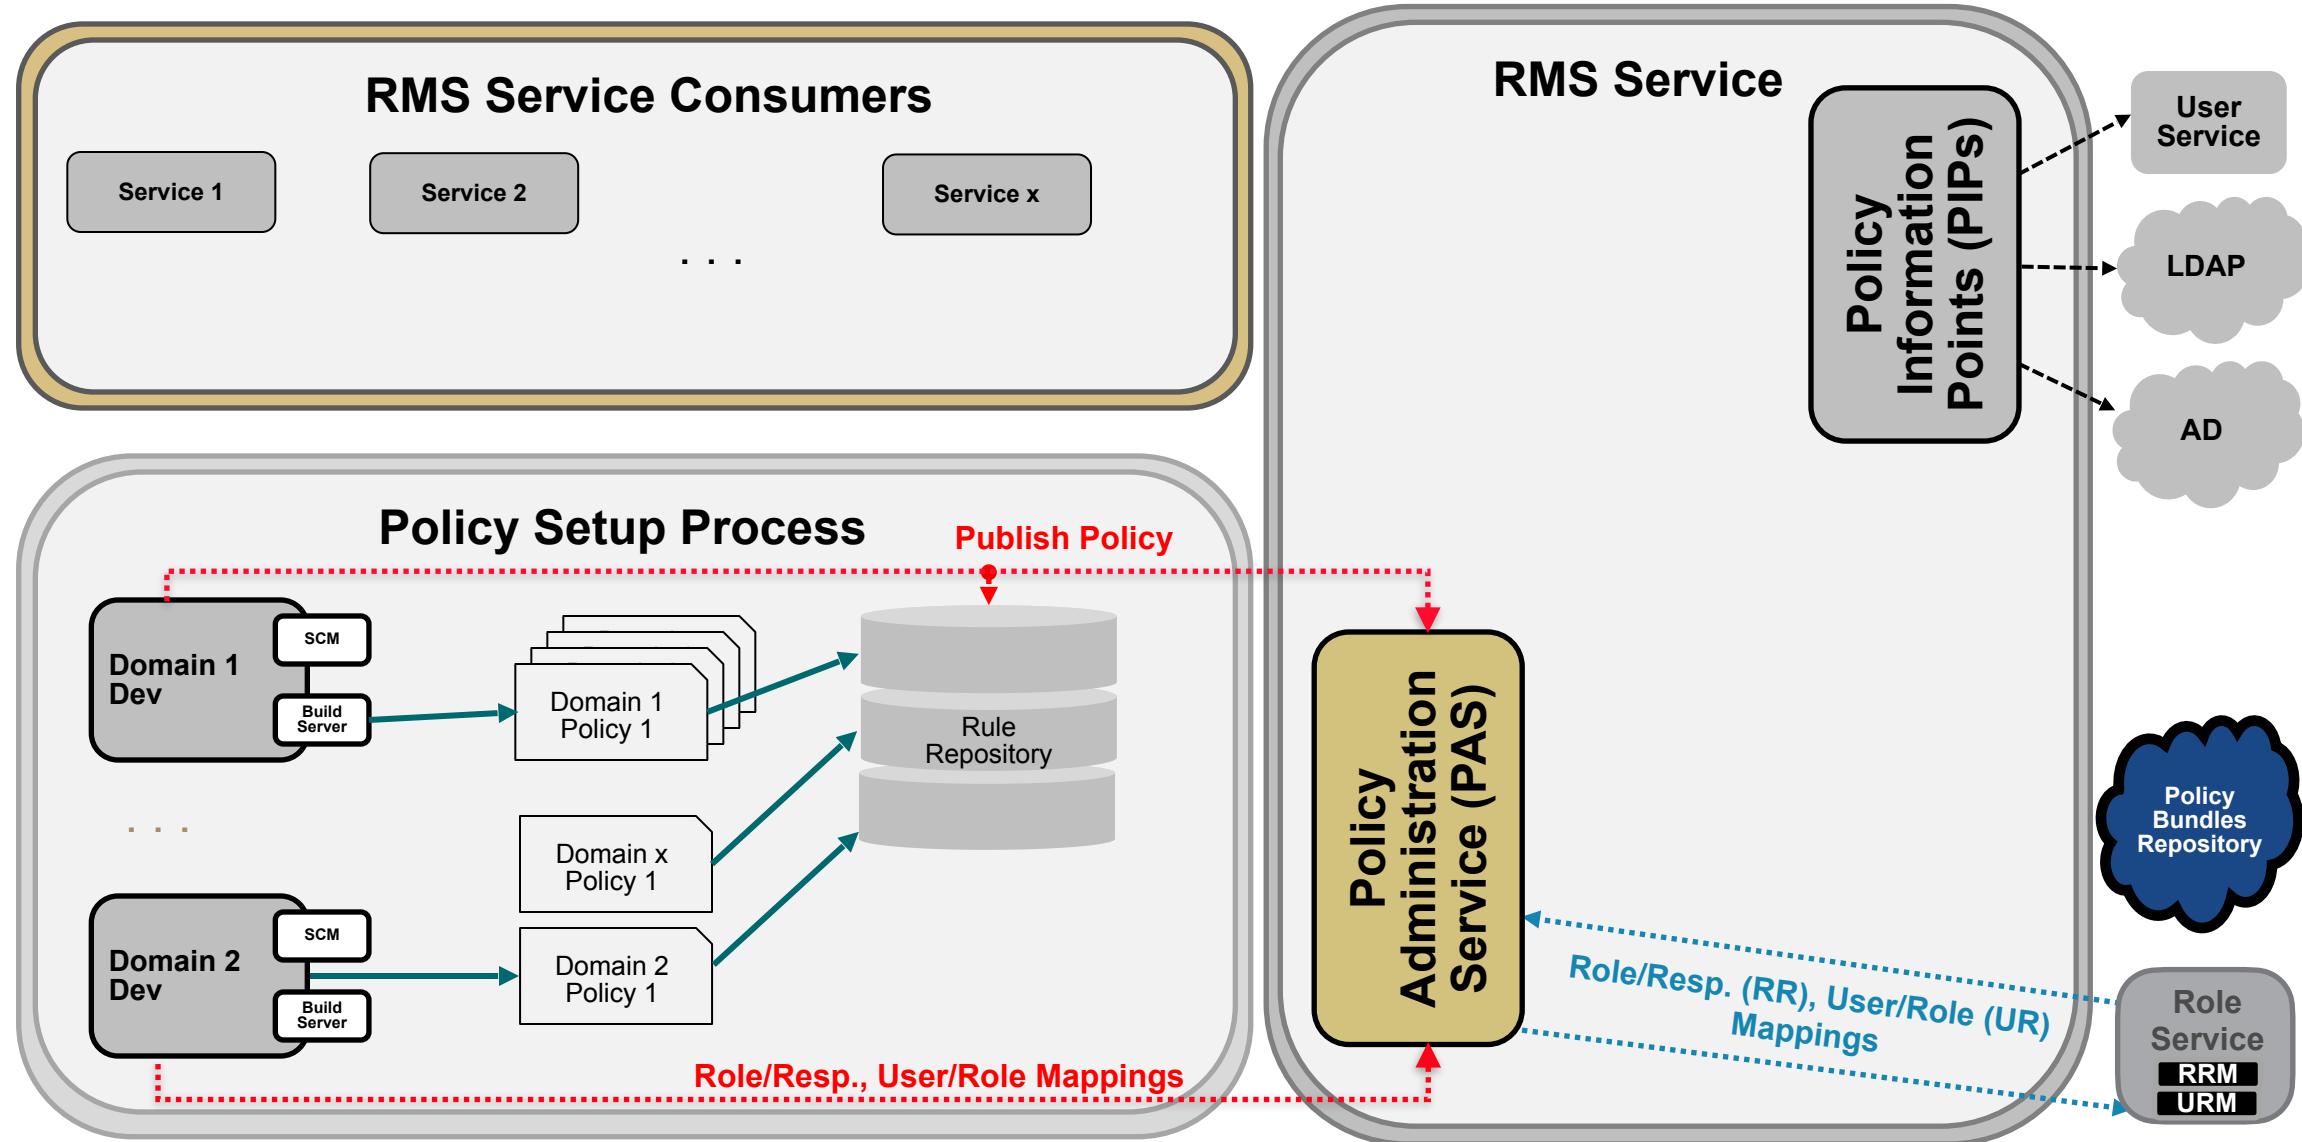
RMS Architecture - Version 2: Distributed



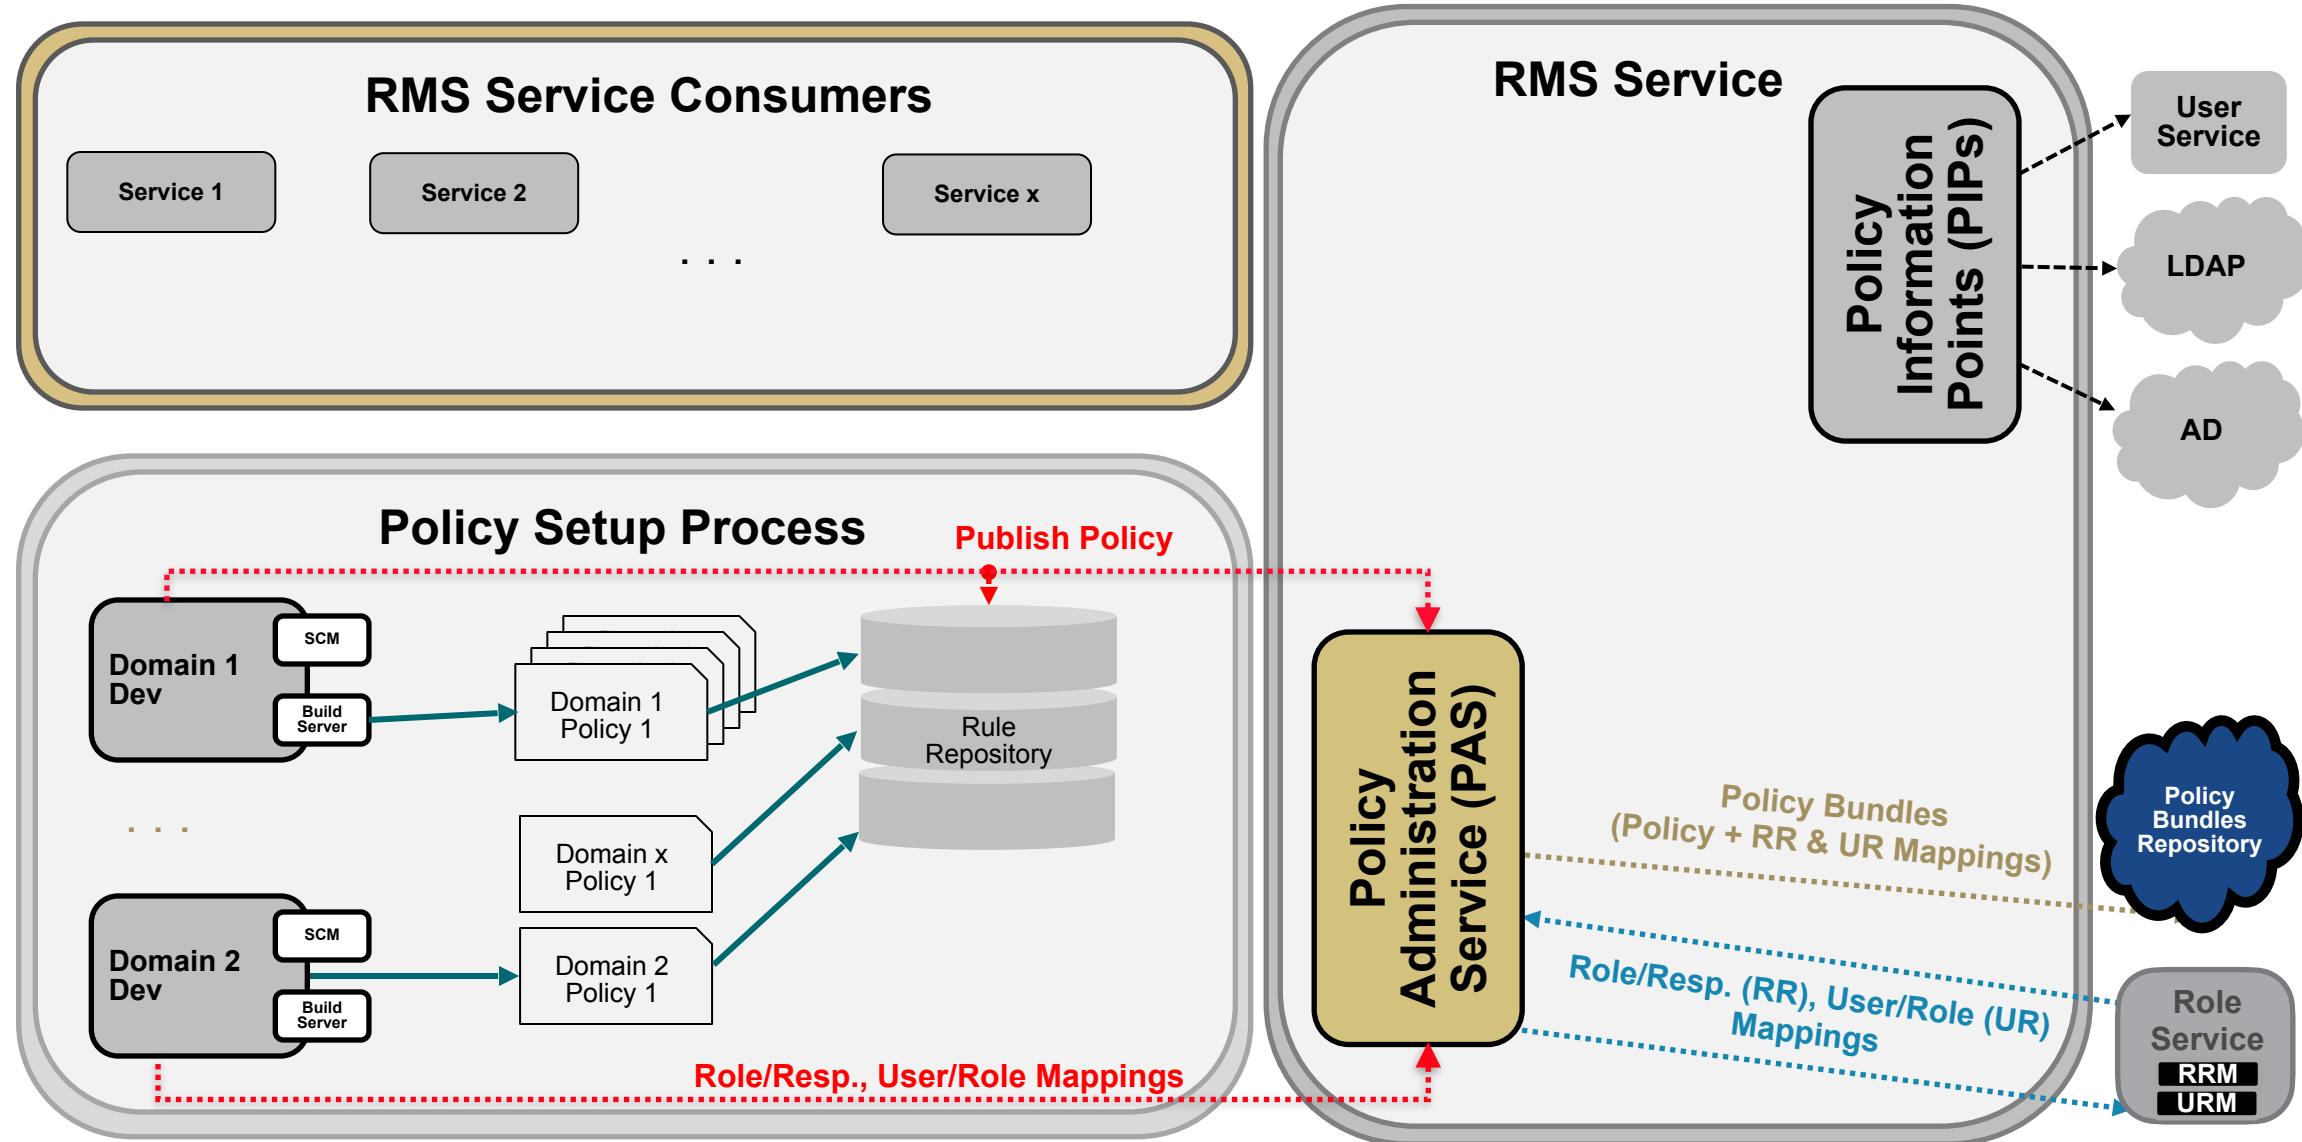
RMS Architecture - Version 2: Distributed



RMS Architecture - Version 2: Distributed



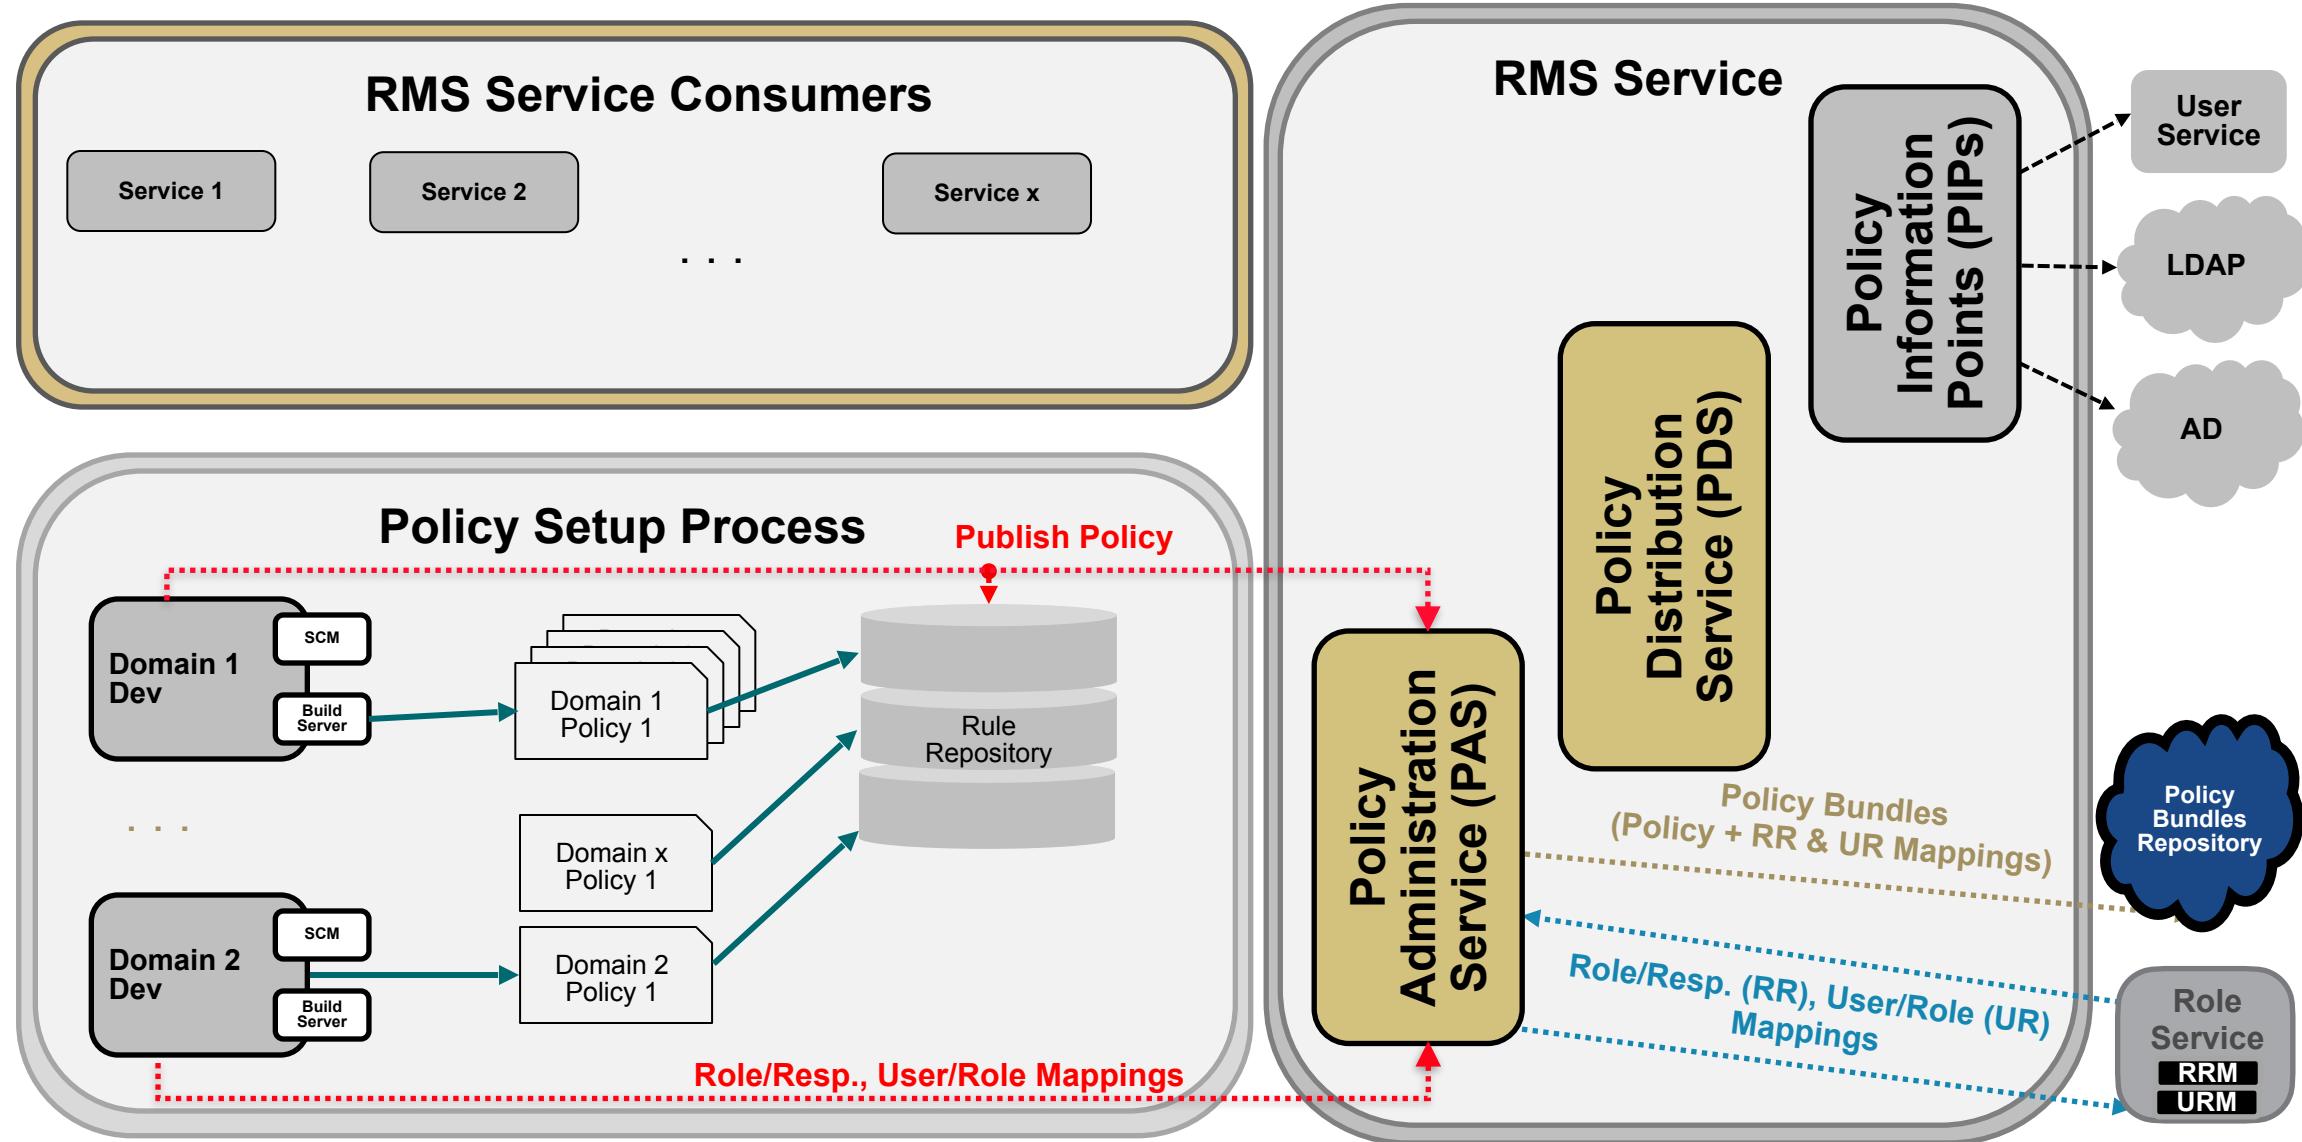
RMS Architecture - Version 2: Distributed



RMS Architecture - Version 2: Distributed

RRM Role Responsibility Mapping

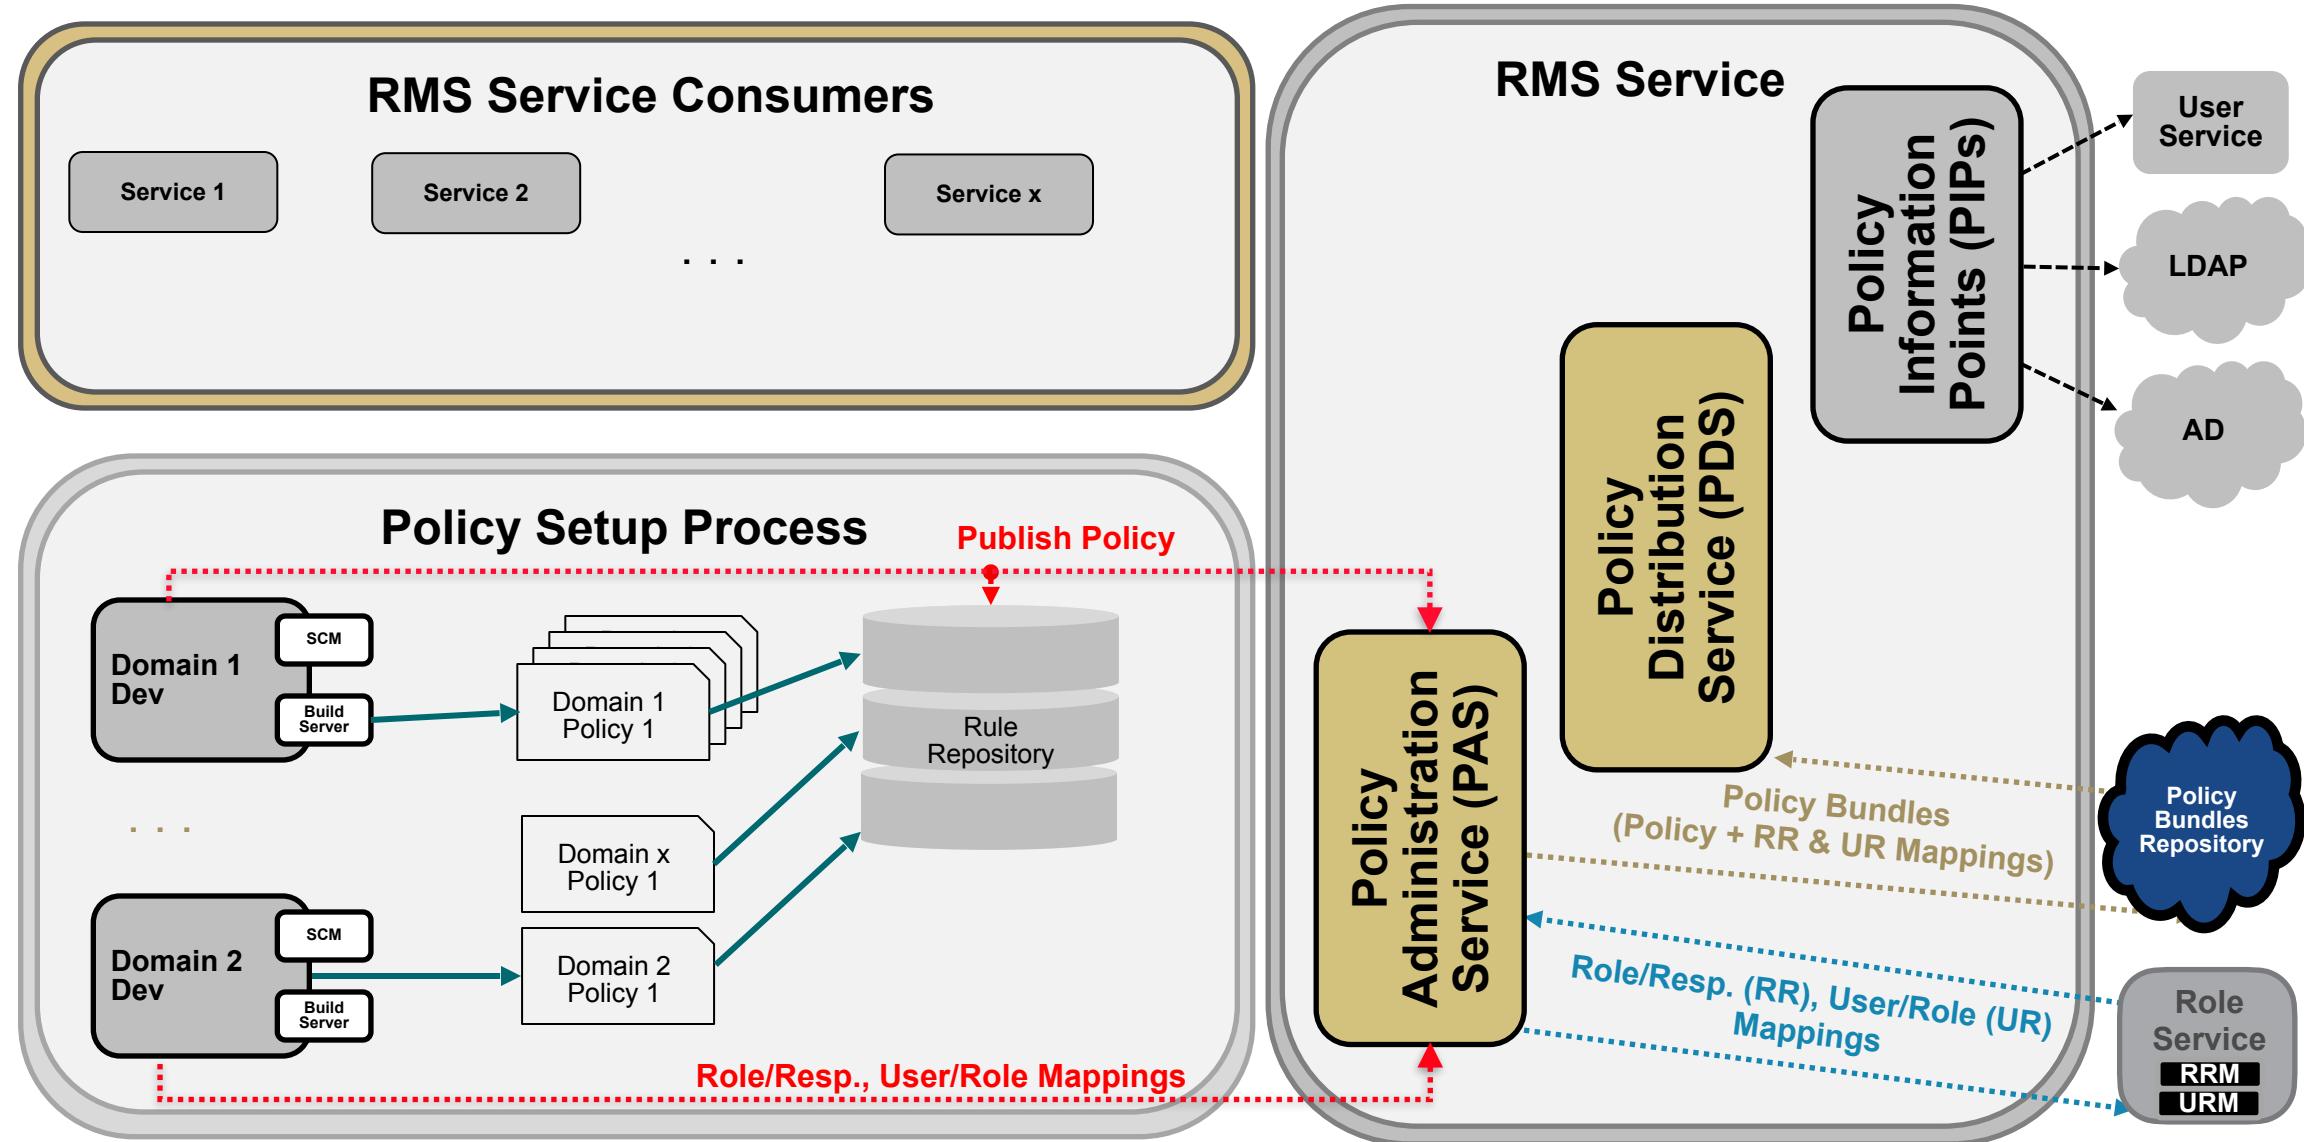
URM User Role Mapping



RMS Architecture - Version 2: Distributed

RRM Role Responsibility Mapping

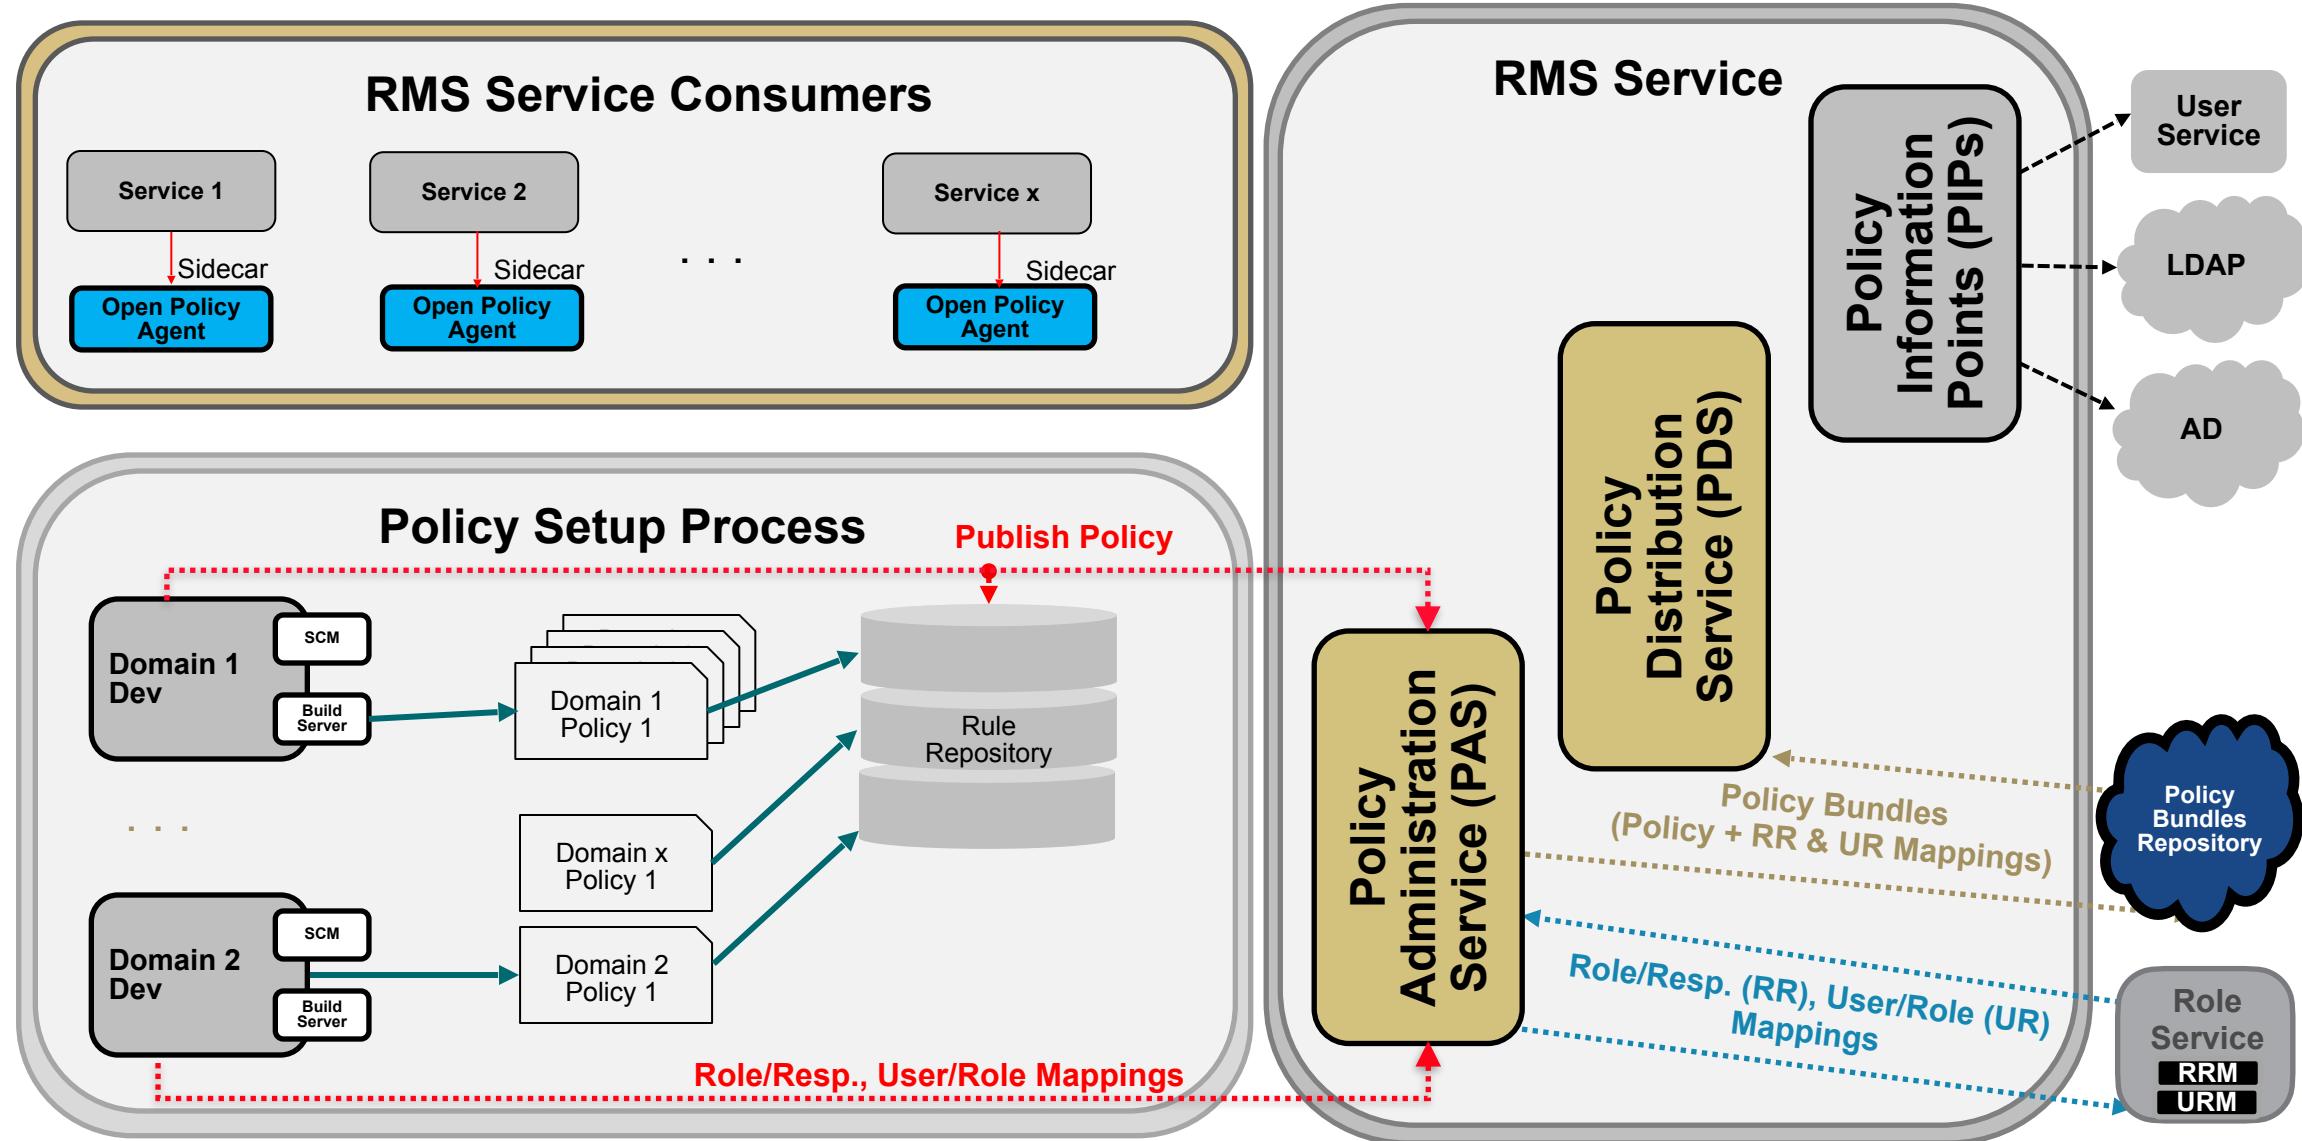
URM User Role Mapping



RMS Architecture - Version 2: Distributed

RRM Role Responsibility Mapping

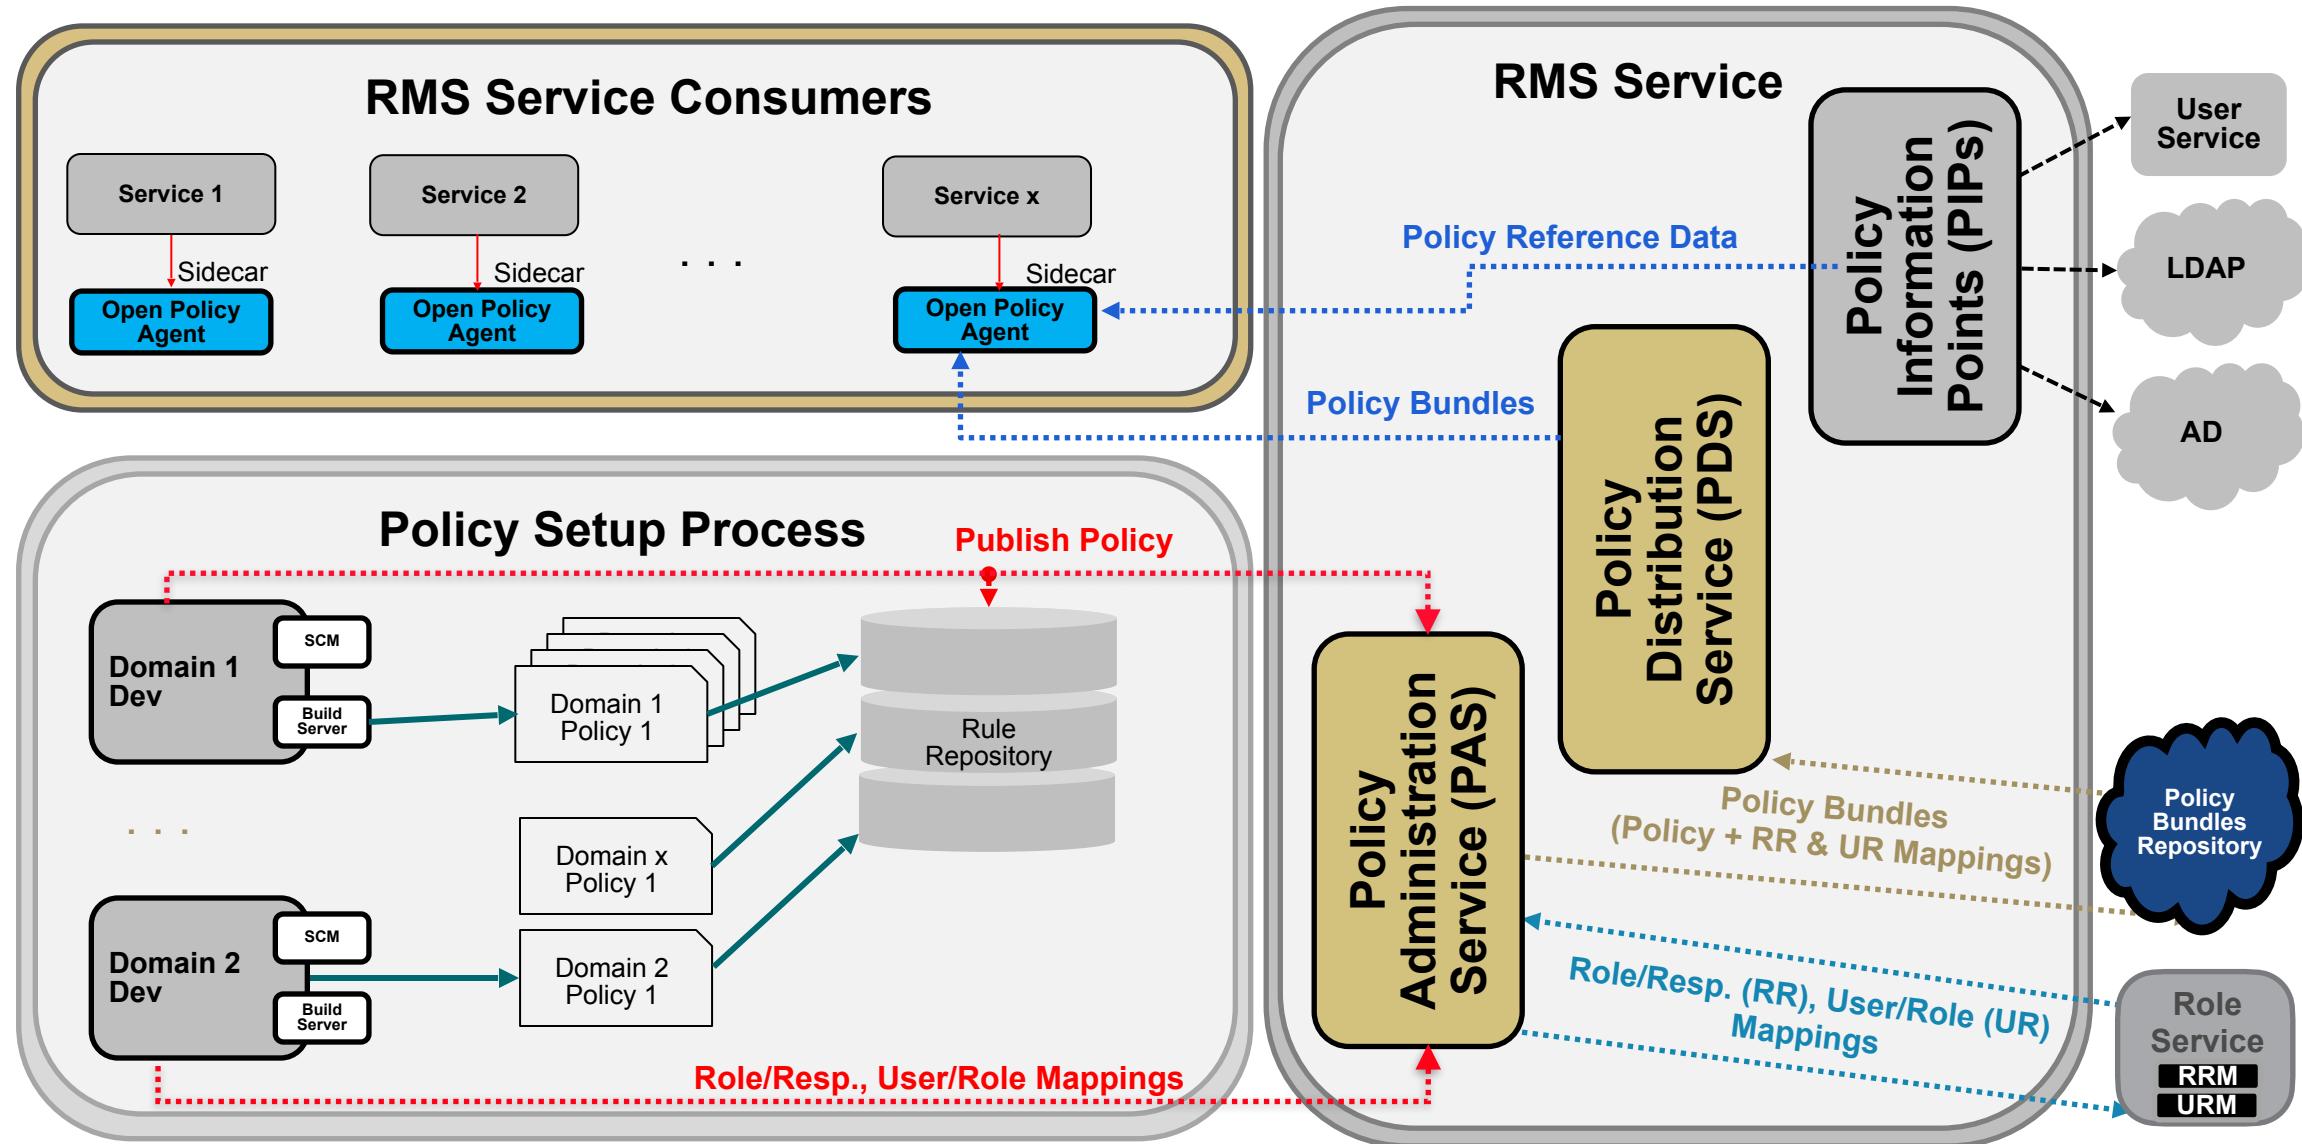
URM User Role Mapping



RMS Architecture - Version 2: Distributed

RRM Role Responsibility Mapping

URM User Role Mapping



Benefits of a Distributed Policy Management Architecture

Comparing Version 1 (federated single policy engine) with Version 2 (distributed policy engines)

	V1 Federated Policy Engine	V2 Distributed Policy Engine
Segregation and Information Barriers	Requires additional work	Is implicit, no additional work
Impact of a rogue policy script	Outage for all domains	Outage only for the specific domain
Gatekeeping for testing and coverage	Requires RMS to be the gatekeeper	Requires domain to be the gatekeeper
Strategy for new and updated policies	Needed a Release Train model	A domain can push policies on-demand
Impact of ad-hoc policy changes	RMS Downtime for all domains	RMS Downtime for the changed domain
Implicit RBAC Support	-	Available

Policy Bundles Repository

Policy bundles repository stored enriched policy archives.

Enriched policy bundles are archives that contain:

- Policy file(s), **specific to the domain**.
- Policy static data, **specific to the domain**.
- Standard RMS OPA policy rego files **common across all domains**.

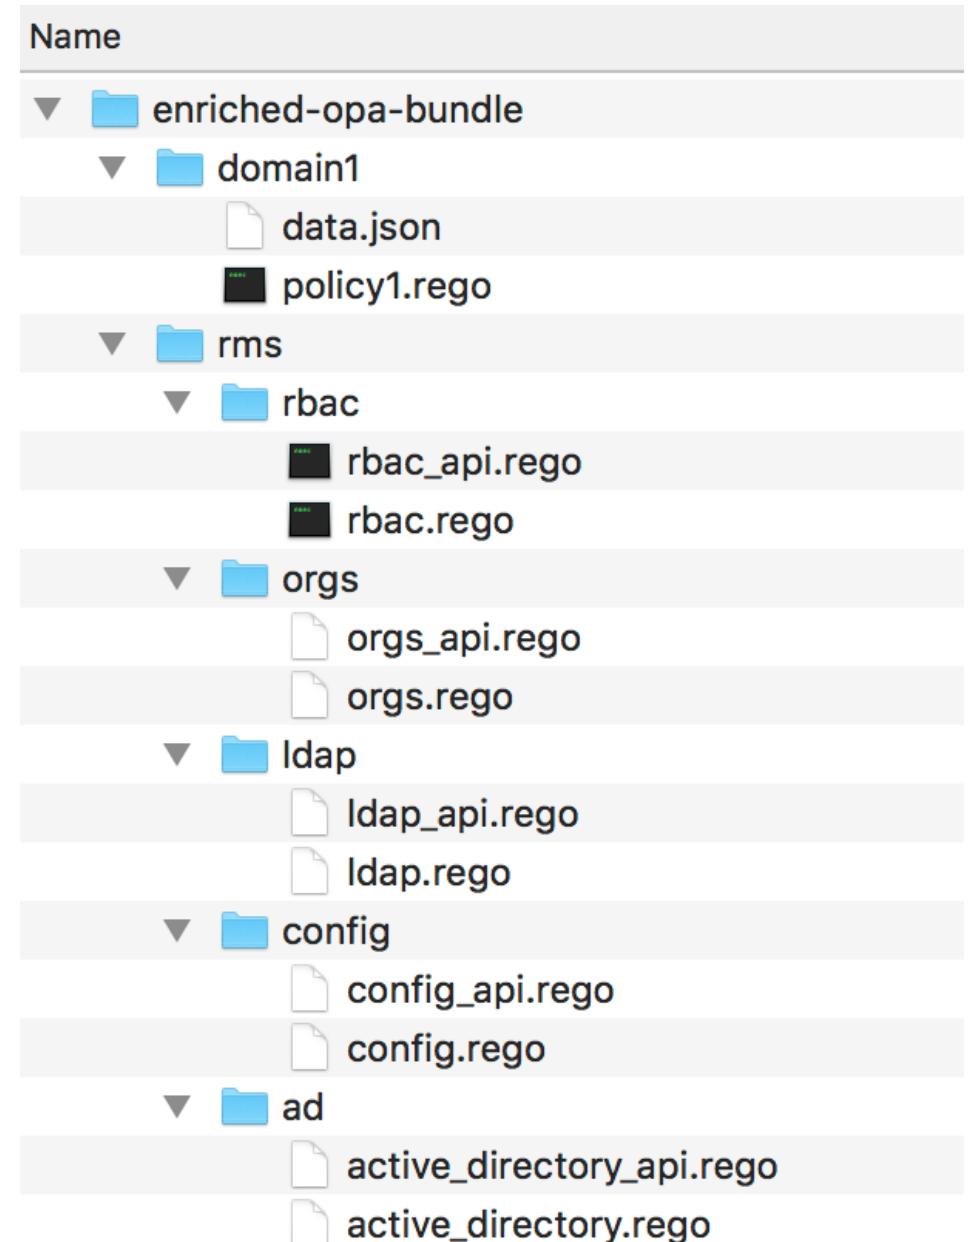
Policy Bundles Repository

Folder structure in policy bundles repository :

- <domain>
 - <policy>
 - <version>
 - <policy bundles>

Example:

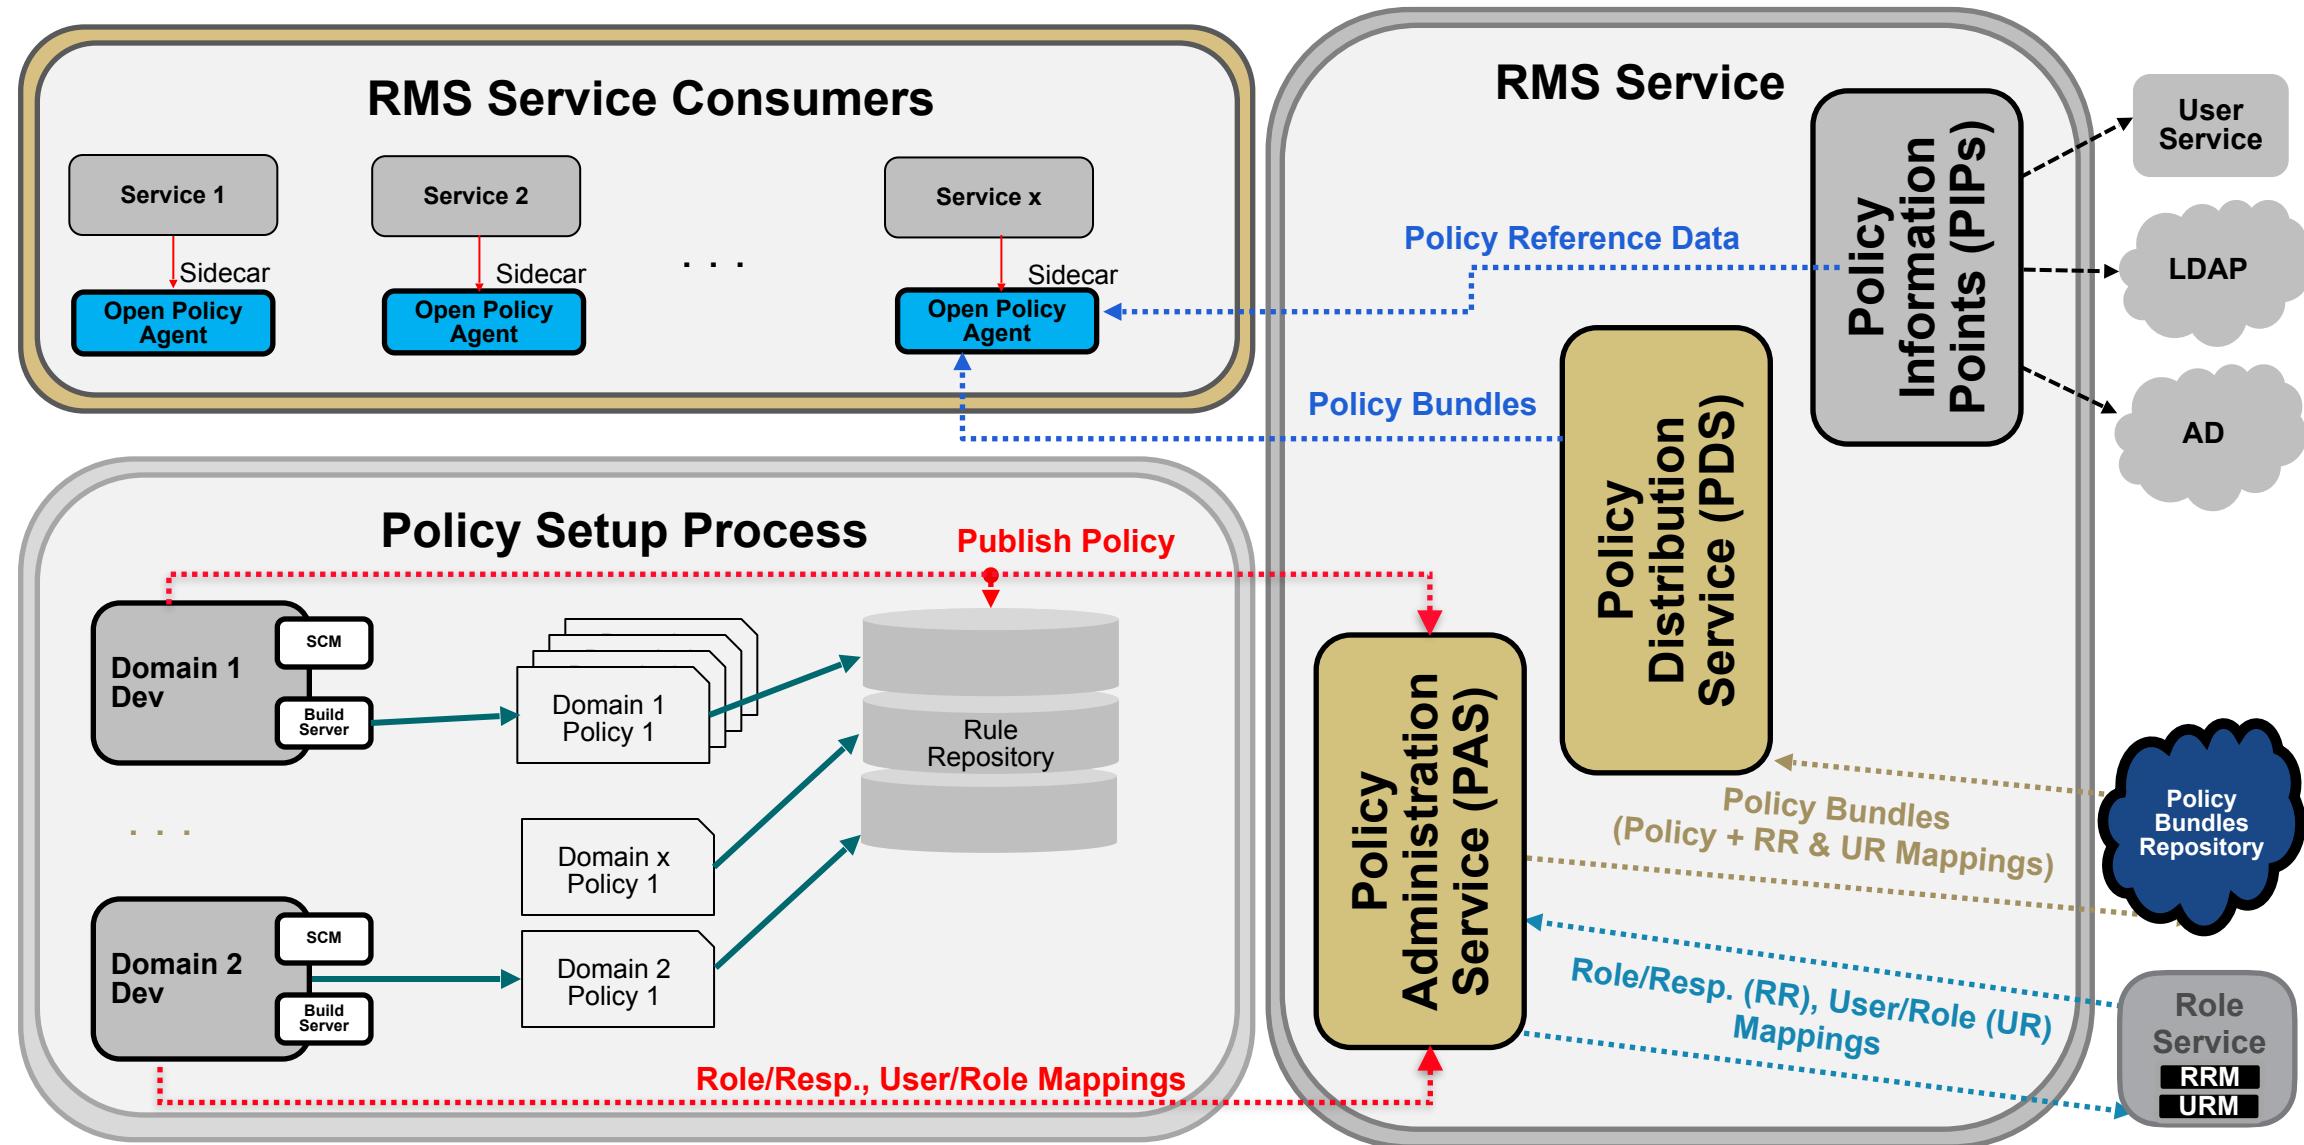
- domain1
 - policy1
 - 1.0.0
 - enriched-opa-bundle.tar.gz



RMS Architecture - Version 2: Distributed

RRM Role Responsibility Mapping

URM User Role Mapping



Policy Agent: Setup

- Open Policy Agent (the executable)
- Open Policy Agent – Configuration
- Open Policy Agent – Dockerfile command

Policy Agent: Setup: Configuration Files

OPA Configuration file (located at \${configPath})

```
services:
  - name: domainPolicies
    url: policyDistributionServiceUrl/
    allow_insecure_tls: true
bundle:
  name: policyDomain/policyName/policyVersion
  service: domainPolicies
  polling:
    min_delay_seconds: minDelaySeconds
    max_delay_seconds: maxDelaySeconds
```

Policy Agent: Setup: Dockerfile Commands

OPA launch command (used in the Dockerfile)

```
exec ./opa run --server --log-level=debug -c ${configPath}
```

RBAC Policy Library

```
package rbac

user_has_responsibility(userId, action, resource) {
    role := roles[_]

    responsibility := role.responsibilities[_]
    does_resource_match(resource, responsibility)

    responsibility.actions[_] = action

    is_user_a_member(userId, role)
}

is_user_a_member(userId, role) {
    ...
}
```

rbac.rego

Application Policy

```
package application1

default allow = false

allow {
    data.rbac.user_has_responsibility(
        input.userid, input.action,
        input.service)
}
```

policy.rego

Sample Role Data Excerpts

```
{
    "name": "App User",
    "responsibilities": [
        {
            "resource": "service.1",
            "actions": [
                "provision"
            ]
        },
        {
            "resource": "service.2",
            "actions": [
                "provision"
            ]
        }
    ],
    "members": [
        "EVERYONE"
    ]
}
```

data1.json

```
{
    "name": "App Admin",
    "responsibilities": [
        {
            "resource": "regexp:service\\..*",
            "actions": [
                "create",
                "update",
                "delete",
                "view"
            ]
        }
    ],
    "members": [
        "org:abc"
    ]
}
```

data2.json

IDE Productivity

An IDE plugin tool for the Open Policy Agent file authoring

Open Policy Agent: IntelliJ Plugin

- OPA IntelliJ Plugin is *functional* work-in-progress policy editor.
- The editor parses and validates OPA policy.
- Relies on the OPA language reference linked * below.
- Can be customized for editor color schemes in IntelliJ.
- Work continues on indentation, run configurations and variable tracking.

<https://www.openpolicyagent.org/docs/latest/language-reference/>

Open Policy Agent: IntelliJ Plugin: Before and After

```
test2.rego
1 package sample.policy
2
3 import datatestdata
4 import data.test_var
5
6 url = test_var("SAMPLE_GROUPS")
7
8 #
9 # Comments: sample policy
10 #
11
12 default authorize = "unknown"
13
14 authorize = "allow" {
15     input.user == "superuser"
16 } else = "deny" {
17     input.path[0] == "admin"
18     input.source_network == "external"
19 }
20
21 same_site[apps[k].name] {
22     #
23     #comments
24     #
25     apps[i].name == "mysql"
26     server := apps[i].servers[_]
27     server == sites[j].servers[_].name
28     other_server := sites[j].servers[_].name      #inline comment
29     server != other_server
30     other_server == apps[k].servers[_]
31 }
```

```
test2_rego
1 package sample.policy
2
3 import datatestdata
4 import data.test_var
5
6 url = test_var("SAMPLE_GROUPS")
7
8 #
9 # Comments: sample policy
10 #
11
12 default authorize = "unknown"
13
14 authorize = "allow" {
15     input.user == "superuser"
16 } else = "deny" {
17     input.path[0] == "admin"
18     input.source_network == "external"
19 }
20
21 same_site[apps[k].name] {
22     #
23     #comments
24     #
25     apps[i].name == "mysql"
26     server := apps[i].servers[_]
27     server == sites[j].servers[_].name
28     other_server := sites[j].servers[_].name      #inline comment
29     server != other_server
30     other_server == apps[k].servers[_]
```

Open Policy Agent: IntelliJ Plugin: Syntax Validation

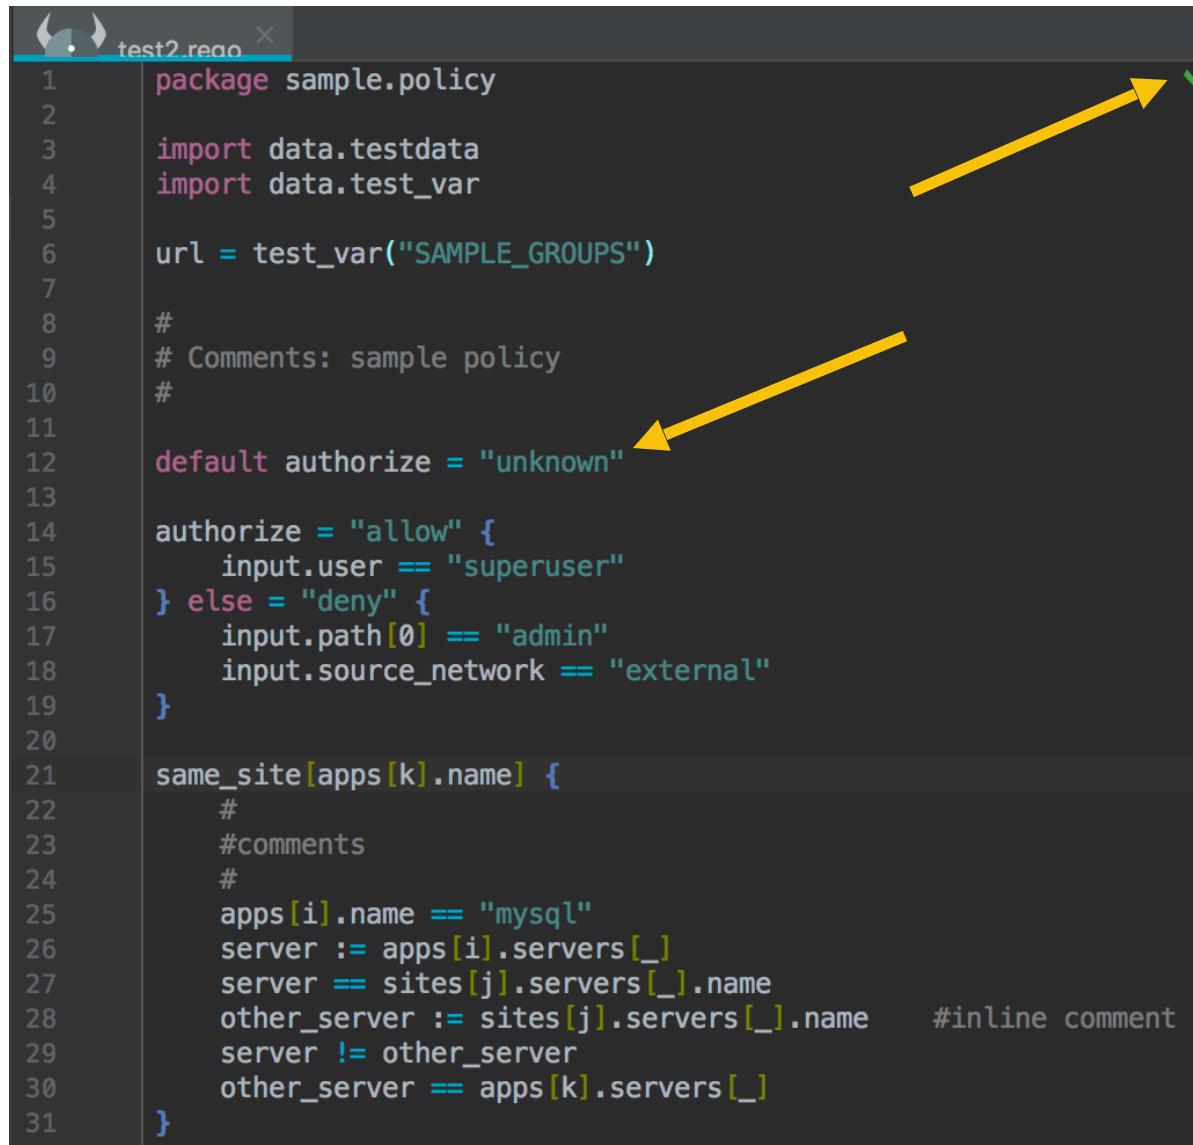
The screenshot shows the IntelliJ IDEA code editor with a dark theme. A file named `test2_reco` is open, containing Java-like pseudocode for an Open Policy Agent (OPA) policy. The code includes imports for `sample.policy`, `testdata`, and `test_var`. It defines a URL using a test variable and contains several comments. Two specific lines are highlighted with red arrows pointing to them from the left:

- Line 12: `default authorize = "unknown"`
- Line 14: `authorize = "allow" {`

A tooltip is displayed over the second highlight, showing the error message: `('[, <rule body>, <rule>, <set sign>, OPATokenType.COMMENTS, OPATokenType.policy_declaration_0_3_0, OPATokenType.var or '[' expected, got "" { input.user =...'`. This indicates that the parser expected a closing brace or bracket but found an unexpected string literal instead.

```
1 package sample.policy
2
3 import data.testdata
4 import data.test_var
5
6 url = test_var("SAMPLE_GROUPS")
7
8 #
9 # Comments: sample policy
10 #
11
12 default authorize = "unknown"
13
14 authorize = "allow" {
15
16     sets - deny {
17         input.path[0] == "admin"
18         input.source_network == "external"
19     }
20
21     same_site[apps[k].name] {
22         #
23         #comments
24         #
25         apps[i].name == "mysql"
26         server := apps[i].servers[_]
27         server == sites[j].servers[_].name
28         other_server := sites[j].servers[_].name      #inline comment
29         server != other_server
30         other_server == apps[k].servers[_]
31     }
32 }
```

Open Policy Agent: IntelliJ Plugin: Syntax Validation - Continued



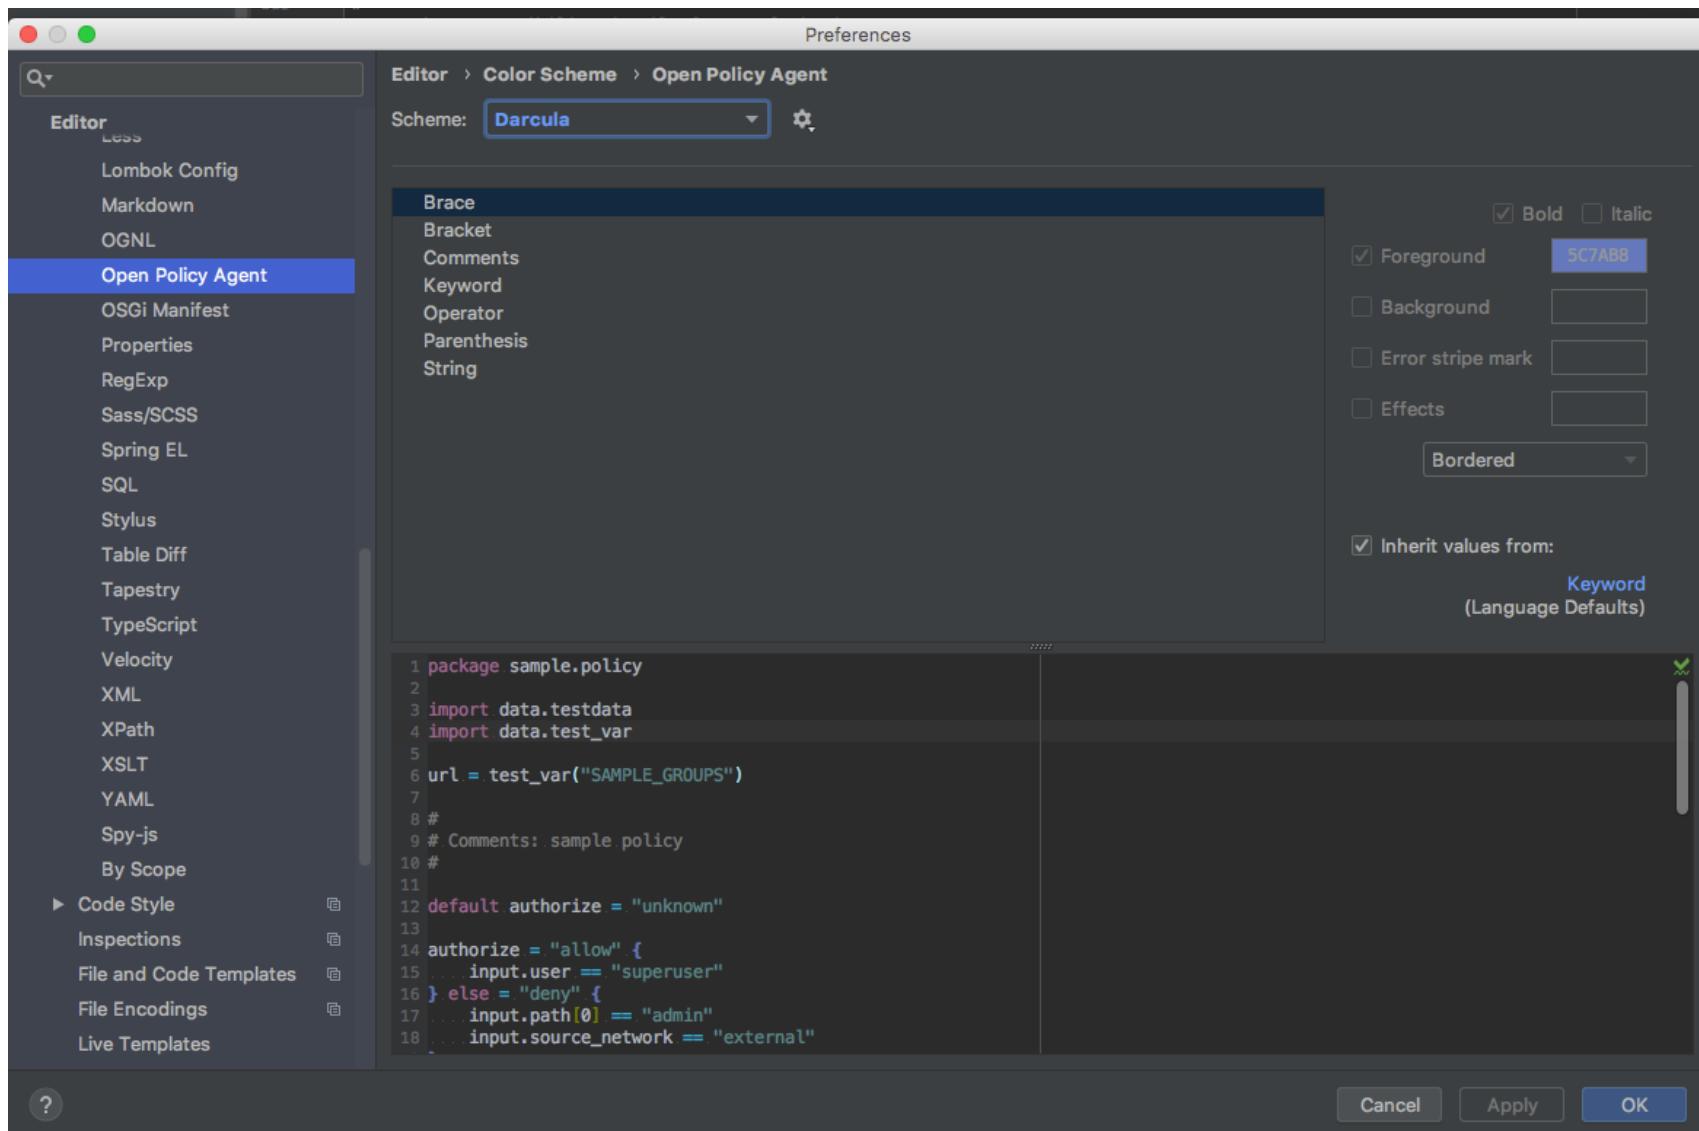
A screenshot of an IntelliJ code editor window titled "test2.reao". The code is written in a policy language, likely Open Policy Agent (OPA) or a similar rule-based system. The code includes imports, URLs, comments, and policy rules. A yellow arrow points from the word "authorize" in line 12 to a green checkmark icon in the top right corner of the editor, indicating that the code is valid according to the plugin's syntax checker.

```
1 package sample.policy
2
3 import datatestdata
4 import data.test_var
5
6 url = test_var("SAMPLE_GROUPS")
7
8 #
9 # Comments: sample policy
10 #
11
12 default authorize = "unknown" ← Yellow arrow points here
13
14 authorize = "allow" {
15     input.user == "superuser"
16 } else = "deny" {
17     input.path[0] == "admin"
18     input.source_network == "external"
19 }
20
21 same_site[apps[k].name] {
22     #
23     #comments
24     #
25     apps[i].name == "mysql"
26     server := apps[i].servers[_]
27     server == sites[j].servers[_].name
28     other_server := sites[j].servers[_].name      #inline comment
29     server != other_server
30     other_server == apps[k].servers[_]
31 }
```

Open Policy Agent: IntelliJ Plugin: Settings

Select

- Preferences
 - Editor
 - > Color Scheme
 - Open Policy Agent



Summary

- Responsibility Management as a Service can resolve issues on several fronts.
- Choice of a payload format (HOCON over JSON or XML) can help control verbosity.
- Choice of architecture (federated versus distributed) can help determine resilience.
- Distributed policy engines can alleviate back-pressure and volume demands.
- Distributed policy engines can reduce outages and maintenance-related downtimes.
- Creating a policy editor plugin can help boost productivity.

Links

- HOCON

<https://github.com/lightbend/config/blob/master/HOCON.md>

- Eclipse Collections

<https://www.eclipse.org/collections/>

- Open Policy Agent

<https://www.openpolicyagent.org/>





Appendix: Responsibility Management Cycle

